

Curriculum Bible and Visual Learning Plan

This pass paused feature work and produced the curriculum architecture package for the next implementation pass. It did not add new games or polish gameplay.

35

proposed Journey lessons

8

new curriculum docs

7

new lesson drafts

What Changed

- Created the v0.7 Curriculum Bible covering current and proposed lessons.
- Proposed a 35-lesson architecture with Essential Path, Deep Path, and Optional Side Tour labels.
- Drafted complete cards for Overfitting and Generalization, Alignment, Hallucinations, Grounding, Effective Prompting, Energy/Water/Compute, and Human-Centered AI Ethics.
- Created the visual-learning framework, lesson-by-lesson visual specs, metaphor/Brain Bridge matrix, misconception map, and source-needs list.
- Updated review notes with files changed, source gaps, recommendations, and build result.

Files Added or Updated

- `scripts/generate-curriculum-v07.mjs`
- `docs/curriculum/PROMPT_LIFE_CURRICULUM_BIBLE_V0_7.md`
- `docs/curriculum/JOURNEY_ARCHITECTURE_V0_7.md`
- `docs/curriculum/NEW_LESSON_DRAFTS_V0_7.md`
- `docs/curriculum/VISUAL_LEARNING_FRAMEWORK_V0_7.md`
- `docs/curriculum/VISUAL_AID_SPECS_V0_7.md`
- `docs/curriculum/METAPHOR_AND_BRAIN_BRIDGE_MATRIX_V0_7.md`
- `docs/curriculum/MISCONCEPTION_MAP_V0_7.md`
- `docs/curriculum/SOURCE_NEEDS_V0_7.md`

- docs/REVIEW_NOTES.md

Verification

- npm run typecheck: passed.
- npm run build: passed.
- Generated docs include 35 Bible lesson sections and 35 visual-aid spec rows.
- No UI files or game logic were changed in this pass.

Screenshots below are current UI reference shots from the existing review/RAG interface. This v0.7 pass changed curriculum documentation rather than app screens.

Reference Screenshots

PROMPT LIFE V0.6.3

Lesson Cards

One accessible curriculum-review card per lesson, built from current app data plus inventory notes and rubric scores.

1 / 26

BEFORE MORNING

What Is an LLM?

A learned next-token prediction system

Stage: architecture

Priority: low

Rubric avg: 4.4/5

The diagram illustrates the LLM architecture. A central light green circle labeled 'weights' is connected by horizontal lines to a light blue rounded rectangle on the left labeled 'prompt' and a light orange rounded rectangle on the right labeled 'next'. Below the 'weights' circle are four smaller light green circles, each with a label underneath: 'grammar', 'facts', 'tone', and 'task'. The entire diagram is set against a dark blue background with a subtle grid pattern.

Prompt to Prediction

A prompt enters the current context; learned weights shape next-token probabilities.

DEFINITION

A large language model is a learned system that predicts likely next tokens from the context it can see.

Review route used as the current curriculum-audit surface.

RAG and Retrieval



Retrieved text joins the current context

Retrieval-augmented generation, or RAG, retrieves outside information and places it into the model's context before generating a response.




VISUAL AID

What to picture



RAG and Retrieval is already a dedicated Journey card and remains central in the v0.7 architecture.



Grounding exists in the glossary today and becomes a dedicated Deep Path lesson draft in v0.7.