

Generated Visual Copy Cleanup

This pass removes internal asset-generation language from learner-facing visual copy while preserving implementation metadata on review routes.

Date: June 6, 2026. Project: `KevinHegg/promptlife`.

What Changed

- Replaced generated-image-backed learner captions for What Is an LLM?, Pretraining, Fine-Tuning, and Alignment.
- Updated generated visual accessible descriptions to describe the concept and image, not the production process.
- Updated the Pretraining callouts and key takeaway to clarify broad durable pattern learning without perfect source recall.
- Changed the learner fallback heading from `Visual asset unavailable.` to `Visual unavailable.`
- Bumped the visible app version to `v0.21.1`.
- Added learner-facing copy guardrails for future visual and generated asset work.

Files Changed

- `src/components/VisualAids.tsx`
- `src/data/visualAssets.ts`
- `src/main.tsx`
- `docs/REVIEW_NOTES.md`
- `docs/design/LEARNER_FACING_COPY_GUARDRAILS_V0_21_1.md`
- `docs/design/screenshots/v0-21-1-*.png`
- `docs/design/screenshots/v0-21-1-screenshot-manifest.json`
- `docs/design/prompt-life-v0-21-1-generated-visual-copy-cleanup-report.html`
- `docs/design/prompt-life-v0-21-1-generated-visual-copy-cleanup-report.pdf`

Phrases Removed

A textless model-cloud scene is paired with HTML callouts

A textless folded landscape is paired with HTML callouts

A textless targeted-path scene is paired with HTML callouts

A textless alignment garden is paired with HTML callouts

textless conceptual image

surrounding HTML callouts

Visual asset unavailable.

New Pretraining Copy

Caption

Many examples flow through the training loop. Repeated updates shape weights into broad patterns the model can use later. This does not make the model a perfect memory of its sources.

Callouts

- **Many examples flow:** Pretraining repeats the training loop across enormous collections of examples.
- **Updates shape weights:** Prediction errors drive repeated weight updates during training.
- **Broad patterns form:** The model learns statistical patterns that can help later prompts.
- **Not perfect recall:** Pretraining does not store every source as a searchable memory.

Key Takeaway

Pretraining builds broad capability, not perfect source recall.

Other Visual Copy Changes

- **What Is an LLM?** now explains current context, learned weights, next-token probabilities, generated token, and append-repeat.
- **Fine-Tuning** now explains pretrained base, targeted examples or preference data, and durable future response shaping.
- **Alignment** now explains preferred behavior, guardrails, feedback and policies, and the no-conscience limit.
- **Review gallery** still exposes asset metadata for QA and documentation.

Verification

- **Passed:** `npm run typecheck`
- **Passed:** `npm run build` with existing Vite large-chunk warning.
- **Passed:** `npm run build:pages` with existing Vite large-chunk warning.
- **Passed:** `npm run audit:answers`
- **Passed:** `git diff --check`
- **Screenshot QA:** learner visual panels reported `badLearnerCopy: false`, `overflowX: false`, and isolated progress storage `[]`.

Known Issues

- The Vite large-chunk warning remains from the existing bundle shape.
- The review route intentionally contains production metadata. That is correct for QA, but it should stay out of normal Journey screens.

Next Step

The cleanup is complete. The next sensible product step is the Decision Room implementation pass, using coded visuals and revised lesson architecture without adding generated assets yet.

Screenshots



Prompt to Prediction

Score, choose, append, repeat

Current context enters the model, learned weights shape next-token probabilities, one token is generated, and the response grows token by token.

- 1 Context enters** The current prompt and response-so-far become the visible input for this run.
- 2 Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- 3 One token is generated** The model chooses one response token from the probability cloud.
- 4 Response grows** That token is appended, then the model runs again.

Key takeaway:

An LLM turns context into next-token probabilities.

Key terms

Tap a term to open the glossary.

LLM

Prompt

Response

Token

Weight

Inference

CORE IDEA

An LLM does not read a prompt and write a whole answer at once. For the prompt "Give



Home



Journey



Play




Glossary



Badge

What Is an LLM: learner-facing visual copy at 390px.



Broad Pretraining

Scale, not perfect recall

Many examples flow through the training loop. Repeated updates shape weights into broad patterns the model can use later. This does not make the model a perfect memory of its sources.

- 1 Many examples flow** Pretraining repeats the training loop across enormous collections of examples.
- 2 Updates shape weights** Prediction errors drive repeated weight updates during training.
- 3 Broad patterns form** The model learns statistical patterns that can help later prompts.
- 4 Not perfect recall** Pretraining does not store every source as a searchable memory.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Key terms

Tap a term to open the glossary.

[Show all terms](#)

Pretraining

Training

Training data

Weight

Loss

CORE IDEA

Pretraining is the same training loop at



Home



Journey



Play



Glossary



Badge

Pretraining: cleaned caption, callouts, and key takeaway at 390px.



Broad Pretraining

Scale, not perfect recall

Many examples flow through the training loop. Repeated updates shape weights into broad patterns the model can use later. This does not make the model a perfect memory of its sources.

- 1 Many examples flow** Pretraining repeats the training loop across enormous collections of examples.
- 2 Updates shape weights** Prediction errors drive repeated weight updates during training.
- 3 Broad patterns form** The model learns statistical patterns that can help later prompts.
- 4 Not perfect recall** Pretraining does not store every source as a searchable memory.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Key terms

[Show all terms](#)

Pretraining

Training

Training data

Weight



Home



Journey



Play



Glossary



Badge

Pretraining: cleaned copy remains readable at 320px.



Fine-Tuning Path

Durable adaptation after pretraining

Fine-tuning starts from a pretrained base, adds targeted examples or preference data, and shapes future responses more durably than one prompt.

- 1 Pretrained base** Fine-tuning begins with a model already shaped by broad pretraining.
- 2 Targeted data** Examples, domain data, or preference data guide the additional training step.
- 3 Future responses** The training changes how the model is likely to behave later.
- 4 More than a prompt** Fine-tuning is durable adaptation, unlike one temporary prompt or retrieved document.

Key takeaway:

Fine-tuning adapts future behavior through additional training.

Key terms

Tap a term to open the glossary.

[Show all terms](#)

Fine-tuning

Adapter

Pretraining

Training

Prompt

CORE IDEA

Fine-tuning happens after pretraining. It adapts a model toward a domain, task, style, or preference by changing weights or



Home



Journey



Play



Glossary



Badge

Fine-Tuning: learner-facing visual copy at 390px.



Alignment Landscape

Shaping behavior, not conscience

Alignment encourages preferred behavior, uses guardrails to reduce risky paths, and relies on feedback and policies without giving the model a conscience.

- 1 Preferred behavior** Training and system design can encourage more useful response patterns.
- 2 Guardrails** Runtime safeguards can reduce risky paths during use.
- 3 Feedback and policies** Human feedback, evaluations, policies, and filters can shape behavior.
- 4 No conscience** Alignment does not give the model moral understanding or a truth guarantee.

Key takeaway:

Alignment shapes behavior, but it is not magic morality.

Key terms

Tap a term to open the glossary.

[Show all terms](#)

Alignment

Instruction tuning

Human feedback learning

RLHF

CORE IDEA

Alignment shapes model behavior toward human instructions, preferences, safety



Home



Journey



Play



Glossary



Badge

Alignment: learner-facing visual copy at 390px.

BEFORE-MORNING-PRETRAINING-LANDSCAPE

Broad Pretraining

Lesson: Pretraining

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-pretraining-landscape.png

Type: generated-image



Broad Pretraining

Scale, not perfect recall

Many examples flow through the training loop. Repeated updates shape weights into broad patterns the model can use later. This does not make the model a perfect memory of its sources.

- 1 Many examples flow** Pretraining repeats the training loop across enormous collections of examples.
- 2 Updates shape weights** Prediction errors drive repeated weight updates during training.
- 3 Broad patterns form** The model learns statistical patterns that can help later prompts.
- 4 Not perfect recall** Pretraining does not store every source as a searchable memory.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Review gallery: asset filename/type metadata remains visible.

LEARNING OBJECTIVE

Show broad data exposure shaping durable model weights before ordinary use.

ASSET PATH

/assets/generated/before-morning/before-morning-pretraining-landscape.png

ALT TEXT

Streams of abstract paper fragments and light rain carve glowing paths into a folded-paper landscape.

ACCESSIBLE DESCRIPTION

Many abstract paper fragments and light streams shape a folded-paper landscape. The explanation emphasizes that many examples flow through training, repeated updates shape weights, broad patterns form, and pretraining is not perfect source memory.

MOBILE PREVIEW

Review at 320px, 390px, and 430px; images must fit the paper frame and callouts must remain readable below the visual.

TEXT HANDLING

Instructional labels and callouts are HTML. The PNG remains textless concept art.

PRINT/PDF NOTES

Provided textless generated PNG. Instructional labels, callouts, captions, accessibility text, and key takeaways live in HTML/SVG/app code, not inside the image file.

OVERFITTING-GENERALIZATION

Overfitting vs Generalization

Lesson: Overfitting and Generalization

Review gallery: asset path, text handling, and print/PDF notes remain available for QA.