

Prompt Life v0.28.3 Journey Exercise Review

Generated: 2026-06-15T15:24:44.843Z

Status: pass

Exercises inventoried: 39

Exercise priority counts: P0: 0, P1: 0, P2: 6, P3: 33

Visual overflow status: pass

Fixed in this pass: Renamed the alignment exercise data item id from policy-filter to runtime-policy-filter while preserving the visible label "policy filter"; this removes a same-exercise id collision risk without changing learner copy.

1. What Is an LLM?

Stage: Before Morning

Interaction: Tap through one prompt tap-step P2


Action: Tap through context, probability, chosen token, and updated context.

Completion: Final step reveals updated context.

Feedback: Step-specific insight text after each tap.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



2. Where LLMs Fit

Stage: Before Morning

Interaction: Tap a branch tap-reveal P3


Action: Tap category branches in the AI family tree.

Completion: At least one branch explanation is revealed.

Feedback: Branch-specific one-sentence explanation.

Recommendation: Keep for human testing.

No P0/P1 issue found.



3. Rationalists vs Empiricists

Stage: Before Morning

Interaction: Sort rules and patterns

sort-buckets

P3

Action: Classify rules, learned patterns, and hybrid ingredients.

Completion: All visible chips are assigned to buckets.

Feedback: Correctness highlighting plus summary insight.

Recommendation: Keep for human testing.

No P0/P1 issue found.



4. Training

Stage: Before Morning

Interaction: Tap through the training loop tap-step P3


Action: Tap each step in the training loop.

Completion: Update weights step is selected or sequence is completed.

Feedback: Durable-update explanation.

Recommendation: Keep for human testing.

No P0/P1 issue found.



5. Pretraining

Stage: Before Morning

Interaction: Pattern learning or recall? toggle-choice

P3

Action: Toggle between broad pattern learning and perfect recall.

Completion: Broad pattern learning state is selected.

Feedback: Boundary feedback against perfect-memory myth.

Recommendation: Keep for human testing.

No P0/P1 issue found.



6. Overfitting and Generalization

Stage: Before Morning

Interaction: Choose the better curve tap-choice P3


Action: Choose the curve that generalizes better.

Completion: Generalizing curve is selected.

Feedback: Validation-example feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



7. Fine-Tuning

Stage: Before Morning

Interaction: Sort the steering moves sort-buckets P3


Action: Sort durable training, temporary context, and decoding moves.

Completion: Each steering move is assigned.

Feedback: Bucket-specific correction.

Recommendation: Keep for human testing.

No P0/P1 issue found.



8. Alignment

Stage: Before Morning

Interaction: Group the alignment methods

sort-buckets

P3


Action: Group alignment methods by durable shaping, runtime steering, and evaluation.

Completion: Each method is assigned.


Feedback: Method-group summary.

Recommendation: Keep for human testing.

No P0/P1 issue found.



320px



390px

9. Inference

Stage: Morning Commute

Interaction: Walk the forward pass tap-step P3

Action: Walk the forward pass from context to next token.

Completion: Forward-pass path is completed.

Feedback: Temporary-versus-durable feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.

TINY INTERACTION

Walk the forward pass

Context

Fixed weights

Temporary activations

Scores

Which part is temporary during ordinary inference?

Activations

Model weights

Training dataset

Pick the part that appears for this run and then disappears.

The current context moves through the model. Temporary activations change, but durable weights do not.

320px

TINY INTERACTION

Walk the forward pass

Context

Fixed weights

Temporary activations

Scores

Which part is temporary during ordinary inference?

Activations

Model weights

Training dataset

Pick the part that appears for this run and then disappears.

The current context moves through the model. Temporary activations change, but durable weights do not.

390px

10. Prompt vs Response

Stage: Morning Commute

Interaction: Label the context pieces label-pieces P3


Action: Tap rows that separate prompt, response so far, next token, and updated context.

Completion: All four rows can be inspected.

Feedback: Row-specific explanation.

Recommendation: Keep for human testing.

No P0/P1 issue found.



11. Tokenization

Stage: Morning Commute

Interaction: Split the sentence tap-reveal P3


Action: Reveal token cards and uneven chunks.

Completion: Uneven examples are shown.

Feedback: Token boundary feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



12. Token IDs

Stage: Morning Commute

Interaction: Match token to ID match-lookup P3

Action: Match token text to token IDs and table rows.

Completion: Lookup relationship is revealed.

Feedback: ID-is-not-meaning feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



13. Embeddings

Stage: Morning Commute

Interaction: Look up a starting vector tap-lookup P3


Action: Tap a token ID to retrieve its embedding row.

Completion: Embedding row is shown.

Feedback: Durable embedding versus temporary hidden-state feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



14. Vectors

Stage: Morning Commute

Interaction: Toggle vector views toggle-view P2

Action: Switch between teaching sliders and distributed feature bars.

Completion: Distributed view is shown.

Feedback: Dimension-label limit feedback.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



15. Tensors

Stage: Morning Commute

Interaction: Inspect the axes tap-axis P2


Action: Inspect token axis, feature axis, and batch sheet.

Completion: All axis notes can be viewed.

Feedback: Axis-specific explanation.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



16. Attention

Stage: Workday

Interaction: Connect the pronoun connect-nodes P3


Action: Connect pronoun token to the most relevant source token.

Completion: A source token is selected.

Feedback: Relevance-not-awareness feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



17. MLP

Stage: Workday

Interaction: Before and after MLP toggle-view P3

Action: Toggle same token before and after MLP.

Completion: After-MLP state is shown.

Feedback: Feature reshaping feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



18. Layers

Stage: Workday

Interaction: Inspect the stack tap-step P3


Action: Tap layers in a transformer stack.

Completion: A layer detail is shown.

Feedback: Repeated numerical transformation feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



19. Hidden States

Stage: Workday

Interaction: Durable or temporary? sort-buckets P3


Action: Sort durable, temporary, and outside-forward-pass items.

Completion: Each item is assigned.

Feedback: Temporary hidden-state feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



20. Logits

Stage: Decision Room

Interaction: Raw scores or probabilities? [toggle-view](#) [P3](#)


Action: Toggle raw scores and probability-ready framing.

Completion: Raw-score state is visible.

Feedback: Logits-are-not-probabilities feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



21. Softmax

Stage: Decision Room

Interaction: Convert the scores tap-transform P3


Action: Convert raw scores into probabilities.

Completion: Softmax conversion is shown.


Feedback: Probability-not-truth feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



320px



390px

22. Sampling

Stage: Decision Room

Interaction: Choose from the probability cloud tap-choice P3


Action: Choose a next token from probability-shaped candidates.

Completion: A candidate token is selected.

Feedback: Sampling uncertainty feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



23. Autoregression

Stage: The Day Repeats

Interaction: Tap-step append loop tap-step P3

Action: Step through choose, append, run again.

Completion: Loop reaches the repeated run state.

Feedback: Append-and-repeat feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



24. Context Window

Stage: The Day Repeats

Interaction: Push the context cards tap-stack P3


Action: Push cards into a bounded context tray.

Completion: A card falls out after the window fills.

Feedback: Temporary-window feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



25. RAG and Retrieval

Stage: The Day Repeats

Interaction: Trace the retrieval lanes tap-step P3

Action: Trace prompt, retriever, retrieved notes, context tray, and response.

Completion: All retrieval lanes can be highlighted.

Feedback: Retrieval-plus-context feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.

TINY INTERACTION

Trace the retrieval lanes

1 Prompt

2 Retriever

3 Notes

4 Context tray

5 Response

Prompt: The user question starts the run.

Tap prompt, retriever, notes, context tray, and response to see retrieval plus context, not training.

320px

TINY INTERACTION

Trace the retrieval lanes

1 Prompt

2 Retriever

3 Notes

4 Context tray

5 Response

Prompt: The user question starts the run.

Tap prompt, retriever, notes, context tray, and response to see retrieval plus context, not training.

390px

26. Grounding

Stage: The Day Repeats

Interaction: Match claims to evidence

connect-nodes

P3

Action: Connect claims to supporting evidence.

Completion: Supported and unsupported claims are visible.

Feedback: Evidence-match feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



27. Hallucinations

Stage: The Day Repeats

Interaction: Mark missing supports tap-multiple P3

Action: Mark fluent claims that need evidence.

Completion: Unsupported claims are marked.

Feedback: Unsupported-output feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



28. How AI Learns

Stage: Twilight: The Wider Landscape

Interaction: Sort the learning modes sort-buckets P3


Action: Sort durable update, temporary steering, retrieval, and evaluation modes.

Completion: Each mode is assigned.


Feedback: What-changed feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



320px



390px

29. Diffusion vs Autoregression

Stage: Twilight: The Wider Landscape

Interaction: Step through two loops tap-compare P3

Action: Step through token append path and denoise path.

Completion: Both paths can be compared.

Feedback: Autoregression-versus-diffusion feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



30. Multimodal AI

Stage: Twilight: The Wider Landscape

Interaction: Match media lanes match-lanes P3

Action: Match media inputs and outputs.

Completion: A media lane pair is selected.

Feedback: Representation-not-human-senses feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



31. The Perfect Storm

Stage: Twilight: The Wider Landscape

Interaction: Tap storm ingredients tap-reveal P3

Action: Tap converging ingredient streams.

Completion: Ingredient explanation is revealed.

Feedback: Convergence-not-spark feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



32. Collective Intelligence, Extracted

Stage: Midnight Ledger

Interaction: Which questions stay human?

sort-buckets

P2

Action: Sort human-rights questions away from model mechanics.

Completion: Questions are classified.

Feedback: Accountability feedback.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



33. Costs We Must Count

Stage: Midnight Ledger

Interaction: Tap the cost ledger tap-reveal P3


Action: Tap cost ledger entries.

Completion: Cost category explanation is revealed.

Feedback: Countable-cost feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



34. Risk vs Myth

Stage: Midnight Ledger

Interaction: Sort risk from myth sort-buckets P3

Action: Sort practical risks from myths.

Completion: Claims are assigned.

Feedback: Mechanism-risk feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



35. Benefits Worth Taking Seriously

Stage: New Dawn

Interaction: Sort benefit strength sort-buckets P2

Action: Sort benefit claims by evidence strength.

Completion: Claims are assigned.

Feedback: Bounded-benefit feedback.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



36. Human-Centered AI

Stage: New Dawn

Interaction: Center the human purpose tap-choice P3

Action: Choose the more human-centered deployment.

Completion: Accountable deployment is selected.

Feedback: Review-and-accountability feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



37. Better AI Is a Choice

Stage: New Dawn

Interaction: Choose the better path tap-multiple P3

Action: Choose design and governance levers.

Completion: Useful levers are selected.

Feedback: Design-choice feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



38. Effective Prompting from Model Literacy

Stage: New Dawn

Interaction: Pack the context tray tap-multiple P3


Action: Pack a context tray with prompt ingredients.

Completion: Prompt tray is completed.

Feedback: Context-design feedback.

Recommendation: Keep for human testing.

No P0/P1 issue found.



39. Model Literate Synthesis

Stage: New Dawn

Interaction: Light the map order-sequence P2

Action: Put the model-literacy chain in order.

Completion: All chain steps are ordered.

Feedback: Mechanics-to-accountability feedback.

Recommendation: Keep for human testing; watch for clarity and replay value.

No P0/P1 issue found.



320px



390px