

Prompt Life v0.28.2 Final Response

Summary

Implemented v0.28.2 Visual Overflow Hardening + Journey Exercise Audit. The app version is now **v0.28.2**. The hardened visual overflow audit now checks actual learner Journey routes at 320px and 390px. A new exercise audit inventories and screenshot-checks all 39 Journey interactions.

Visual overflow status

pass; 0 visual row issues remain.

Journey exercises

39 inventoried; 0 audit issues remain.

Exercise priority counts

P0 0, P1 0, P2 6, P3 33

Readiness

Ready for small human testing, with six P2 interactions to watch.


What Changed

- Fixed learner-route visual label/marker overflow in Overfitting, Softmax, Context Window, Grounding, Hallucinations, and Collective Intelligence.
- Moved former temporary visuals Overfitting vs Generalization, Feature Vector, and Tensor Block to fixed/testable status.
- Added `npm run audit:exercises` and generated JSON, CSV, Markdown, HTML, PDF, and screenshots.
- Repaired mobile interaction layout in Token IDs, Attention, Inference, MLP, Diffusion, Human-Centered AI, and Synthesis.
- Fixed the alignment exercise internal item id collision risk while preserving learner-facing copy.

Verification

- `npm run typecheck`: passed.
- `npm run build`: passed with existing Vite large-chunk warning.
- `npm run build:pages`: passed with existing Vite large-chunk warning.
- `npm run audit:answers`, `audit:checkpoints`, `audit:learner-copy`, `audit:language`, `audit:visual-aids`, `audit:visual-overflow`, and `audit:exercises`: passed.
- `npm run audit:question-clues`: completed with existing warning counts for human review.
- In-app browser smoke QA: Home, Journey, Play, Glossary, and Badge navigated with no horizontal overflow; Badge showed the PNG and v0.28.2.

Screenshots



Overfitting vs Generalization


Memorizing training examples is not the same as learning patterns that transfer to set-aside validation examples.

- 1 Training examples** Old dots are examples the model fit during training.
- 2 Validation examples** Set-aside examples are kept out of training to test transfer.
- 3 Overfit curve** The overfit curve traces old dots too tightly.
- 4 Generalizing curve** The smoother curve reaches new examples better.

Key takeaway:

A model should learn transferable patterns, not just memorize the training dots.

Overfitting visual, 320px



Score to Probability

Normalize before choosing


Softmax turns raw scores into probabilities that sum to one.

- Logits go in** Raw scores for floor, mat, kitchen, and banana enter softmax.
- Softmax converts** The function turns score differences into a probability distribution; it does not pick the next token.
- Probabilities come out** The rounded teaching values sum to 100%.
- Sampling comes next** The probability distribution is ready for a separate choice step.

Key takeaway:

Softmax normalizes raw scores; sampling chooses later.

Softmax visual, 320px



Temporary Context Window

Limited visible input

Only the cards still inside the context tray can influence the next token.

- 1 **Visible now** System instructions, the user prompt, retrieved notes, and response-so-far may all be current context.
- 2 **Limited tray** When the tray fills, older material can fall out or need summarizing by the surrounding system.
- 3 **Temporary** Context can feel like memory, but it does not normally update weights or create permanent memory.
- 4 **Next token** The model can use only what remains visible in the current run.

Key takeaway:

Context is temporary visible input, not durable memory.

Context Window visual, 320px



Claim Support Map

Tying answers to evidence


Grounding asks whether generated claims are actually connected to available evidence.

- 1 **Claim** A generated answer can contain several claims.
- 2 **Evidence** Retrieved passages, data, citations, or tool results can support some claims.
- 3 **Support check** A citation-looking answer is not grounded unless the evidence actually matches the claim.
- 4 **Limit** Grounding helps, but the evidence and its use still need review.

Key takeaway:

Grounding ties claims to evidence; it does not make every claim true.

Grounding visual, 320px



Unsupported Bridge

Fluent is not always grounded

A response can look smooth while one claim has support, another is uncertain, and another lacks evidence.

- 1 **Fluent surface** The output may read smoothly and confidently.
- 2 **Supported claim** One claim can be tied to evidence.
- 3 **Missing support** Another claim may lack evidence, citation, or retrieved context.
- 4 **Review** Grounding, uncertainty, and human review reduce risk but do not erase it.

Key takeaway:

Fluency is not evidence.

Hallucination visual, 320px



Borrowed Flames

No creators, no model

Human-created traces can become training data and model patterns, while rights and responsibility remain human questions.

- **Human expression** Books, sites, code, art, journalism, forums, documentation, and research leave learnable traces.
- **Collection questions** Provenance, consent, copyright, attribution, and compensation matter.
- **Model limit** The model absorbs statistical patterns; it does not understand gratitude or responsibility.

Key takeaway:

The model did not create its abilities alone.

Collective Intelligence visual, 320px

TINY INTERACTION

Match token to ID

dog	cat	
floor		
dog		
ID 421	ID 982	ID 1576
embedding row ?		
Match dog to its lookup number.		

Match dog, cat, and floor to their lookup numbers and table rows.

Token IDs interaction, 390px

TINY INTERACTION

Connect the pronoun

Which earlier token helps interpret "it"?

dog	cat
because	it
hissed	
dog	cat

Tap the token that "it" most likely depends on.

Choose which earlier token helps interpret "it." Attention is relevance, not awareness.

Attention interaction, 320px

TINY INTERACTION

Step through two loops

Token path

Denoise path

score next
token

choose token

append
token

run again

Show choose token

Current step: score next
token.

Tap Token path or Denoise path to
compare append-and-repeat with
denoise-and-refine.

Diffusion interaction, 320px

Skip to content

7 waiting

8 waiting

prompt enters context

text becomes tokens
and IDs

embeddings become
hidden states

logits become
probabilities

one token is sampled

token is appended and
the model runs again

RAG/grounding can add
evidence

humans review and
remain accountable

Reset chain

Begin with what enters
the model run.

Put the model story in order, from
prompt context to human
accountability.

Synthesis interaction, 320px

BADGE

Prompt Life: Model Literate

Under construction while the evidence model is reviewed.



Under construction

Pending human review

Prompt Life: Model Literate is under construction. The Journey, checkpoint bank, and evidence model are being tested before any badge is issued.



Home



Journey



Play



Glossary



Badge

Badge page smoke, 390px