

# Prompt Life v0.27.17 Visual Aid Quality Pass 1

Final response generated 2026-06-11. This pass repaired priority model-mechanism visuals, prepared Image 2 prompts for social/ethical candidates, kept exact diagrams coded, and preserved Journey, Play, checkpoint, progress, and badge behavior.

**39**

visual aids inspected

**8**

priority visuals changed

**4**

Image 2 candidates

**0**

marker mismatches

## Summary

- Changed Fine-Tuning callouts and the coded Attention, MLP, Layers, Hidden States, Logits, Softmax, and Sampling visuals.
- Left Collective Intelligence, Benefits, Costs, and Human-Centered AI stable in-app and documented them as future Image 2 candidates.
- Created v0.27.17 visual-aid quality report, prompt sheet, backlog JSON/CSV, consistency audit files, screenshots, and this final PDF.
- Version is now `0.27.17`.

## Priority Visuals Changed

Fine-Tuning Path; Attention Weave; MLP Feature Workshop; Transformer Stack; Hidden State Flow; Raw Scoreboard; Score to Probability; Weighted Token Choice.

## Priority Left Unchanged

Collective Intelligence, Benefits, Costs, and Human-Centered AI were left unchanged because they are better handled as future textless Image 2 assets with HTML callouts, not rushed replacements.

## Verification

- `npm run typecheck`: passed
- `npm run build`: passed; existing Vite large chunk warning remains
- `npm run build:pages`: passed; existing Vite large chunk warning remains
- `npm run audit:answers`: passed
- `npm run audit:checkpoints`: passed

- npm run audit:question-clues: completed with 36 term-echo candidates, 43 answer-length warnings, 0 vague-stem warnings
- npm run audit:learner-copy: passed
- npm run audit:visual-aids: passed, 39 aids, 0 mismatches

## Mobile/Desktop QA

- 390px and 320px visual review screenshots captured.
- Home, Journey, Play, Glossary, and Badge loaded at 390px with no horizontal overflow.
- All eight Journey stage links focused the expected section.
- Desktop visual aid sample captured at 1024px.

## Files Changed

- package.json
- package-lock.json
- src/main.tsx
- src/components/VisualAids.tsx
- scripts/audit-visual-aids.mjs
- scripts/generate-visual-aid-quality-v02717.mjs
- docs/REVIEW\_NOTES.md
- docs/journey/visual-aids/image-2-visual-prompts-v0-27-17.md
- docs/journey/visual-aids/visual-aid-quality-backlog-v0-27-17.json
- docs/journey/visual-aids/visual-aid-quality-backlog-v0-27-17.csv
- docs/journey/visual-aids/visual-aid-consistency-audit-v0-27-17.json
- docs/journey/visual-aids/visual-aid-consistency-audit-v0-27-17.csv
- docs/journey/visual-aids/visual-aid-consistency-audit-v0-27-17.md
- docs/journey/visual-aids/prompt-life-v0-27-17-visual-aid-quality-pass-1-report.html
- docs/journey/visual-aids/prompt-life-v0-27-17-visual-aid-quality-pass-1-report.pdf
- docs/journey/screenshots/v0-27-17/\*.png

# Screenshots

COLLECTIVE-INTELLIGENCE-LANTERN

## Borrowed Flames

Lesson: Collective Intelligence, Extracted

Pattern: collectiveLantern

Variant: zen-garden-map

Callouts: bullets

Markers: 0

```
graph LR; Writing --- TD[Training data]; Code --- TD; Art --- TD; Research --- TD; Forums --- TD; Docs --- TD; rights_questions[rights questions] --- TD;
```

### Borrowed Flames

**No creators, no model**

Human-created traces can become training data and model patterns, while rights and responsibility remain human questions.

- Human expression** Books, sites, code, art, journalism, forums, documentation, and research leave learnable traces.
- Collection questions** Provenance, consent, copyright, attribution, and compensation matter.
- Model limit** The model absorbs statistical patterns; it does not understand gratitude or responsibility.

Key takeaway:

The model did not create its abilities alone.

LEARNING OBJECTIVE

Collective Intelligence 390px

BENEFITS-TOOL-GARDEN

## Benefit Tiers

Lesson: Benefits Worth Taking Seriously

Pattern: benefitsGarden

Variant: zen-garden-map

Callouts: bullets

Markers: 0

| Benefit confidence tiers                                                     |                                                                 |                                                                          |
|------------------------------------------------------------------------------|-----------------------------------------------------------------|--------------------------------------------------------------------------|
| <b>Useful now + review</b><br>Access<br>Translate<br>Summarize<br>Search/RAG | <b>Plausible + safeguards</b><br>Tutoring<br>Triage<br>Workflow | <b>Speculative / hype</b><br>Solves all<br>Replaces judgment<br>No costs |
| Real help needs evidence + review                                            |                                                                 |                                                                          |

### Benefit Tiers

**Benefits without utopia**

AI benefits are easiest to trust when claims are tiered by evidence, context, safeguards, and human review.

- Useful now with review** Accessibility, translation support, summarization, search/RAG, drafting, and coding assistance can help under human review.
- Plausible with safeguards** Tutoring support, research triage, brainstorming, and workflow support should be evaluated in context.
- Speculative / hype** Broad utopia, replacement, or no-cost claims should not be stated as fact.

Key takeaway:

Benefits can be real and still bounded.

LEARNING OBJECTIVE

Benefits 390px

## Cost Ledger


Lesson: Costs We Must Count

Pattern: costsFactory

Variant: zen-garden-map

Callouts: bullets

Markers: 0



### Cost Ledger

**The answer is not weightless**

AI systems depend on infrastructure, human systems, and governance choices that need honest accounting.

- **Infrastructure** Energy, water, carbon, data centers, chips, and e-waste vary by system and workload.
- **Human systems** Labor disruption, deskilling, privacy, bias, and information pollution depend on deployment choices.
- **Power** Concentration of data, compute, and capital affects who benefits and who decides.

#### Key takeaway:

Costs vary, but they are real enough to count.

#### LEARNING OBJECTIVE

Make AI costs visible as a ledger without inventing statistics or fear-heavy imagery.

## Accountability Flow

Lesson: Human-Centered AI

Pattern: humanGarden

Variant: zen-garden-map

Callouts: bullets

Markers: 0



### Accountability Flow

**Tools should serve dignity**

A human-centered deployment treats AI output as support inside a decision context, then keeps review, governance, and accountable action with people.

- **Human judgment** People remain accountable for high-stakes decisions and institutional use.
- **Dignity** Speed, profit, and automation should not outrank persons, learning, or relationships.
- **Model limit** A model can sound ethical without moral understanding.

#### Key takeaway:

Powerful tools still need human purpose.

#### LEARNING OBJECTIVE

Make human-centered AI concrete through a support-note accountability scenario.

Costs 390px

Human-Centered AI 390px



## Fine-Tuning Path

Lesson: Fine-Tuning Pattern: generatedImage

Variant: generated-image Callouts: bullets

Markers: 0

Asset: before-morning-finetuning-path.png

Type: generated-image



### Fine-Tuning Path

#### Durable adaptation, not current-run steering

Fine-tuning can durably adapt model behavior through additional training or adapters; prompts, RAG, and sampling steer only the current run.


- **Durable adaptation** Additional training or adapters can change how future runs behave.
- **Prompt steering** A prompt can steer this run without changing weights.
- **RAG and context** Retrieved text can enter the current context without training the model.

## Attention Weave

Lesson: Attention Pattern: attention

Variant: paper-diagram Callouts: bullets

Markers: 0



### Attention Weave

#### Relevance, not awareness

Attention assigns temporary relevance weights from one token position to other positions in the current context.

- **Target position** The token it is the position currently gathering context.
- **Source positions** Other tokens, such as cat and dog, can contribute different amounts of relevance.
- **Weighted relevance** Here, cat receives the strongest link because it fits the local sentence relationship.
- **Limit** Attention is a calculation over positions. It is not human noticing or awareness.

#### Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

Fine-Tuning 390px

Attention 390px

## MLP

## MLP Feature Workshop


Lesson: MLP

Pattern: mlp

Variant: origami-object

Callouts: bullets

Markers: 0



## MLP Feature Workshop

## Same token, reshaped features

After attention mixes across token positions, the MLP reshapes features within each token position.

- **Attention mixes** Attention lets token positions pull in context from other positions.
- **Same position** The example token cat stays in its own position as it enters the MLP.
- **MLP reshapes** The feature bars change inside that token position.
- **Temporary values** The MLP uses durable weights, but the feature values in this run are temporary activations.

## Key takeaway:

Attention mixes across positions; the MLP reshapes features within one position.

## LEARNING OBJECTIVE

Show the attention-versus-MLP distinction:

## LAYERS

## Transformer Stack

Lesson: Layers

Pattern: layers

Variant: origami-object

Callouts: bullets

Markers: 0



## Transformer Stack

## Repeated attention plus MLP

Layers repeat attention and MLP transformations to refine temporary hidden states.

- **Input states** Current-context hidden states enter the stack.
- **Repeated block** Each layer applies attention and an MLP to the temporary states.
- **Refined states** The states keep changing as they move toward next-token scores.
- **Limit** Layers are repeated numerical transformations, not conscious reasoning steps.

## Key takeaway:

Layers refine hidden states during the run; they do not store permanent memory.

## LEARNING OBJECTIVE

Show repeated transformer blocks without

MLP 390px

Layers 390px

## HIDDEN-STATES

# Hidden State Flow


Lesson: Hidden States

Pattern: hidden

Variant: origami-object

Callouts: bullets

Markers: 0



## Hidden State Flow

**Temporary during this forward pass**

A token starts with an embedding and becomes context-shaped hidden states as it moves through layers.

- **Embedding** A token ID retrieves a learned starting vector.
- **Hidden states** Layers reshape that vector into temporary context-shaped states.
- **Scores later** Final hidden states help produce raw next-token scores.
- **Not memory** A hidden state is not visible text, permanent memory, or a stored fact.

### Key takeaway:

Hidden states are temporary internal vectors shaped by current context.

### LEARNING OBJECTIVE

Contrast embedding, hidden state, weight, memory, and visible text without making hidden

## LOGITS

# Raw Scoreboard

Lesson: Logits

Pattern: logits

Variant: paper-diagram

Callouts: bullets

Markers: 0



## Raw Scoreboard

**Before probabilities**

The final hidden state feeds candidate next-token scores. These logits rank candidates, but they are raw scores, not probabilities or truth.

- **Hidden state arrives** The final context-shaped vector reaches the output scoring step.
- **Candidates get scores** Possible next tokens such as floor, room, counter, and quantum receive raw logits.
- **Scores rank candidates** Higher score means stronger local candidate, not guaranteed truth.
- **Softmax comes next** Softmax converts raw scores into probabilities.

### Key takeaway:

Logits are raw next-token scores, not probabilities and not truth.

Hidden States 390px

Logits 390px



## SOFTMAX

### Score to Probability


Lesson: Softmax

Pattern: softmax

Variant: neon-flow

Callouts: bullets

Markers: 0



#### Score to Probability

##### Normalize before choosing

Softmax converts raw logits into probabilities that sum to 100%. It prepares the distribution; it does not choose the token.

- **Logits go in** Raw scores such as 8.2, 5.1, 2.0, and -1.0 enter softmax.
- **Softmax converts** The function turns score differences into a probability distribution; it does not pick the next token.
- **Probabilities come out** The rounded teaching values sum to 100%.
- **Not truth** Probability is about model likelihood, not evidence or correctness.

##### Key takeaway:

Softmax turns raw scores into probabilities; sampling chooses later.

## LEARNING OBJECTIVE

Softmax 390px

## SAMPLING

### Weighted Token Choice


Lesson: Sampling

Pattern: sampling

Variant: neon-flow

Callouts: bullets

Markers: 0



#### Weighted Token Choice

##### One token chosen

Sampling chooses one next response token from probabilities. Probability is not truth: a low-probability token can still be possible, and a likely token can still be unsupported.

- **Probability candidates** floor is strongest, room is plausible, counter is weak, and quantum is barely likely in this context.
- **One token chosen** The decoding step chooses one next response token.
- **Append next** The chosen token is appended before the model runs again.
- **Fit is not proof** A likely token can fit the model pattern and still be unsupported in a factual claim.

##### Key takeaway:

Sampling chooses one token from probabilities; fit is not proof.

Sampling 390px



## ATTENTION

# Attention Weave


Lesson: Attention

Pattern: attention

Variant: paper-diagram

Callouts: bullets

Markers: 0



## Attention Weave

### Relevance, not awareness

Attention assigns temporary relevance weights from one token position to other positions in the current context.

- **Target position** The token it is the position currently gathering context.
- **Source positions** Other tokens, such as cat and dog, can contribute different amounts of relevance.
- **Weighted relevance** Here, cat receives the strongest link because it fits the local sentence relationship.
- **Limit** Attention is a calculation over positions. It is not human noticing or awareness.


### Key takeaway:

Attention is weighted relevance between token positions, not

## ATTENTION

### Attention Weave

Lesson: Attention Pattern: attention Variant: paper-diagram Callouts: bullets Markers: 0



### Attention Weave

#### Relevance, not awareness

Attention assigns temporary relevance weights from one token position to other positions in the current context.

- **Target position** The token it is the position currently gathering context.
- **Source positions** Other tokens, such as cat and dog, can contribute different amounts of relevance.

## Desktop visual aid sample