Prompt Life v0.27.17 Visual Aid Quality Pass 1

Generated 2026-06-12T03:21:11.544Z. This pass inspected all 39 visual-aid catalog entries, reviewed the 12 priority visuals deeply, repaired the coded mechanism visuals that needed clearer model boundaries, and prepared Image 2 prompts for the four social/ethical candidates. No generated PNG assets were added.

39
catalog aids inspected

8
priority visuals changed

4
priority visuals left stable

4
Image 2 candidates

Executive Summary

The pass keeps exact mechanisms coded and moves atmospheric social concepts into a future Image 2 queue. Fine-Tuning received stronger durable-versus-temporary callouts. Attention, MLP, Layers, Hidden States, Logits, Softmax, and Sampling were refined as coded SVGs/HTML so learners see source-target relevance, per-token feature reshaping, repeated transformer blocks, temporary hidden-state flow, raw-score boundaries, softmax normalization, and one-token sampling more clearly.

Priority Visual Table

Learning card	Current type	Concept	Issue	Action taken	Remaining concern	Next recommendation
Collective Intelligence, Extracted	coded SVG/HTML	Human-created traces can enter training data; rights questions remain human and institutional.	Social/provenance idea can feel too chart-like if overdiagrammed.	Left learner UI unchanged and created a production Image 2 prompt with HTML callouts.	Needs a future textless PNG before replacing the coded scene.	Generate Image 2 asset in a dedicated asset pass.

Learning card	Current type	Concept	Issue	Action taken	Remaining concern	Next recommendation
Benefits Worth Taking Seriously	coded SVG/HTML	Benefits can be useful while still bounded by evidence, context, safeguards, and review.	Tier concept is more judgment landscape than exact mechanism.	Left learner UI unchanged and created a production Image 2 prompt with outside callouts.	Future PNG must avoid hype, miracle imagery, and baked-in labels.	Generate Image 2 asset after visual tone review.
Costs We Must Count	coded SVG/HTML	AI outputs depend on infrastructure, human systems, power, and governance choices.	Invisible costs benefit from atmosphere, but numbers and claims must stay out of the image.	Left learner UI unchanged and created a production Image 2 prompt.	Generated asset must avoid fake statistics and disaster imagery.	Generate Image 2 asset with strict negative prompt.
Human-Centered AI	coded SVG/HTML	AI can support decisions, but review, governance, dignity, and accountability remain human.	Accountability concept needs warmth without robot judges or moral-halo metaphors.	Left learner UI unchanged and created a production Image 2 prompt.	Future image must stay textless and avoid faces, halos, and automation fantasy.	Generate Image 2 asset after the social concept art set is approved.
Fine-Tuning	generated PNG plus HTML callouts	Fine-tuning can durably adapt behavior; prompts, RAG/context, and sampling steer the current run.	The previous caption emphasized durable adaptation but did not make the neighboring boundaries explicit enough.	Updated caption, subtitle, callouts, key takeaway, accessible description, and print note.	Future coded overlay markers may help if the generated art is reused in a high-stakes review packet.	Keep current PNG and improved HTML callouts for v0.27.17.
Attention	coded SVG/HTML	Attention computes weighted relevance between token positions, not awareness.	A web-like scene can accidentally imply broad awareness or mind-like noticing.	Redrew the diagram around a target token, source tokens, one strong link, one weak link, and a not-awareness note.	No major concern after mobile QA; keep labels short.	Keep coded.
MLP	coded SVG/HTML	Attention mixes across positions; the MLP reshapes features within one token position.	The visual needed a cleaner contrast between cross-position mixing and same-position reshaping.	Redrew with a compact attention panel, same-position cat token, MLP panel, and before/after feature bars.	Feature bars are still simplified teaching symbols.	Keep coded.
Layers	coded SVG/HTML	Repeated transformer blocks refine temporary hidden states.	The stack needed stronger "repeated block" structure without implying thought steps.	Redrew each layer as Layer N plus attention and MLP chips, with refined output states and temporary-this-run note.	Residual paths and normalization remain in callouts, not the diagram.	Keep coded.

Learning card	Current type	Concept	Issue	Action taken	Remaining concern	Next recommendation
Hidden States	coded SVG/HTML	Hidden states are temporary context-shaped vectors, distinct from embeddings and durable weights.	Needed a simpler contrast among ID, embedding, hidden states, scores, weights, and memory.	Redrew the flow and added durable weights versus temporary states language inside the diagram.	Still simplified; details stay in the HTML callouts.	Keep coded.
Logits	coded SVG/HTML	Logits are raw next-token scores before probabilities, not confidence or truth.	Needed stronger "raw score" and "not probability/truth" framing.	Adjusted top labels and bottom boundary chip while preserving the candidate-score bars.	Scores are teaching values, not actual model internals.	Keep coded.
Softmax	coded SVG/HTML	Softmax turns raw scores into probabilities that sum to one; it does not choose the token.	Previous language risked letting learners merge softmax and sampling.	Updated caption/callouts and diagram note to say softmax is not the chooser.	Rounded probabilities are teaching values.	Keep coded.
Sampling	coded SVG/HTML	Sampling chooses one next token from probability-shaped options; fit is not proof.	Needed the truth boundary in callouts, not squeezed into the diagram.	Updated callout/key takeaway and simplified the chosen-token label area.	Keep temperature/top-k/top-p out of this intro visual.	Keep coded.

Image 2 Prompt Section

The production prompt sheet is `image-2-visual-prompts-v0-27-17.md`. Recommended Image 2 candidates:

Collective Intelligence, Extracted **Benefits Worth Taking Seriously** **Costs We Must Count** **Human-Centered AI**

All four prompts require textless ZenTron Origami imagery and keep essential explanations in HTML callouts outside the image.

Cards Kept Coded

These are explicitly kept as coded SVG/HTML because their labels, positions, and arrows teach exact model mechanics or category relationships:

Attention | MLP | Layers | Hidden States | Logits | Softmax | Sampling | Tokenization | Token IDs | Embeddings | Vectors | Tensors | RAG and Retrieval | Grounding | Context Window

Before / After Notes

Fine-Tuning

Before: The previous caption emphasized durable adaptation but did not make the neighboring boundaries explicit enough.

After: Updated caption, subtitle, callouts, key takeaway, accessible description, and print note.

Remaining concern: Future coded overlay markers may help if the generated art is reused in a high-stakes review packet.

Attention

Before: A web-like scene can accidentally imply broad awareness or mind-like noticing.

After: Redrew the diagram around a target token, source tokens, one strong link, one weak link, and a not-awareness note.

Remaining concern: No major concern after mobile QA; keep labels short.

MLP

Before: The visual needed a cleaner contrast between cross-position mixing and same-position reshaping.

After: Redrew with a compact attention panel, same-position cat token, MLP panel, and before/after feature bars.

Remaining concern: Feature bars are still simplified teaching symbols.

Layers

Before: The stack needed stronger "repeated block" structure without implying thought steps.

After: Redrew each layer as Layer N plus attention and MLP chips, with refined output states and temporary-this-run note.

Remaining concern: Residual paths and normalization remain in callouts, not the diagram.

Hidden States

Before: Needed a simpler contrast among ID, embedding, hidden states, scores, weights, and memory.

After: Redrew the flow and added durable weights versus temporary states language inside the diagram.

Remaining concern: Still simplified; details stay in the HTML callouts.

Logits

Before: Needed stronger "raw score" and "not probability/truth" framing.

After: Adjusted top labels and bottom boundary chip while preserving the candidate-score bars.

Remaining concern: Scores are teaching values, not actual model internals.

Softmax

Before: Previous language risked letting learners merge softmax and sampling.

After: Updated caption/callouts and diagram note to say softmax is not the chooser.

Remaining concern: Rounded probabilities are teaching values.

Sampling

Before: Needed the truth boundary in callouts, not squeezed into the diagram.

After: Updated callout/key takeaway and simplified the chosen-token label area.

Remaining concern: Keep temperature/top-k/top-p out of this intro visual.

Screenshots

Screenshots are captured from the review route after the repairs. They are intentionally mobile-heavy because Prompt Life is mobile-first.

COLLECTIVE-INTELLIGENCE-LANTERN

Borrowed Flames

Lesson: Collective Intelligence, Extracted

Pattern: collectiveLantern

Variant: zen-garden-map

Callouts: bullets

Markers: 0

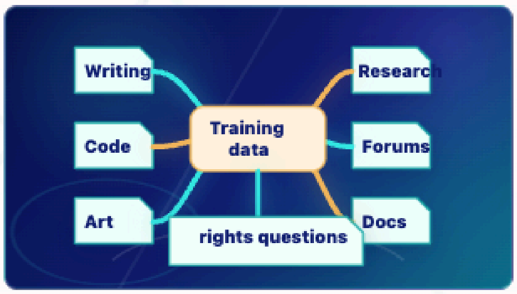

Borrowed Flames

No creators, no model

Human-created traces can become training data and model patterns, while rights and responsibility remain human questions.

- **Human expression** Books, sites, code, art, journalism, forums, documentation, and research leave learnable traces.
- **Collection questions** Provenance, consent, copyright, attribution, and compensation matter.
- **Model limit** The model absorbs statistical patterns; it does not understand gratitude or responsibility.

> **Key takeaway:**
> The model did not create its abilities alone.

LEARNING OBJECTIVE

Collective Intelligence 390px

BENEFITS-TOOL-GARDEN

Benefit Tiers

Lesson: Benefits Worth Taking Seriously

Pattern: benefitsGarden

Variant: zen-garden-map

Callouts: bullets

Markers: 0

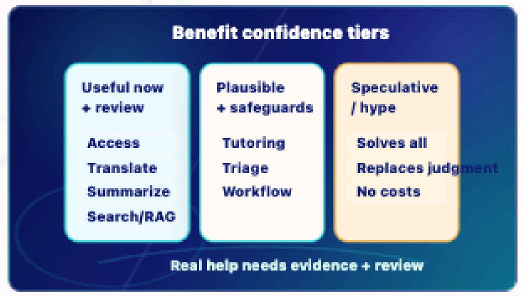

Benefit Tiers

Benefits without utopia

AI benefits are easiest to trust when claims are tiered by evidence, context, safeguards, and human review.

- **Useful now with review** Accessibility, translation support, summarization, search/RAG, drafting, and coding assistance can help under human review.
- **Plausible with safeguards** Tutoring support, research triage, brainstorming, and workflow support should be evaluated in context.
- **Speculative / hype** Broad utopia, replacement, or no-cost claims should not be stated as fact.

> **Key takeaway:**
> Benefits can be real and still bounded.

LEARNING OBJECTIVE

Benefits 390px

COSTS-INVISIBLE-FACTORY

Cost Ledger

Lesson: Costs We Must Count

Pattern: costsFactory

Variant: zen-garden-map

Callouts: bullets

Markers: 0

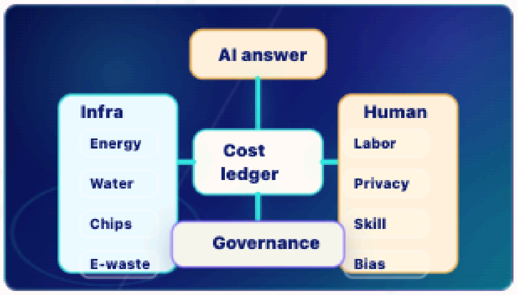

Cost Ledger

The answer is not weightless

AI systems depend on infrastructure, human systems, and governance choices that need honest accounting.

- **Infrastructure** Energy, water, carbon, data centers, chips, and e-waste vary by system and workload.
- **Human systems** Labor disruption, deskilling, privacy, bias, and information pollution depend on deployment choices.
- **Power** Concentration of data, compute, and capital affects who benefits and who decides.

Key takeaway:
Costs vary, but they are real enough to count.

LEARNING OBJECTIVE

Make AI costs visible as a ledger without inventing statistics or fear-heavy imagery.

Costs 390px

HUMAN-CENTERED-AI-GARDEN

Accountability Flow

Lesson: Human-Centered AI

Pattern: humanGarden

Variant: zen-garden-map

Callouts: bullets

Markers: 0

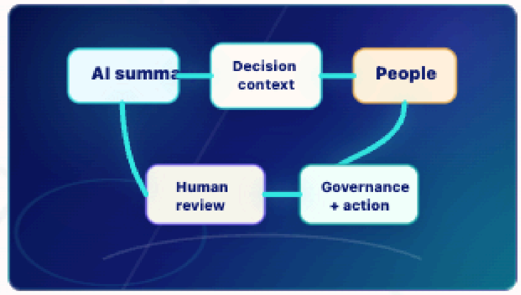

Accountability Flow

Tools should serve dignity

A human-centered deployment treats AI output as support inside a decision context, then keeps review, governance, and accountable action with people.

- **Human judgment** People remain accountable for high-stakes decisions and institutional use.
- **Dignity** Speed, profit, and automation should not outrank persons, learning, or relationships.
- **Model limit** A model can sound ethical without moral understanding.

Key takeaway:
Powerful tools still need human purpose.

LEARNING OBJECTIVE

Make human-centered AI concrete through a support-note accountability scenario.

Human-Centered AI 390px

BEFORE-MORNING-FINETUNING-PATH

Fine-Tuning Path

Lesson: Fine-Tuning | Pattern: generatedImage | Variant: generated-image | Callouts: bullets | Markers: 0 | Asset: before-morning-finetuning-path.png | Type: generated-image

Fine-Tuning Path

Durable adaptation, not current-run steering

Fine-tuning can durably adapt model behavior through additional training or adapters; prompts, RAG, and sampling steer only the current run.

- **Durable adaptation** Additional training or adapters can change how future runs behave.
- **Prompt steering** A prompt can steer this run without changing weights.
- **RAG and context** Retrieved text can enter the current context without training the model.

Fine-Tuning 390px

ATTENTION

Attention Weave

Lesson: Attention | Pattern: attention | Variant: paper-diagram | Callouts: bullets | Markers: 0

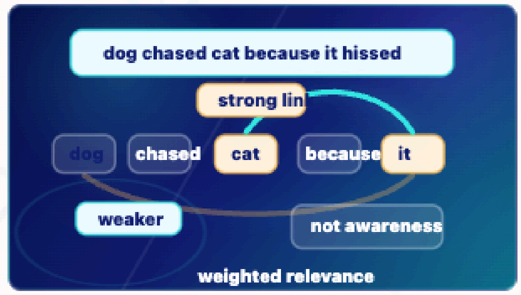

Attention Weave

Relevance, not awareness

Attention assigns temporary relevance weights from one token position to other positions in the current context.

- **Target position** The token it is the position currently gathering context.
- **Source positions** Other tokens, such as cat and dog, can contribute different amounts of relevance.
- **Weighted relevance** Here, cat receives the strongest link because it fits the local sentence relationship.
- **Limit** Attention is a calculation over positions. It is not human noticing or awareness.

> **Key takeaway:**
> Attention is weighted relevance between token positions, not consciousness.

Attention 390px

MLP

MLP Feature Workshop

Lesson: MLP | Pattern: mlp

Variant: origami-object | Callouts: bullets

Markers: 0

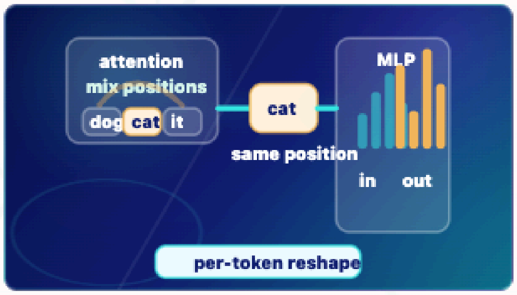

MLP Feature Workshop

Same token, reshaped features

After attention mixes across token positions, the MLP reshapes features within each token position.

- **Attention mixes** Attention lets token positions pull in context from other positions.
- **Same position** The example token cat stays in its own position as it enters the MLP.
- **MLP reshapes** The feature bars change inside that token position.
- **Temporary values** The MLP uses durable weights, but the feature values in this run are temporary activations.

Key takeaway:
Attention mixes across positions; the MLP reshapes features within one position.

LEARNING OBJECTIVE

Show the attention-versus-MLP distinction:

MLP 390px

LAYERS

Transformer Stack

Lesson: Layers | Pattern: layers

Variant: origami-object | Callouts: bullets

Markers: 0

Transformer Stack

Repeated attention plus MLP

Layers repeat attention and MLP transformations to refine temporary hidden states.

- **Input states** Current-context hidden states enter the stack.
- **Repeated block** Each layer applies attention and an MLP to the temporary states.
- **Refined states** The states keep changing as they move toward next-token scores.
- **Limit** Layers are repeated numerical transformations, not conscious reasoning steps.

Key takeaway:
Layers refine hidden states during the run; they do not store permanent memory.

LEARNING OBJECTIVE

Show repeated transformer blocks without

Layers 390px

HIDDEN-STATES

Hidden State Flow

Lesson: Hidden States

Pattern: hidden

Variant: origami-object

Callouts: bullets

Markers: 0

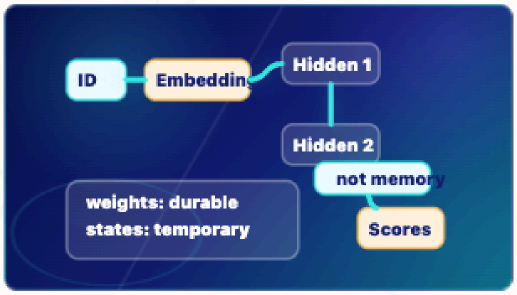

Hidden State Flow

Temporary during this forward pass

A token starts with an embedding and becomes context-shaped hidden states as it moves through layers.

- **Embedding** A token ID retrieves a learned starting vector.
- **Hidden states** Layers reshape that vector into temporary context-shaped states.
- **Scores later** Final hidden states help produce raw next-token scores.
- **Not memory** A hidden state is not visible text, permanent memory, or a stored fact.

> **Key takeaway:**
> Hidden states are temporary internal vectors shaped by current context.

LEARNING OBJECTIVE

Contrast embedding, hidden state, weight, memory, and visible text without making hidden

Hidden States 390px

LOGITS

Raw Scoreboard

Lesson: Logits

Pattern: logits

Variant: paper-diagram

Callouts: bullets

Markers: 0

Raw Scoreboard

Before probabilities

The final hidden state feeds candidate next-token scores. These logits rank candidates, but they are raw scores, not probabilities or truth.

- **Hidden state arrives** The final context-shaped vector reaches the output scoring step.
- **Candidates get scores** Possible next tokens such as floor, room, counter, and quantum receive raw logits.
- **Scores rank candidates** Higher score means stronger local candidate, not guaranteed truth.
- **Softmax comes next** Softmax converts raw scores into probabilities.

> **Key takeaway:**
> Logits are raw next-token scores, not probabilities and not truth.

Logits 390px

SOFTMAX

Score to Probability

Lesson: Softmax | Pattern: softmax | Variant: neon-flow | Callouts: bullets | Markers: 0

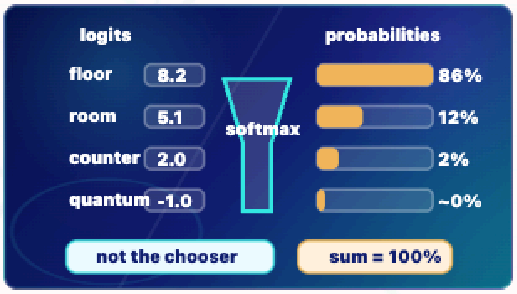

Score to Probability

Normalize before choosing

Softmax converts raw logits into probabilities that sum to 100%. It prepares the distribution; it does not choose the token.

- **Logits go in** Raw scores such as 8.2, 5.1, 2.0, and -1.0 enter softmax.
- **Softmax converts** The function turns score differences into a probability distribution; it does not pick the next token.
- **Probabilities come out** The rounded teaching values sum to 100%.
- **Not truth** Probability is about model likelihood, not evidence or correctness.

> **Key takeaway:**
> Softmax turns raw scores into probabilities; sampling chooses later.

LEARNING OBJECTIVE

Softmax 390px

SAMPLING

Weighted Token Choice

Lesson: Sampling | Pattern: sampling | Variant: neon-flow | Callouts: bullets | Markers: 0

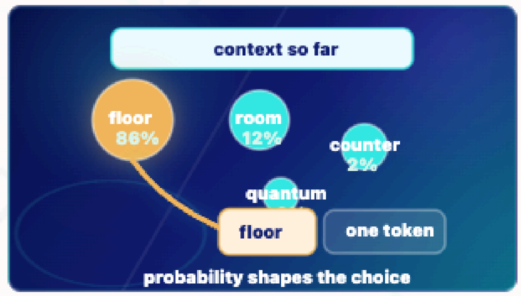

Weighted Token Choice

One token chosen

Sampling chooses one next response token from probabilities. Probability is not truth: a low-probability token can still be possible, and a likely token can still be unsupported.

- **Probability candidates** floor is strongest, room is plausible, counter is weak, and quantum is barely likely in this context.
- **One token chosen** The decoding step chooses one next response token.
- **Append next** The chosen token is appended before the model runs again.
- **Fit is not proof** A likely token can fit the model pattern and still be unsupported in a factual claim.

> **Key takeaway:**
> Sampling chooses one token from probabilities; fit is not proof.

Sampling 390px

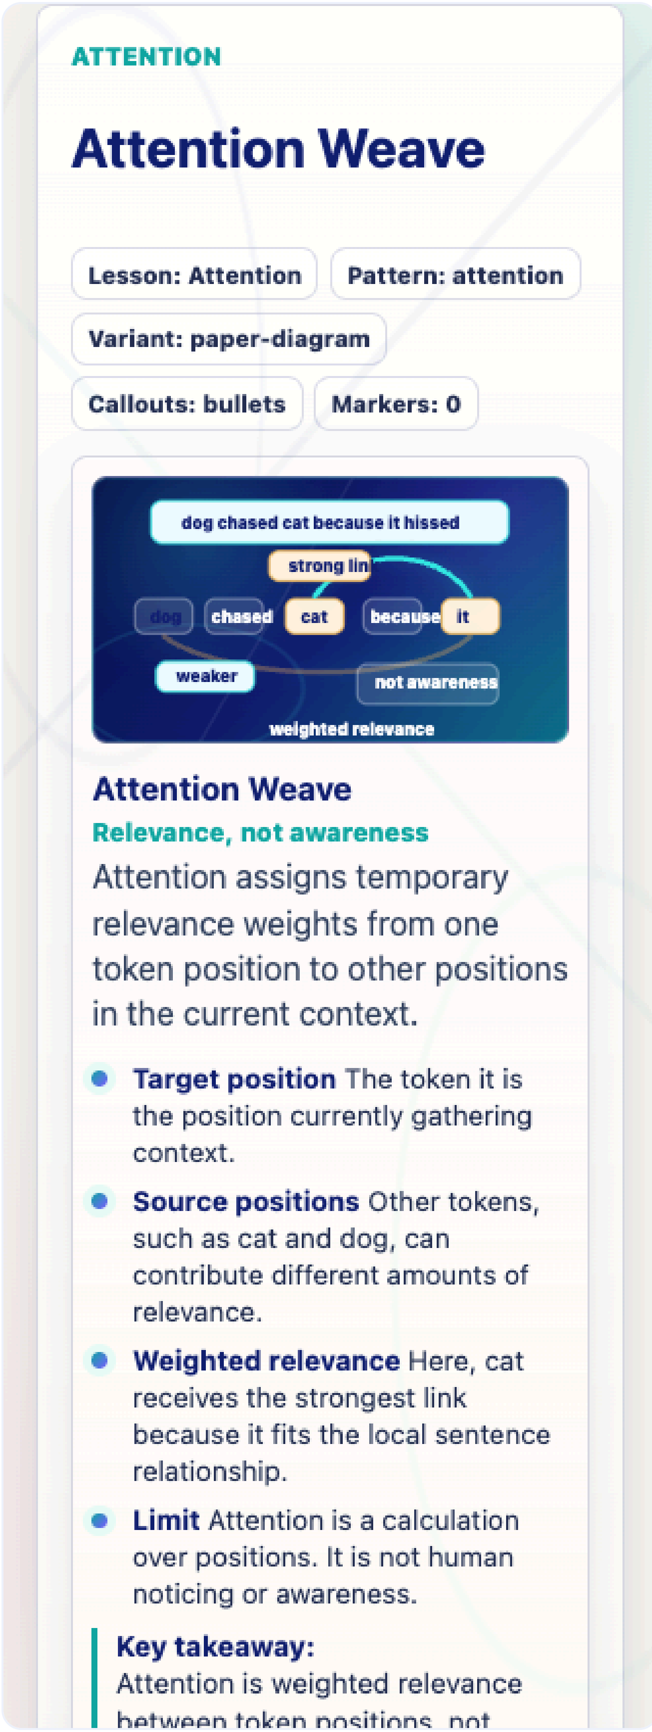

Attention 320px

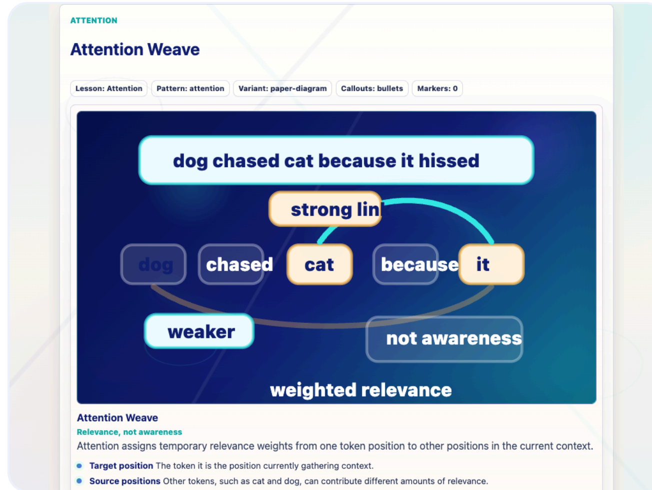

Visual aid desktop sample

Backlog

The machine-readable backlog is available as `visual-aid-quality-backlog-v0-27-17.json` and `visual-aid-quality-backlog-v0-27-17.csv`. It preserves the decision for every priority visual: change now, keep coded, or prepare future Image 2 generation.