

Prompt Life v0.27.18 Final Response With Screenshots

Generated June 12, 2026. Latest app version: 0.27.18.

39

visual aids audited

4

visual aids
structurally
changed

0

numbered callout
mismatches

5

remaining backlog
items

Summary

- Removed learner-facing renderer/debug metadata chips from visual-aid review UI.
- Renamed and rebuilt the Attention visual as **Relevance Between Tokens**.
- Raised diagram readability rules and removed compressed review-route visual text.
- Kept MLP, Layers, and Hidden States structurally, with readability repairs.
- Updated Logits, Softmax, and Sampling around the same candidate tokens: floor, mat, kitchen, banana.
- Created the renderer strategy inventory, readability audit, report PDF, and screenshot set.

Attention Visual

Title: Relevance Between Tokens

Caption: Attention assigns temporary relevance weights between token positions. It is not awareness.

Callouts: Target token: "it" is the token being updated. Source clue: "cat" fits the local sentence relationship. Weighted relevance: attention changes how strongly one token position uses information from another. Limit: attention is calculation over positions, not awareness.

Key takeaway: Attention is weighted relevance between token positions, not consciousness.

Rendering Strategy


Preferred: shared visual-aid frame with HTML callouts plus coded SVG/HTML for exact mechanisms. Generated PNGs stay useful for atmospheric visuals only, with instructional text in HTML.

Deprecated: renderer metadata chip wrapper in learner-facing UI.

Screenshots

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

LEARNING OBJECTIVE

Use a concrete sentence to show that attention weights token relevance without implying awareness.


ACCESSIBLE DESCRIPTION

The diagram shows sentence chips for dog

Attention, 390px

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.


LEARNING OBJECTIVE

Use a concrete sentence to show that attention weights token relevance without implying awareness.

Attention, 320px

MLP

MLP Feature Workshop



MLP Feature Workshop

Same token, reshaped features

After attention mixes across token positions, the MLP reshapes features within each token position.

- **Attention mixes** Attention lets token positions pull in context from other positions.
- **Same position** The example token cat stays in its own position as it enters the MLP.
- **MLP reshapes** The feature bars change inside that token position.
- **Temporary values** The MLP uses durable weights, but the feature values in this run are temporary activations.

Key takeaway:

Attention mixes across positions; the MLP reshapes features within one position.

LEARNING OBJECTIVE

Show the attention-versus-MLP distinction: attention mixes token positions; MLP reshapes each token vector independently.

ACCESSIBLE DESCRIPTION

The diagram contrasts an attention panel where

MLP, 390px

LAYERS

Transformer Stack



Transformer Stack

Repeated attention plus MLP

Layers repeat attention and MLP transformations to refine temporary hidden states.

- **Input states** Current-context hidden states enter the stack.
- **Repeated block** Each layer applies attention and an MLP to the temporary states.
- **Refined states** The states keep changing as they move toward next-token scores.
- **Limit** Layers are repeated numerical transformations, not conscious reasoning steps.

Key takeaway:

Layers refine hidden states during the run; they do not store permanent memory.

LEARNING OBJECTIVE

Show repeated transformer blocks without implying human thought steps.


ACCESSIBLE DESCRIPTION

A vertical stack shows input hidden states moving through three transformer blocks, each

Layers, 390px

HIDDEN STATES

Hidden State Flow



Hidden State Flow

Temporary during this forward pass

A token starts with an embedding and becomes context-shaped hidden states as it moves through layers.

- **Embedding** A token ID retrieves a learned starting vector.
- **Hidden states** Layers reshape that vector into temporary context-shaped states.
- **Scores later** Final hidden states help produce raw next-token scores.
- **Not memory** A hidden state is not visible text, permanent memory, or a stored fact.

Key takeaway:

Hidden states are temporary internal vectors shaped by current context.

LEARNING OBJECTIVE

Contrast embedding, hidden state, weight, memory, and visible text without making hidden states mysterious.

ACCESSIBLE DESCRIPTION

A flow moves from token ID to embedding, then through two temporary hidden-state panels

Hidden States, 390px

LOGITS

Raw Scoreboard



Raw Scoreboard

Before probabilities

Logits are raw next-token scores before probabilities.

- **Hidden state arrives** The final context-shaped vector reaches the output scoring step.
- **Candidates get scores** Possible next tokens such as floor, mat, kitchen, and banana receive raw logits.
- **Scores rank candidates** Higher score means a stronger local candidate in this context.
- **Softmax comes next** Softmax converts raw scores into probabilities.

Key takeaway:

Logits are raw next-token scores, not probabilities.

LEARNING OBJECTIVE

Show final hidden state becoming raw vocabulary scores before softmax.


ACCESSIBLE DESCRIPTION

The diagram starts with a final hidden state flowing into raw score bars for floor, mat,

Logits, 390px

SOFTMAX

Score to Probability



Score to Probability

Normalize before choosing

Softmax turns raw scores into probabilities that sum to one.

- **Logits go in** Raw scores for floor, mat, kitchen, and banana enter softmax.
- **Softmax converts** The function turns score differences into a probability distribution; it does not pick the next token.
- **Probabilities come out** The rounded teaching values sum to 100%.
- **Sampling comes next** The probability distribution is ready for a separate choice step.

Key takeaway:

Softmax normalizes raw scores; sampling chooses later.

LEARNING OBJECTIVE

Show raw-score values becoming rounded probabilities that sum to one.

ACCESSIBLE DESCRIPTION

The diagram shows raw score bars for floor, mat, kitchen, and banana on the left, a softmax step in the middle, and probabilities on the right.

Softmax, 390px

SAMPLING

Weighted Token Choice



Weighted Token Choice

One token chosen

Sampling chooses one next token from the probability-shaped options. Likely is not the same as true.

- **Probability candidates** floor is strongest, mat is plausible, kitchen is weaker, and banana is unlikely in this context.
- **One token chosen** The decoding step chooses one next response token.
- **Append next** The chosen token is appended before the model runs again.
- **Fit is not proof** A likely token can fit the model pattern and still need evidence in a factual claim.

Key takeaway:

Sampling chooses one token from probabilities; likely is not the same as true.

LEARNING OBJECTIVE

Show probability weights and the single selected next token before autoregression repeats.

ACCESSIBLE DESCRIPTION

Sampling, 390px

FINE-TUNING

Fine-Tuning Path



Fine-Tuning Path

Durable adaptation, not current-run steering

Fine-tuning can durably adapt model behavior through additional training or adapters; prompts, RAG, and sampling steer only the current run.

- **Durable adaptation** Additional training or adapters can change how future runs behave.
- **Prompt steering** A prompt can steer this run without changing weights.
- **RAG and context** Retrieved text can enter the current context without training the model.
- **Sampling later** Decoding still chooses one next token from probabilities during inference.


Key takeaway:

Fine-tuning changes future behavior; prompting, RAG, and sampling operate during the current run.

Atmospheric PNG sample, 390px

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Desktop visual aid sample

PLAY LAB

Play to learn

Short, calm challenges that make model mechanics visible.

Progress saved on this device. 0 of 5 practice challenges show saved completion or review progress.

Each challenge practices one model-literacy move.

More below ↓

Play challenges

Suggested order follows the Journey.



Glossary Dojo

Ready

Practice concept discrimination.

Practice move: Name concepts

Recommended: Anytime

Anytime

Glossary

Learning cards: All learning cards

Action: Choose, compare, review.

Time: 4 min



Home



Journey



Play



Glossary



Badge

Play smoke, 390px

BADGE

Prompt Life: Model Literate

Under construction while the evidence model is reviewed.



Under construction

Pending human review

Prompt Life: Model Literate is under construction. The Journey, checkpoint bank, and evidence model are being tested before any badge is issued.



Home



Journey



Play



Glossary



Badge

Badge smoke, 390px

Verification

- `npm run typecheck` : passed.
- `npm run build` : passed sequentially.
- `npm run build:pages` : passed sequentially.
- `npm run audit:answers` : passed.
- `npm run audit:checkpoints` : passed.
- `npm run audit:question-clues` : completed; informational warnings remain for human review.
- `npm run audit:learner-copy` : passed.
- `npm run audit:visual-aids` : passed with 39 aids, 0 mismatches, 0 readability issues.
- Browser QA: 320px/390px mobile and desktop visual review passed; all eight Journey stage links still route/focus correctly.

