

Prompt Life v0.27.18 Visual Renderer Readability Report

Generated 2026-06-12T17:15:49.773Z. This pass removes learner-facing renderer metadata, raises label readability standards, rebuilds the Attention visual, and repairs the Decision Room mechanism sequence while preserving Journey, Play, checkpoints, progress, and the one-badge model.

8

renderer strategies found

4

visual aids structurally changed

39

visual aids affected by metadata cleanup

5

remaining backlog items

Executive Summary

- Renderer strategies kept: generated PNG with HTML callouts, inline coded SVG, HTML/CSS diagram cards, DiagramKit SVG examples, shared visual-aid frame, and custom SVG primitives.
- Deprecated for learner UI: renderer metadata chip wrapper.
- Metadata UI removed from visual aid review cards: no Pattern, Variant, Callouts, Markers, renderer, or implementation chips.
- Attention was renamed from Attention Weave to Relevance Between Tokens and rebuilt as a clean coded SVG.
- Workday cluster status: MLP, Layers, and Hidden States were kept structurally but benefit from larger label standards and metadata removal.
- Decision Room cluster status: Logits, Softmax, and Sampling now share floor, mat, kitchen, banana as candidate tokens and clearer captions.

Renderer Strategy Inventory

Strategy	Status	Used by	Recommendation
Generated PNG image asset with HTML callouts	keep-with-limits	Prompt to Prediction, Broad Pretraining, Fine-Tuning Path, Alignment Landscape	Keep for atmospheric aids only. Do not use for exact mechanism diagrams.
Inline coded SVG	preferred-for-mechanisms	Forward Pass, Token IDs, Embeddings, Vectors, Tensors, Relevance Between Tokens, MLP, Layers, Hidden States, Logits, Softmax, Sampling, Autoregression, Context Window, RAG, Grounding, Hallucination, Diffusion,	Preferred for model-mechanism diagrams, with strict label limits and HTML callouts for nuance.

Strategy	Status	Used by	Recommendation
		Multimodal, Perfect Storm, Twilight ethics/risk visuals	
HTML/CSS diagram card	keep	Prompt vs Response, Text to Tokens	Keep for text-heavy token/card layouts. Consider for future token-stream visuals.
DiagramKit React SVG examples	keep	AI Family Tree, Rules and Learned Patterns, Training Loop, Overfitting vs Generalization	Keep, but follow the v0.27.18 minimum label-size rules.
Shared visual aid frame with HTML caption, callouts, key takeaway, and accessible description	preferred-wrapper	All 39 Journey visual aids	Preferred wrapper for all Journey visual aids.
Renderer metadata chip wrapper	deprecated-for-ui	Former review-route visual aid cards	Deprecated from learner-facing UI. Keep renderer metadata in docs, JSON, and reports only.
Custom SVG primitives	keep-and-standardize	Label, Arrow, Callout, TokenChip, StepNode, diagram-specific bars/arcs/cards	Keep and consolidate into shared styles; avoid new one-off tiny labels.
Third-party visualization helper	not-used	None	Do not introduce for current Journey visual aids.


Attention Before / After

Old issue summary: Attention Weave used legacy wording, visible renderer metadata cluttered the review screen, and the diagram carried too many competing labels.

New diagram summary: Relevance Between Tokens uses six sentence chips, a highlighted target token, a highlighted source clue, one stronger arc, one weaker arc, and four HTML callouts below the visual.

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

LEARNING OBJECTIVE

Use a concrete sentence to show that attention weights token relevance without implying awareness.

ACCESSIBLE DESCRIPTION

The diagram shows sentence chips for dog

Attention visual at 390px

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

LEARNING OBJECTIVE

Use a concrete sentence to show that attention weights token relevance without implying awareness.

Attention visual at 320px


Cluster Status

Visual	Current title	Status	Action
Attention	Relevance Between Tokens	changed	Rebuilt coded SVG and replaced legacy Attention Weave title.
MLP	MLP Feature Workshop	kept-with-readability-repair	Kept structure; larger label standard and metadata removal improve readability.
Layers	Transformer Stack	kept-with-readability-repair	Kept repeated-block layout; larger labels and no metadata chips.
Hidden States	Hidden State Flow	kept-with-readability-repair	Kept temporary-state flow; larger labels and no metadata chips.
Logits	Raw Scoreboard	changed	Updated caption and token set; simplified raw-score labels.
Softmax	Score to Probability	changed	Updated caption and token set; kept softmax separate from sampling.
Sampling	Weighted Token Choice	changed	Replaced bubble cluster with readable probability bars and selected-token card.

Mechanism Screenshots

MLP

MLP Feature Workshop



MLP Feature Workshop

Same token, reshaped features

After attention mixes across token positions, the MLP reshapes features within each token position.

- **Attention mixes** Attention lets token positions pull in context from other positions.
- **Same position** The example token cat stays in its own position as it enters the MLP.
- **MLP reshapes** The feature bars change inside that token position.
- **Temporary values** The MLP uses durable weights, but the feature values in this run are temporary activations.

Key takeaway:

Attention mixes across positions; the MLP reshapes features within one position.

LEARNING OBJECTIVE

Show the attention-versus-MLP distinction: attention mixes token positions; MLP reshapes each token vector independently.

ACCESSIBLE DESCRIPTION

The diagram contrasts an attention panel where

MLP at 390px

LAYERS

Transformer Stack



Transformer Stack

Repeated attention plus MLP

Layers repeat attention and MLP transformations to refine temporary hidden states.

- **Input states** Current-context hidden states enter the stack.
- **Repeated block** Each layer applies attention and an MLP to the temporary states.
- **Refined states** The states keep changing as they move toward next-token scores.
- **Limit** Layers are repeated numerical transformations, not conscious reasoning steps.

Key takeaway:

Layers refine hidden states during the run; they do not store permanent memory.

LEARNING OBJECTIVE

Show repeated transformer blocks without implying human thought steps.


ACCESSIBLE DESCRIPTION

A vertical stack shows input hidden states moving through three transformer blocks, each

Layers at 390px

HIDDEN STATES

Hidden State Flow



Hidden State Flow

Temporary during this forward pass

A token starts with an embedding and becomes context-shaped hidden states as it moves through layers.

- **Embedding** A token ID retrieves a learned starting vector.
- **Hidden states** Layers reshape that vector into temporary context-shaped states.
- **Scores later** Final hidden states help produce raw next-token scores.
- **Not memory** A hidden state is not visible text, permanent memory, or a stored fact.

Key takeaway:

Hidden states are temporary internal vectors shaped by current context.

LEARNING OBJECTIVE

Contrast embedding, hidden state, weight, memory, and visible text without making hidden states mysterious.

ACCESSIBLE DESCRIPTION

A flow moves from token ID to embedding, then through two temporary hidden-state panels

Hidden States at 390px

LOGITS

Raw Scoreboard



Raw Scoreboard

Before probabilities

Logits are raw next-token scores before probabilities.

- **Hidden state arrives** The final context-shaped vector reaches the output scoring step.
- **Candidates get scores** Possible next tokens such as floor, mat, kitchen, and banana receive raw logits.
- **Scores rank candidates** Higher score means a stronger local candidate in this context.
- **Softmax comes next** Softmax converts raw scores into probabilities.

Key takeaway:

Logits are raw next-token scores, not probabilities.

LEARNING OBJECTIVE

Show final hidden state becoming raw vocabulary scores before softmax.


ACCESSIBLE DESCRIPTION

The diagram starts with a final hidden state flowing into raw scores here for floor, mat,

Logits at 390px

SOFTMAX

Score to Probability



Score to Probability

Normalize before choosing

Softmax turns raw scores into probabilities that sum to one.

- **Logits go in** Raw scores for floor, mat, kitchen, and banana enter softmax.
- **Softmax converts** The function turns score differences into a probability distribution; it does not pick the next token.
- **Probabilities come out** The rounded teaching values sum to 100%.
- **Sampling comes next** The probability distribution is ready for a separate choice step.

Key takeaway:

Softmax normalizes raw scores; sampling chooses later.

LEARNING OBJECTIVE

Show raw-score values becoming rounded probabilities that sum to one.

ACCESSIBLE DESCRIPTION

The diagram shows raw score bars for floor, mat, kitchen, and banana on the left, a softmax step in the middle, and probabilities on the right.

Softmax at 390px

SAMPLING

Weighted Token Choice



Weighted Token Choice

One token chosen

Sampling chooses one next token from the probability-shaped options. Likely is not the same as true.

- **Probability candidates** floor is strongest, mat is plausible, kitchen is weaker, and banana is unlikely in this context.
- **One token chosen** The decoding step chooses one next response token.
- **Append next** The chosen token is appended before the model runs again.
- **Fit is not proof** A likely token can fit the model pattern and still need evidence in a factual claim.

Key takeaway:

Sampling chooses one token from probabilities; likely is not the same as true.

LEARNING OBJECTIVE

Show probability weights and the single selected next token before autoregression repeats.

ACCESSIBLE DESCRIPTION

The diagram shows the probability distribution and the selected token.

Sampling at 390px

FINE-TUNING

Fine-Tuning Path



Fine-Tuning Path

Durable adaptation, not current-run steering

Fine-tuning can durably adapt model behavior through additional training or adapters; prompts, RAG, and sampling steer only the current run.

- **Durable adaptation** Additional training or adapters can change how future runs behave.
- **Prompt steering** A prompt can steer this run without changing weights.
- **RAG and context** Retrieved text can enter the current context without training the model.
- **Sampling later** Decoding still chooses one next token from probabilities during inference.


Key takeaway:

Fine-tuning changes future behavior; prompting, RAG, and sampling operate during the current run.

Atmospheric PNG sample at 390px

ATTENTION

Relevance Between Tokens



Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

Desktop visual aid sample

Remaining Backlog

- **Collective Intelligence, Extracted:** future Image 2 atmospheric replacement (medium)
- **Benefits Worth Taking Seriously:** future Image 2 atmospheric replacement (medium)
- **Costs We Must Count:** future Image 2 atmospheric replacement (medium)
- **Human-Centered AI:** future Image 2 atmospheric replacement (medium)
- **Responsible AI Control Panel:** coded simplification/backlog from v0.27.15 (low)

Audit Results

- Visual aid consistency: 0 numbered marker/list mismatches.
- Visual aid readability: 0 readability issues.
- Learner-copy audit: run in verification; no normal learner UI debug/developer labels expected.
- Stage-link regression: verified by browser QA in this pass.
- Mobile QA: 320px and 390px screenshots captured.