Prompt Life v0.28.0 Visual Aid Readiness Pass

Generated June 13, 2026. This packet wraps the final response and representative screenshots for the Journey visual-aid human-testing pass.

39 visual aids reviewed	**18** P0/P1 before	**0** P0/P1 after	**6** templates used

Summary

- Added six canonical visual templates to the style guide.
- Scored all 39 Journey visual aids using a seven-part human-readability rubric.
- Marked 13 visuals as temporary but acceptable for small human testing.
- Created the full visual-aid human-test review PDF and feedback sheet PDF.
- Updated `audit:visual-aids` to catch readiness issues, not just marker mismatches.
- Bumped the app/package version to `0.28.0`.

Template Distribution

- **Atmospheric Scene**: 9
- **Taxonomy Map**: 6
- **Comparison Board**: 8
- **Mechanism Flow**: 9
- **Probability Bars**: 3
- **Context Tray / Stack**: 4

Human-Testing Watch List

- **Full Chain Map** (24/35): This is the top later-coded redesign candidate; testers should flag overload.
- **Borrowed Flames** (25/35): Avoid cramming source labels; social meaning belongs in callouts.
- **Cost Ledger** (25/35): No fake statistics; keep categories in HTML.
- **Benefit Tiers** (25/35): Avoid implying all benefits are proven or automatic.
- **Accountability Flow** (25/35): Avoid generic AI-person imagery that implies the model has agency.
- **Learning Modes Matrix** (26/35): Watch whether matrix labels feel too small at 320px.
- **Media Lane Map** (26/35): Reduce nodes later if testers call it busy.
- **Better-AI Control Panel** (26/35): Panel may feel dense; testers should flag unreadable lever labels.
- **Fine-Tuning Path** (27/35): Tester should check whether the metaphor needs a stronger coded comparison later.
- **Feature Vector** (27/35): Distributed side may feel abstract; tester feedback should guide simplification.

Representative Screenshots

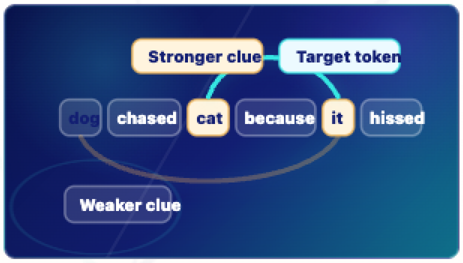

Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

> **Key takeaway:**
> Attention is weighted relevance between token positions, not consciousness.

Attention / Relevance Between Tokens, 390px

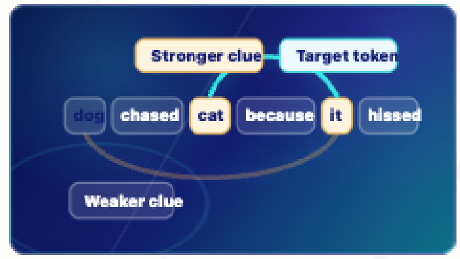

Relevance Between Tokens

Weighted relevance, not awareness

Attention assigns temporary relevance weights between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Source clue** "cat" fits the local sentence relationship.
- **Weighted relevance** Attention changes how strongly one token position uses information from another.
- **Limit** Attention is calculation over positions, not awareness.

> **Key takeaway:**
> Attention is weighted relevance between token positions, not consciousness.

Attention, 320px

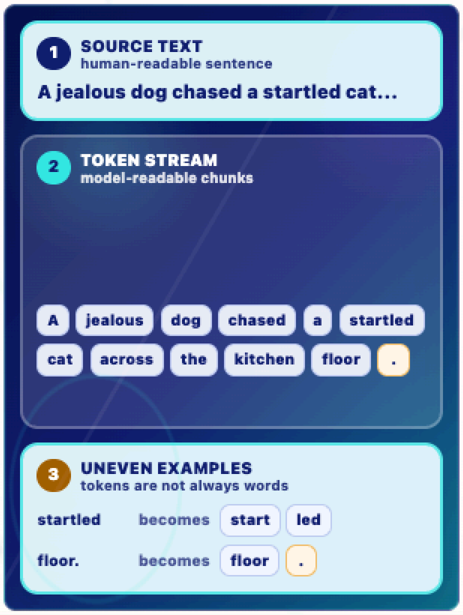

Text to Tokens

Uneven chunks, punctuation included

Text is split into model-readable chunks before token IDs and embedding lookup.

1. **Simplified split** A | jealous | dog | chased | a | startled | cat | across | the | kitchen | floor | .
2. **Uneven chunks** Some tokenizers may split words or punctuation, such as start | led or floor | .
3. **Teaching note** This app uses a simplified split; real tokenizer output can differ.

> **Key takeaway:**
> Tokens are model-readable chunks, not always human words.

Mechanism Flow: Tokenization

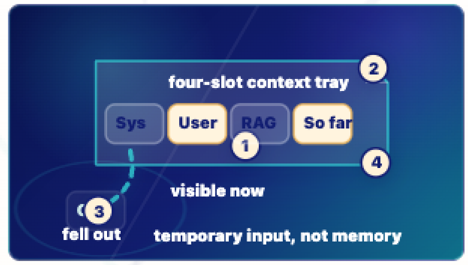

Temporary Context Window

Limited visible input

Only the cards still inside the context tray can influence the next token.

1. **Visible now** System instructions, the user prompt, retrieved notes, and response-so-far may all be current context.
2. **Limited tray** When the tray fills, older material can fall out or need summarizing by the surrounding system.
3. **Temporary** Context can feel like memory, but it does not normally update weights or create permanent memory.
4. **Next token** The model can use only what remains visible in the current run.

> **Key takeaway:**
> Context is temporary visible input, not durable memory.

Context Tray / Stack

Weighted Token Choice

One token chosen

Sampling chooses one next token from the probability-shaped options. Likely is not the same as true.

- **Probability candidates** floor is strongest, mat is plausible, kitchen is weaker, and banana is unlikely in this context.
- **One token chosen** The decoding step chooses one next response token.
- **Append next** The chosen token is appended before the model runs again.
- **Fit is not proof** A likely token can fit the model pattern and still need evidence in a factual claim.

> **Key takeaway:**
> Sampling chooses one token from probabilities; likely is not the same as true.

Probability Bars: Sampling

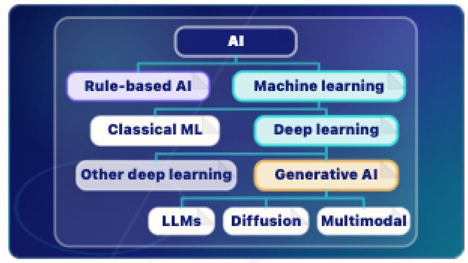

AI Family Tree

Where LLMs fit

A clean taxonomy tree shows AI as the broad field, machine learning as one branch inside AI, and LLMs as one branch inside generative AI.

- **AI** AI is the broad field.
- **Machine learning** Machine learning systems learn patterns from data.
- **Deep learning** Deep learning is a neural-network branch of machine learning.
- **Generative AI** Generative AI creates new media, such as text, images, audio, code, or video.
- **LLMs** LLMs generate language/code; diffusion and multimodal systems are neighboring generative AI branches.

> **Key takeaway:**
> An LLM is one kind of generative AI inside the broader AI family.

Taxonomy Map: AI Family Tree

Prompt to Prediction

Score, choose, append, repeat

Current context enters the model, learned weights shape next-token probabilities, one token is generated, and the response grows token by token.

- **Context enters** The current prompt and response-so-far become the visible input for this run.
- **Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- **One token is generated** The model chooses one response token from the probability cloud.
- **Response grows** That token is appended, then the model runs again.

Key takeaway:
An LLM turns context into next-token probabilities.

Atmospheric Scene

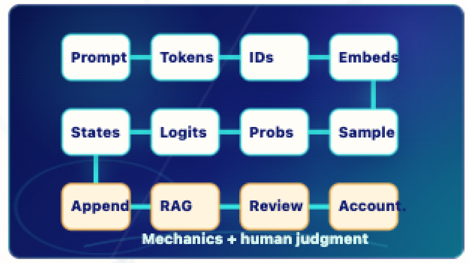

Full Chain Map

Mechanics plus judgment

Model literacy connects prompt, tokens, IDs, embeddings, hidden states, logits, probabilities, sampling, append-and-repeat generation, RAG, grounding, review, and accountability.

- **Machine chain** Prompt context becomes tokens and IDs, then embeddings, hidden states, logits, probabilities, and sampled response tokens.
- **Evidence chain** Context windows, RAG, and grounding can add evidence for this run, but they do not guarantee truth.
- **Human chain** Benefits, costs, review, governance, and accountability remain human responsibilities.

Key takeaway:
Mechanics matter, and humans remain responsible.

Top watch-list visual: Full Chain Map

PLAY LAB

Play to learn

Short, calm challenges that make model mechanics visible.

Progress saved on this device. 0 of 5 practice challenges show saved completion or review progress.

Each challenge practices one model-literacy move.

More below ↓

Play challenges

Suggested order follows the Journey.

Glossary Dojo

Ready

Practice concept discrimination.

Practice move: Name concepts

Recommended: Anytime

Anytime Glossary

Learning cards: All learning cards

Action: Choose, compare, review.

Time: 4 min

Played: 0

Progress: Not started

Home Journey Play Glossary Badge

Play smoke

BADGE

Prompt Life: Model Literate

Under construction while the evidence model is reviewed.

Under construction

Pending human review

Prompt Life: Model Literate is under construction. The Journey, checkpoint bank, and evidence model are being tested before any badge is issued.

Your progress below is practice evidence, not an issued credential.

Home Journey Play Glossary Badge

Badge smoke