Prompt Life v0.28.3 Flowchart Replacement Report

Generated 2026-06-14T00:22:42.590Z. This packet documents the hand-specified mobile diagrams that replaced the target flowchart-like Journey visuals.

Summary

- Visuals changed: 10
- Flowchart-like target visuals replaced: 10
- Templates used: Boundary Board, Vertical Mechanism Strip, Tray / Stack / Bars
- Flowchart replacement audit: ok
- Visual overflow audit: pass

Before / After Table

VISUAL	OLD ISSUE	NEW TEMPLATE	NEW TITLE	NEW CAPTION	NEW CALLOUTS
Fine-Tuning	Dense durable/current-run comparison risked reading like a flowchart.	Boundary Board	Durable vs Temporary Steering	Fine-tuning or adapters can shape future responses. Prompting, retrieval, and sampling shape the current run.	Additional training or adapters can carry a pattern into later runs. A prompt steers the current context. Outside material helps only after it enters context. Sampling chooses one next token; it does not train the model.
Inference	Multiple boxes and arrows made fixed weights versus temporary activations feel cramped.	Vertical Mechanism Strip	Forward Pass	Inference uses fixed learned weights and current context to produce next-token scores.	What the model can currently use. Learned numbers from training. Internal values created for this run. Raw next-token scores. One token is selected and appended.

VISUAL	OLD ISSUE	NEW TEMPLATE	NEW TITLE	NEW CAPTION	NEW CALLOUTS
Attention	SVG labels and arcs were vulnerable to endpoint drift on narrow screens.	Boundary Board	Relevance Between Tokens	Attention weights relevance between token positions. It is not awareness.	"it" is the token being updated. "cat" fits the local sentence relationship. "dog" is nearby context but less useful here. Attention is calculation over positions, not consciousness.
MLP	Workshop-style visual tried to show too much mechanism inside one small panel.	Boundary Board	Attention vs MLP	Attention mixes information across token positions. The MLP reshapes each token position's features.	Asks which positions matter to this token. Transforms the token's feature vector at that position. Both help update hidden states inside a layer.
Layers	Stack flow compressed repeated layer behavior into small labels.	Vertical Mechanism Strip	Transformer Layer Strip	A transformer layer refines hidden states; many layers repeat the same kind of update.	Temporary internal vectors. Mixes relevant token-position information. Reshapes features. Later layers refine the representation again.
Hidden States	State-flow diagram blurred embedding, hidden state, and durable weight boundaries.	Boundary Board	Embedding, Hidden State, Weight	A hidden state is temporary and context-shaped. It is not the same as an embedding or a durable weight.	The learned starting vector for a token ID. The token's temporary representation after context processing. Learned model parameter reused across runs.

VISUAL	OLD ISSUE	NEW TEMPLATE	NEW TITLE	NEW CAPTION	NEW CALLOUTS
Vectors	Teaching-slider comparison risked implying each vector dimension has a neat human label.	Tray / Stack / Bars	Feature Vector	A vector represents something with many numbers at once.	Meaning is not stored in one neat dimension. Useful patterns are spread across many values. Vectors let the model compute with text-like inputs.
Tensors	Axis-heavy grid carried more labels than the mobile diagram needed.	Tray / Stack / Bars	Tensor Block	A tensor is a shaped block of numbers, often holding many token vectors at once.	One list of numbers. Many vectors arranged in a shape. Rows, columns, and batches organize computation.
Context Window	Tray metaphor worked, but SVG labels and markers added avoidable density.	Tray / Stack / Bars	Temporary Context Window	The context window is the limited visible workspace for the current run.	Only current context can shape the next token. Old content can stop being visible. Outside evidence helps only after it enters context.
Model Literate Synthesis	Full-chain map tried to recap too many mechanisms at once.	Boundary Board	Model Literacy Board	Model literacy connects how the system works with how humans should use it.	Prompts become tokens, scores, probabilities, and generated tokens. Fluent output still needs checking. Humans remain accountable for use, review, and decisions.

Screenshots

Fine-Tuning: Durable vs Temporary Steering

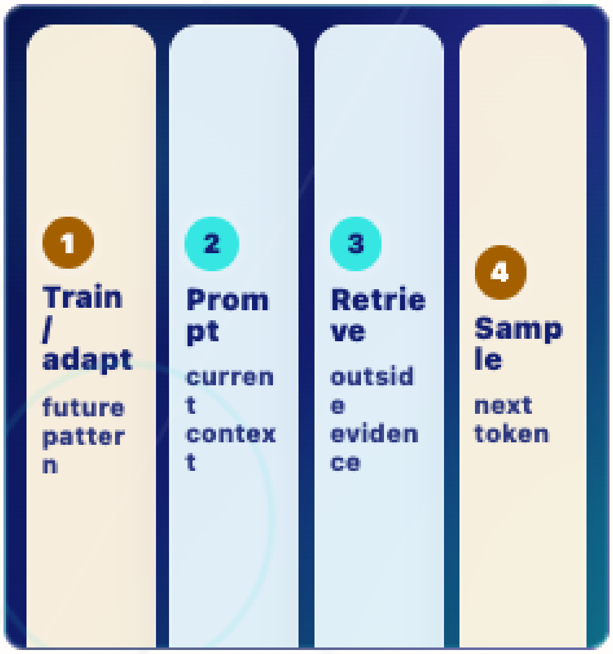

Durable vs Temporary Steering

Fine-tuning is not prompting

Fine-tuning or adapters can shape future responses. Prompting, retrieval, and sampling shape the current run.

1. **Durable adaptation** Additional training or adapters can carry a pattern into later runs.
2. **Prompting** A prompt steers the current context.
3. **Retrieval** Outside material helps only after it enters context.
4. **Sampling** Sampling chooses one next token; it does not train the model.

> **Key takeaway:**
> Fine-tuning can shape future behavior; prompting, retrieval, and sampling shape this run.

320px

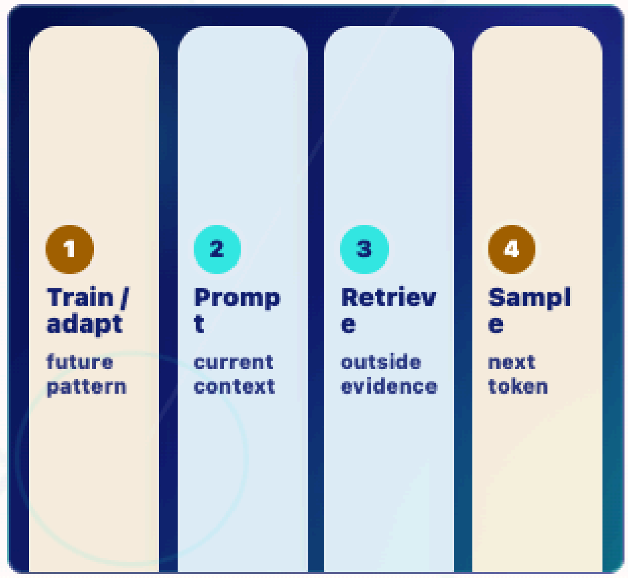

Durable vs Temporary Steering

Fine-tuning is not prompting

Fine-tuning or adapters can shape future responses. Prompting, retrieval, and sampling shape the current run.

1. **Durable adaptation** Additional training or adapters can carry a pattern into later runs.
2. **Prompting** A prompt steers the current context.
3. **Retrieval** Outside material helps only after it enters context.
4. **Sampling** Sampling chooses one next token; it does not train the model.

> **Key takeaway:**
> Fine-tuning can shape future behavior; prompting, retrieval, and sampling shape this run.

390px

Inference: Forward Pass

1 Context
2 Fixed weights
3 Temporary activations
4 Scores
5 Next token

Forward Pass

Fixed weights, temporary activations

Inference uses fixed learned weights and current context to produce next-token scores.

1. **Context** What the model can currently use.
2. **Fixed weights** Learned numbers from training.
3. **Temporary activations** Internal values created for this run.
4. **Scores** Raw next-token scores.
5. **Next token** One token is selected and appended.

> **Key takeaway:**
> Inference uses weights without normally updating them.

320px

1 Context
2 Fixed weights
3 Temporary activations
4 Scores
5 Next token

Forward Pass

Fixed weights, temporary activations

Inference uses fixed learned weights and current context to produce next-token scores.

1. **Context** What the model can currently use.
2. **Fixed weights** Learned numbers from training.
3. **Temporary activations** Internal values created for this run.
4. **Scores** Raw next-token scores.
5. **Next token** One token is selected and appended.

> **Key takeaway:**
> Inference uses weights without normally updating them.

390px

Attention: Relevance Between Tokens

Relevance Between Tokens

Weighted relevance, not awareness

Attention weights relevance between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Strong clue** "cat" fits the local sentence relationship.
- **Weaker clue** "dog" is nearby context but less useful here.
- **Limit** Attention is calculation over positions, not consciousness.

> **Key takeaway:**
> Attention is weighted relevance between token positions, not consciousness.

320px

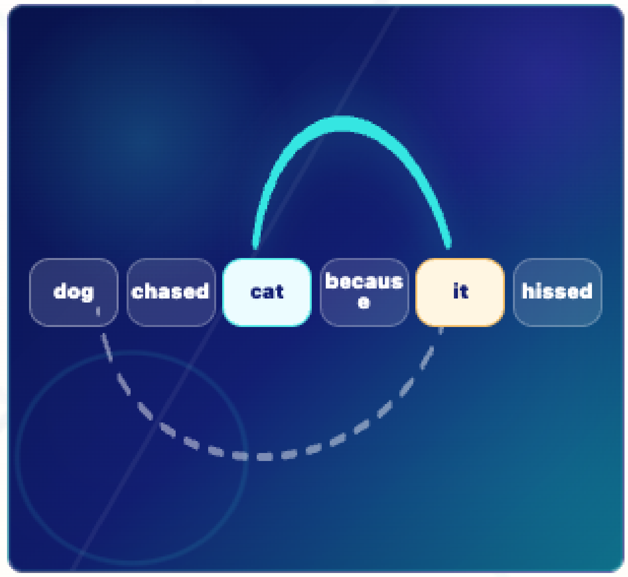

Relevance Between Tokens

Weighted relevance, not awareness

Attention weights relevance between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Strong clue** "cat" fits the local sentence relationship.
- **Weaker clue** "dog" is nearby context but less useful here.
- **Limit** Attention is calculation over positions, not consciousness.

> **Key takeaway:**
> Attention is weighted relevance between token positions, not consciousness.

390px

MLP: Attention vs MLP

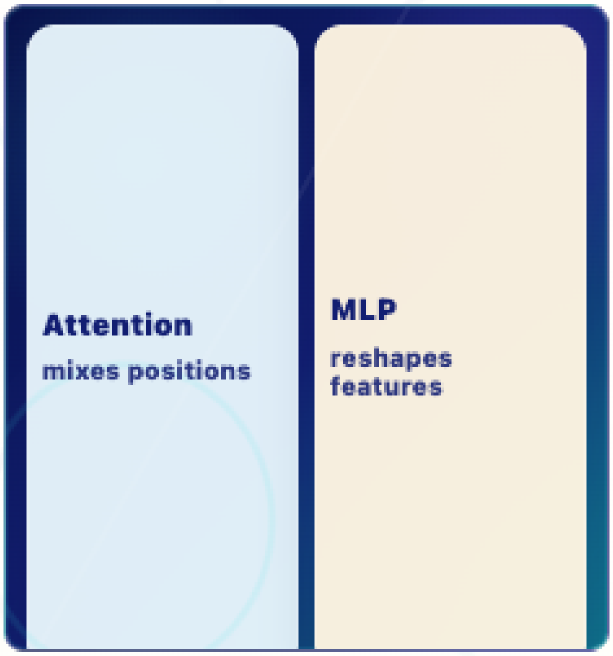

Attention vs MLP

Mix positions or reshape features

Attention mixes information across token positions. The MLP reshapes each token position's features.

- **Attention** Asks which positions matter to this token.
- **MLP** Transforms the token's feature vector at that position.
- **Together** Both help update hidden states inside a layer.

> **Key takeaway:**
> Attention mixes positions; MLP reshapes features.

320px

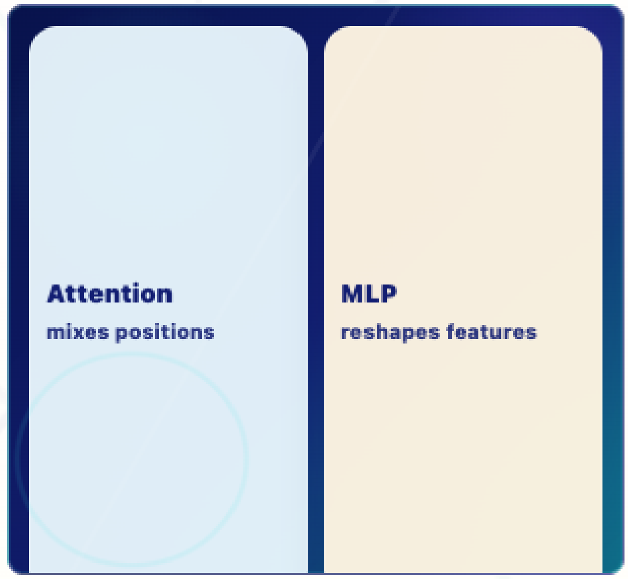

Attention vs MLP

Mix positions or reshape features

Attention mixes information across token positions. The MLP reshapes each token position's features.

- **Attention** Asks which positions matter to this token.
- **MLP** Transforms the token's feature vector at that position.
- **Together** Both help update hidden states inside a layer.

> **Key takeaway:**
> Attention mixes positions; MLP reshapes features.

390px

Layers: Transformer Layer Strip

1. Hidden states
2. Attention
3. MLP
4. Updated states
5. Repeat layer

Transformer Layer Strip

Refine, then repeat

A transformer layer refines hidden states; many layers repeat the same kind of update.

- **Hidden states** Temporary internal vectors.
- **Attention** Mixes relevant token-position information.
- **MLP** Reshapes features.
- **Repeat** Later layers refine the representation again.

> **Key takeaway:**
> Layers refine hidden states during the run; they do not store permanent memory.

320px

1. Hidden states
2. Attention
3. MLP
4. Updated states
5. Repeat layer

Transformer Layer Strip

Refine, then repeat

A transformer layer refines hidden states; many layers repeat the same kind of update.

- **Hidden states** Temporary internal vectors.
- **Attention** Mixes relevant token-position information.
- **MLP** Reshapes features.
- **Repeat** Later layers refine the representation again.

> **Key takeaway:**
> Layers refine hidden states during the run; they do not store permanent memory.

390px

Hidden States: Embedding, Hidden State, Weight

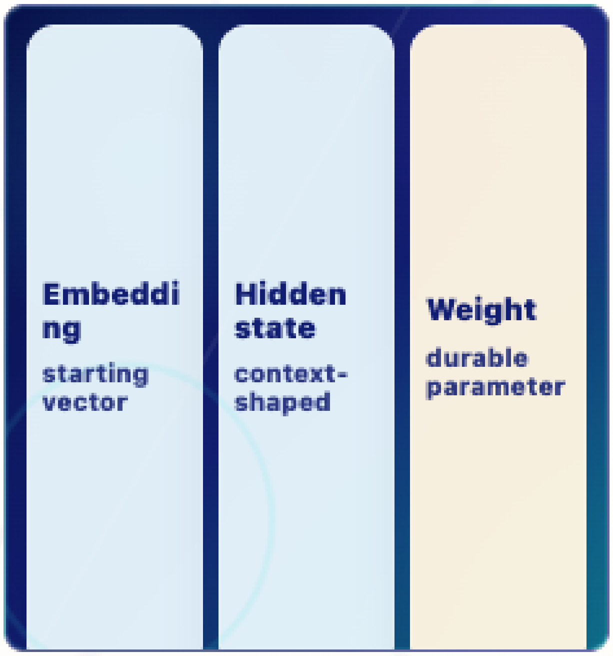

Embedding, Hidden State, Weight

Temporary is not durable

A hidden state is temporary and context-shaped. It is not the same as an embedding or a durable weight.

- **Embedding** The learned starting vector for a token ID.
- **Hidden state** The token's temporary representation after context processing.
- **Weight** Learned model parameter reused across runs.

> **Key takeaway:**
> Hidden states are temporary internal vectors shaped by current context.

320px

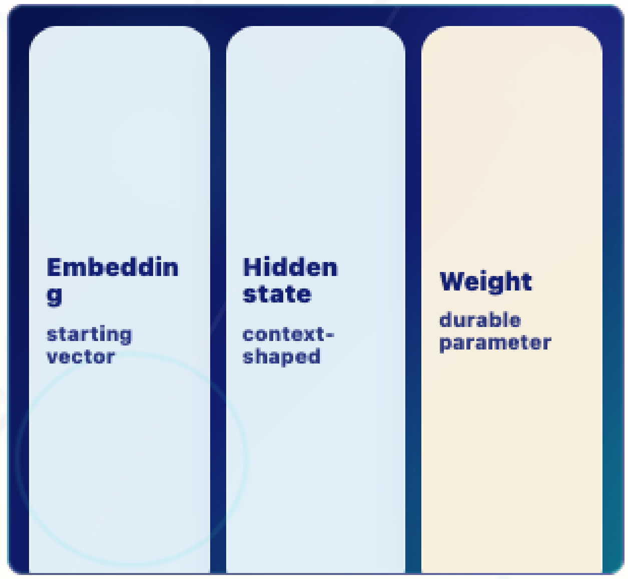

Embedding, Hidden State, Weight

Temporary is not durable

A hidden state is temporary and context-shaped. It is not the same as an embedding or a durable weight.

- **Embedding** The learned starting vector for a token ID.
- **Hidden state** The token's temporary representation after context processing.
- **Weight** Learned model parameter reused across runs.

> **Key takeaway:**
> Hidden states are temporary internal vectors shaped by current context.

390px

Vectors: Feature Vector

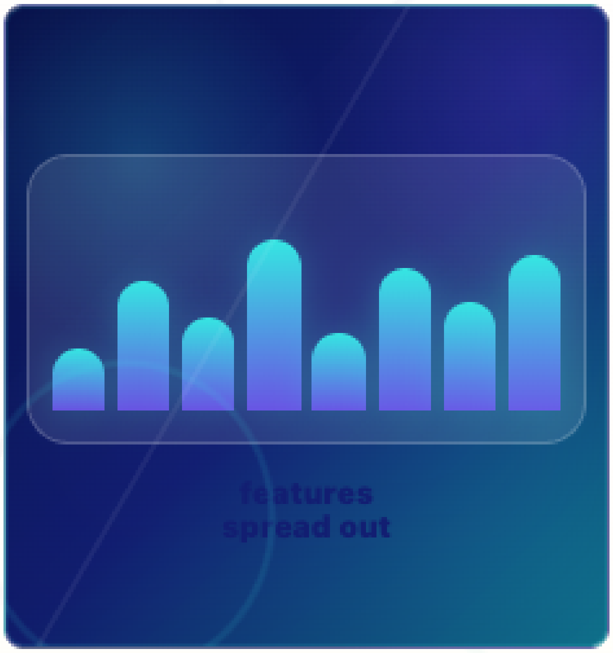

Feature Vector

Many numbers at once

A vector represents something with many numbers at once.

1. **Not one slot** Meaning is not stored in one neat dimension.
2. **Distributed features** Useful patterns are spread across many values.
3. **Model-readable** Vectors let the model compute with text-like inputs.

> **Key takeaway:**
> A vector carries distributed numerical features.

320px

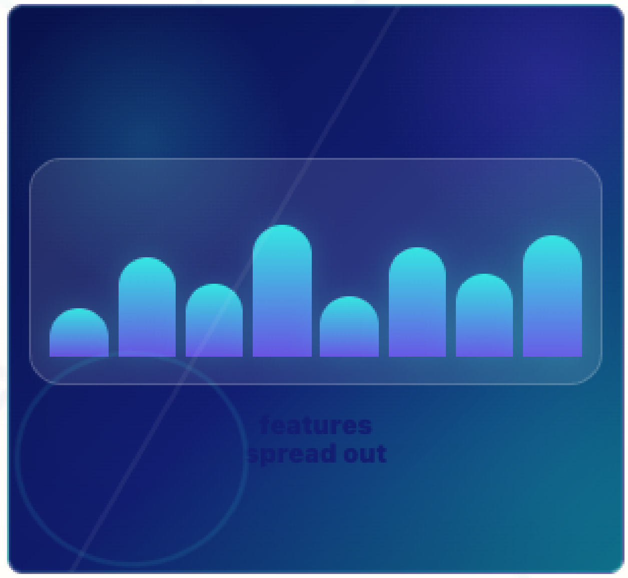

Feature Vector

Many numbers at once

A vector represents something with many numbers at once.

1. **Not one slot** Meaning is not stored in one neat dimension.
2. **Distributed features** Useful patterns are spread across many values.
3. **Model-readable** Vectors let the model compute with text-like inputs.

> **Key takeaway:**
> A vector carries distributed numerical features.

390px

Tensors: Tensor Block

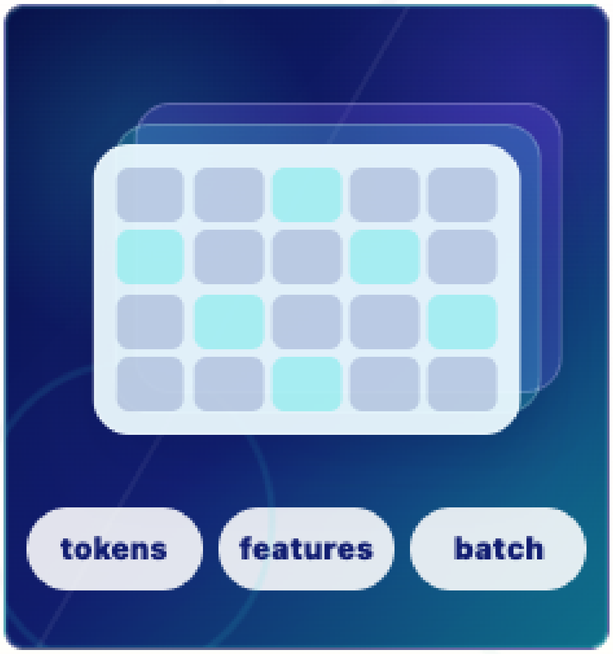

Tensor Block

Rows, columns, and batches

A tensor is a shaped block of numbers, often holding many token vectors at once.

1. **Vector** One list of numbers.
2. **Tensor** Many vectors arranged in a shape.
3. **Shape matters** Rows, columns, and batches organize computation.

> **Key takeaway:**
> Tensors are shaped blocks of numbers.

320px

Tensor Block

Rows, columns, and batches

A tensor is a shaped block of numbers, often holding many token vectors at once.

1. **Vector** One list of numbers.
2. **Tensor** Many vectors arranged in a shape.
3. **Shape matters** Rows, columns, and batches organize computation.

> **Key takeaway:**
> Tensors are shaped blocks of numbers.

390px

Context Window: Temporary Context Window

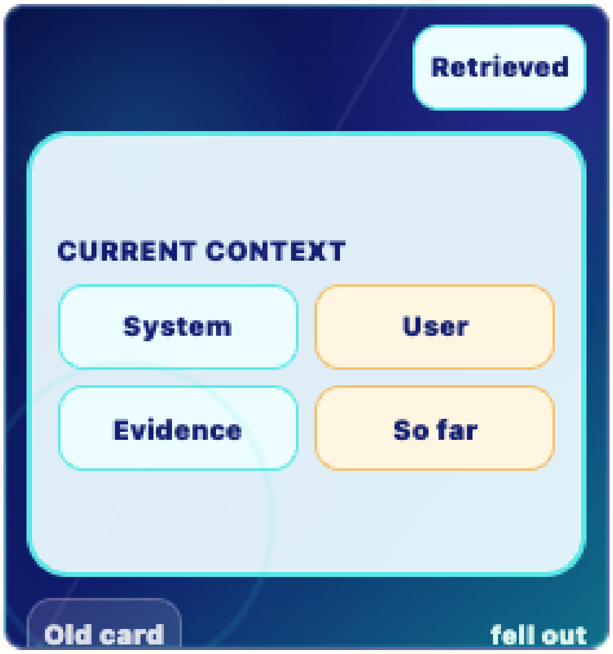

Temporary Context Window

Limited visible input

The context window is the limited visible workspace for the current run.

1. **Visible now** Only current context can shape the next token.
2. **Falls out** Old content can stop being visible.
3. **Retrieved material** Outside evidence helps only after it enters context.

> **Key takeaway:**
> Context is temporary visible input, not durable memory.

320px

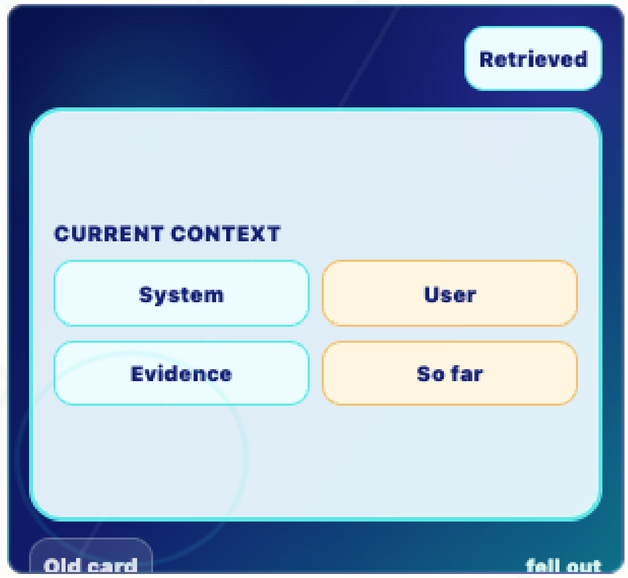

Temporary Context Window

Limited visible input

The context window is the limited visible workspace for the current run.

1. **Visible now** Only current context can shape the next token.
2. **Falls out** Old content can stop being visible.
3. **Retrieved material** Outside evidence helps only after it enters context.

> **Key takeaway:**
> Context is temporary visible input, not durable memory.

390px

Model Literate Synthesis: Model Literacy Board

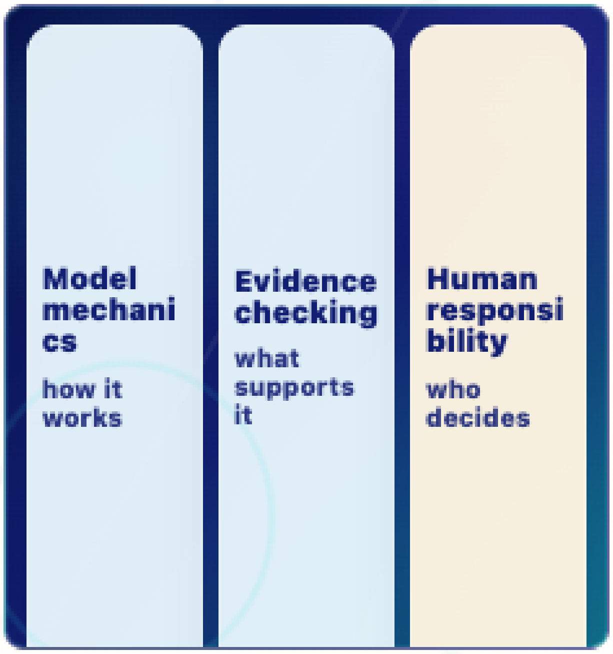

Model Literacy Board

Mechanics plus judgment

Model literacy connects how the system works with how humans should use it.

- **Mechanics** Prompts become tokens, scores, probabilities, and generated tokens.
- **Evidence** Fluent output still needs checking.
- **Responsibility** Humans remain accountable for use, review, and decisions.

> **Key takeaway:**
> Mechanics matter, and humans remain responsible.

320px

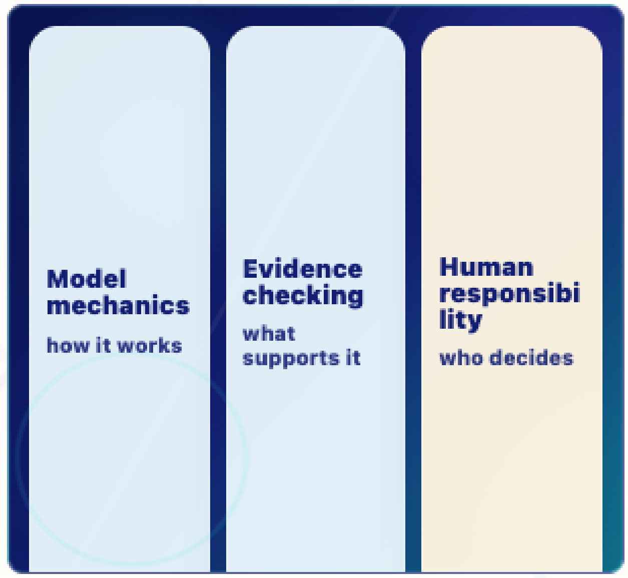

Model Literacy Board

Mechanics plus judgment

Model literacy connects how the system works with how humans should use it.

- **Mechanics** Prompts become tokens, scores, probabilities, and generated tokens.
- **Evidence** Fluent output still needs checking.
- **Responsibility** Humans remain accountable for use, review, and decisions.

> **Key takeaway:**
> Mechanics matter, and humans remain responsible.

390px

Remaining Concerns

The ten target visuals now use hand-specified mobile templates. Remaining acceptable-for-testing concerns are pedagogical rather than overflow related: testers should still report whether the vector and tensor visuals feel too abstract, and whether the final synthesis board is memorable enough after simplification.