

Prompt Life v0.28.4 Word-Wrap Guardrails Report

Generated 2026-06-14T03:02:55.482Z. This packet documents the typography and visual-label repair pass for mobile Journey visuals.

Audit Result

- audit:word-wrap status: pass
- Unsafe wrapping CSS issues: 0
- Known bad split strings: 0
- Long internal labels: 0
- Paragraphs inside diagram bodies: 0

CSS Rules Changed

- Replaced arbitrary overflow wrapping in current app CSS with natural wrapping.
- Added shared guardrail classes: `.visual-label`, `.visual-chip`, `.visual-node-label`, `.visual-callout-title`, and `.visual-callout-body`.
- Visual labels now use `word-break: normal`, `overflow-wrap: normal`, and `hyphens: none`.
- Four-column boundary boards collapse to a 2-by-2 layout at very narrow widths so labels have room to wrap at spaces.

Visual Labels Shortened

LOCATION	BEFORE	AFTER
Prompt vs Response internal panel subtitle	provided before generation	given first
Prompt vs Response internal panel subtitle	what the next run sees	next run sees
Tokenization source panel subtitle	human-readable sentence	sentence
Tokenization stream panel subtitle	model-readable chunks	token chunks
Tokenization uneven-example panel subtitle	tokens are not always words	not always words
Tokenization source panel text	A jealous dog chased a startled cat...	A jealous dog...
Inference mechanism-strip row	Temporary activations	Activations
Fine-Tuning boundary-board label	Retrieve	Find

LOCATION	BEFORE	AFTER
Hidden States boundary-board label	Embedding	Embed
Model Literate Synthesis boundary-board label	Model mechanics	Model
Model Literate Synthesis boundary-board label	Evidence checking	Evidence
Model Literate Synthesis boundary-board label	Human responsibility	Human role
Attention token chips	chased / because / hissed	saw / as / ran

Moved To Callouts

- Temporary activations remains fully explained in the Inference HTML callout.
- Human responsibility remains fully explained in the Model Literate Synthesis caption and callout.
- Full retrieval, embedding, mechanics, and evidence-checking terms remain in captions, callouts, glossary chips, and accessible descriptions.
- The full tokenization sentence stays in the lesson/callouts rather than inside a small diagram panel.
- Prompt/response timing detail stays in the callouts rather than long panel subtitles.

320px And 390px Screenshots

Prompt vs Response

1 GIVEN PROMPT given first

Describe two pets conflict

2 RESPONSE SO FAR generated tokens

A jealous dog

3 NEXT TOKEN newly selected

chased

4 UPDATED CONTEXT next run sees

prompt A jealous dog
chased

Prompt vs Response

Given context versus generated tokens

Prompt tokens are given to the model; response tokens are generated one at a time and appended to the next context.

- 1 **Given prompt** The request "Describe two pets in conflict" is already provided input.
- 2 **Generated response** The response-so-far "A jealous dog" was generated by the model.
- 3 **Next token** The newly selected

1 GIVEN PROMPT given first

Describe two pets conflict

2 RESPONSE SO FAR generated tokens

A jealous dog

3 NEXT TOKEN newly selected

chased

4 UPDATED CONTEXT next run sees

prompt A jealous dog chased

Prompt vs Response

Given context versus generated tokens

Prompt tokens are given to the model; response tokens are generated one at a time and appended to the next context.

- 1 **Given prompt** The request "Describe two pets in conflict" is already provided input.
- 2 **Generated response** The response-so-far "A jealous dog" was generated by the model.
- 3 **Next token** The newly selected token "chased" is appended to the response.
- 4 **Updated context** The next run sees prompt plus response so far plus the new token.

Key takeaway:

Prompt is given; response tokens are generated and appended.

token "chased" is appended to the response.

- 4 Updated context** The next run sees prompt plus response so far plus the new token.

Key takeaway:

Prompt is given; response tokens are generated and appended.

320px

Text to Tokens

1 SOURCE TEXT sentence

A jealous dog...

2 TOKEN STREAM token chunks

A jealous dog chased a
startled cat across the
kitchen floor .

3 UNEVEN EXAMPLES not always words

startled becomes start led
floor. becomes floor .

Text to Tokens

Uneven chunks, punctuation included

Text is split into model-readable chunks before token IDs and embedding lookup.

- 1 **Simplified split** A | jealous | dog | chased | a | startled | cat | across | the | kitchen | floor | .
- 2 **Uneven chunks** Some tokenizers may split words or punctuation, such as start | led or floor | .
- 3 **Teaching note** This app uses a simplified split; real tokenizer output can differ.

1 SOURCE TEXT sentence

A jealous dog...

2 TOKEN STREAM token chunks

A jealous dog chased a startled
cat across the kitchen floor .

3 UNEVEN EXAMPLES not always words

startled becomes start led
floor. becomes floor .

Text to Tokens

Uneven chunks, punctuation included

Text is split into model-readable chunks before token IDs and embedding lookup.

- 1 **Simplified split** A | jealous | dog | chased | a | startled | cat | across | the | kitchen | floor | .
- 2 **Uneven chunks** Some tokenizers may split words or punctuation, such as start | led or floor | .
- 3 **Teaching note** This app uses a simplified split; real tokenizer output can differ.

Key takeaway:

Tokens are model-readable chunks, not always human words.

390px

Key takeaway:

Tokens are model-readable
chunks, not always human words.

320px

Fine-Tuning



Durable vs Temporary Steering

Fine-tuning is not prompting

Fine-tuning or adapters can shape future responses. Prompting, retrieval, and sampling shape the current run.

- 1 Durable adaptation** Additional training or adapters can carry a pattern into later runs.
- 2 Prompting** A prompt steers the current context.
- 3 Retrieval** Outside material helps only after it enters context.
- 4 Sampling** Sampling chooses one next token; it does not train the model.

Key takeaway:

Fine-tuning can shape future behavior; prompting, retrieval, and sampling shape this run.



Durable vs Temporary Steering

Fine-tuning is not prompting

Fine-tuning or adapters can shape future responses. Prompting, retrieval, and sampling shape the current run.

- 1 Durable adaptation** Additional training or adapters can carry a pattern into later runs.
- 2 Prompting** A prompt steers the current context.
- 3 Retrieval** Outside material helps only after it enters context.
- 4 Sampling** Sampling chooses one next token; it does not train the model.

Key takeaway:

Fine-tuning can shape future behavior; prompting, retrieval, and sampling shape this run.

390px

Inference



Forward Pass

Fixed weights, temporary activations

Inference uses fixed learned weights and current context to produce next-token scores.

- 1 **Context** What the model can currently use.
- 2 **Fixed weights** Learned numbers from training.
- 3 **Temporary activations** Internal values created for this run.
- 4 **Scores** Raw next-token scores.
- 5 **Next token** One token is selected and appended.

Key takeaway:

Inference uses weights without normally updating them.



Forward Pass

Fixed weights, temporary activations

Inference uses fixed learned weights and current context to produce next-token scores.

- 1 **Context** What the model can currently use.
- 2 **Fixed weights** Learned numbers from training.
- 3 **Temporary activations** Internal values created for this run.
- 4 **Scores** Raw next-token scores.
- 5 **Next token** One token is selected and appended.

Key takeaway:

Inference uses weights without normally updating them.

390px

320px

Attention



Relevance Between Tokens

Weighted relevance, not awareness

Attention weights relevance between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Strong clue** "cat" fits the local sentence relationship.
- **Weaker clue** "dog" is nearby context but less useful here.
- **Limit** Attention is calculation over positions, not consciousness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

320px



Relevance Between Tokens

Weighted relevance, not awareness

Attention weights relevance between token positions. It is not awareness.

- **Target token** "it" is the token being updated.
- **Strong clue** "cat" fits the local sentence relationship.
- **Weaker clue** "dog" is nearby context but less useful here.
- **Limit** Attention is calculation over positions, not consciousness.

Key takeaway:

Attention is weighted relevance between token positions, not consciousness.

390px

MLP



Attention vs MLP

Mix positions or reshape features

Attention mixes information across token positions. The MLP reshapes each token position's features.

- **Attention** Asks which positions matter to this token.
- **MLP** Transforms the token's feature vector at that position.
- **Together** Both help update hidden states inside a layer.

Key takeaway:

Attention mixes positions; MLP reshapes features.

320px



Attention vs MLP

Mix positions or reshape features

Attention mixes information across token positions. The MLP reshapes each token position's features.


- **Attention** Asks which positions matter to this token.
- **MLP** Transforms the token's feature vector at that position.
- **Together** Both help update hidden states inside a layer.

Key takeaway:

Attention mixes positions; MLP reshapes features.

390px

Layers



Transformer Layer Strip

Refine, then repeat

A transformer layer refines hidden states; many layers repeat the same kind of update.

- **Hidden states** Temporary internal vectors.
- **Attention** Mixes relevant token-position information.
- **MLP** Reshapes features.
- **Repeat** Later layers refine the representation again.

Key takeaway:

Layers refine hidden states during the run; they do not store permanent memory.

320px



Transformer Layer Strip

Refine, then repeat

A transformer layer refines hidden states; many layers repeat the same kind of update.


- **Hidden states** Temporary internal vectors.
- **Attention** Mixes relevant token-position information.
- **MLP** Reshapes features.
- **Repeat** Later layers refine the representation again.

Key takeaway:

Layers refine hidden states during the run; they do not store permanent memory.

390px

Hidden States



Embedding, Hidden State, Weight


Temporary is not durable

A hidden state is temporary and context-shaped. It is not the same as an embedding or a durable weight.

- **Embedding** The learned starting vector for a token ID.
- **Hidden state** The token's temporary representation after context processing.
- **Weight** Learned model parameter reused across runs.

Key takeaway:

Hidden states are temporary internal vectors shaped by current context.



Embedding, Hidden State, Weight

Temporary is not durable

A hidden state is temporary and context-shaped. It is not the same as an embedding or a durable weight.

- **Embedding** The learned starting vector for a token ID.
- **Hidden state** The token's temporary representation after context processing.
- **Weight** Learned model parameter reused across runs.


Key takeaway:

Hidden states are temporary internal vectors shaped by current context.

390px

320px

Vectors



Feature Vector

Many numbers at once


A vector represents something with many numbers at once.

- 1 **Not one slot** Meaning is not stored in one neat dimension.
- 2 **Distributed features** Useful patterns are spread across many values.
- 3 **Model-readable** Vectors let the model compute with text-like inputs.

Key takeaway:

A vector carries distributed numerical features.

320px



Feature Vector

Many numbers at once

A vector represents something with many numbers at once.


- 1 **Not one slot** Meaning is not stored in one neat dimension.
- 2 **Distributed features** Useful patterns are spread across many values.
- 3 **Model-readable** Vectors let the model compute with text-like inputs.

Key takeaway:

A vector carries distributed numerical features.

390px

Tensors



Tensor Block

Rows, columns, and batches

A tensor is a shaped block of numbers, often holding many token vectors at once.

- 1 **Vector** One list of numbers.
- 2 **Tensor** Many vectors arranged in a shape.
- 3 **Shape matters** Rows, columns, and batches organize computation.

Key takeaway:

Tensors are shaped blocks of numbers.

320px



Tensor Block

Rows, columns, and batches

A tensor is a shaped block of numbers, often holding many token vectors at once.


- 1 **Vector** One list of numbers.
- 2 **Tensor** Many vectors arranged in a shape.
- 3 **Shape matters** Rows, columns, and batches organize computation.

Key takeaway:

Tensors are shaped blocks of numbers.

390px

Context Window



Temporary Context Window

Limited visible input


The context window is the limited visible workspace for the current run.

- 1 **Visible now** Only current context can shape the next token.
- 2 **Falls out** Old content can stop being visible.
- 3 **Retrieved material** Outside evidence helps only after it enters context.

Key takeaway:

Context is temporary visible input, not durable memory.

320px



Temporary Context Window

Limited visible input

The context window is the limited visible workspace for the current run.


- 1 **Visible now** Only current context can shape the next token.
- 2 **Falls out** Old content can stop being visible.
- 3 **Retrieved material** Outside evidence helps only after it enters context.

Key takeaway:

Context is temporary visible input, not durable memory.

390px

Model Literate Synthesis



Model Literacy Board

Mechanics plus judgment

Model literacy connects how the system works with how humans should use it.

- **Mechanics** Prompts become tokens, scores, probabilities, and generated tokens.
- **Evidence** Fluent output still needs checking.
- **Responsibility** Humans remain accountable for use, review, and decisions.

Key takeaway:

Mechanics matter, and humans remain responsible.

320px



Model Literacy Board

Mechanics plus judgment

Model literacy connects how the system works with how humans should use it.

- **Mechanics** Prompts become tokens, scores, probabilities, and generated tokens.
- **Evidence** Fluent output still needs checking.
- **Responsibility** Humans remain accountable for use, review, and decisions.

Key takeaway:

Mechanics matter, and humans remain responsible.

390px