

Prompt Life v0.28.7b Image 2 Source Processing

Generated 2026-06-15T02:46:15.034Z. This pass processed five newly supplied Image 2 source files, created optimized and thumbnail derivatives, and made verified assets live through the Journey visual manifest. Missing entries continue to render their existing visuals in normal learner mode.

12
MANIFEST ENTRIES

6
LIVE ASSETS

5
NEWLY LIVE

6
STILL MISSING

0
REJECTED

SUMMARY

Five new assets were accepted

The new live assets are **Costs We Must Count**, **Risk vs Myth**, **Model Literate Synthesis**, **The Perfect Storm**, **Human-Centered AI**. Collective Intelligence remains live from v0.28.7. No assets were rejected, and no normal learner-facing implementation placeholder appears in this pass.

SOURCE MAPPING

Source Files, Optimized Files, and Thumbnails

CARD	SOURCE	OPTIMIZED	THUMBNAIL	STATUS	REVIEW NOTES
Costs We Must Count costs-we-must-count	costs-we-must-count.png 1448x1086 2204 KB	costs-we-must-count.png 1448x1086 2383 KB	costs-we-must-count.png 480x360 280 KB	live	Accepted live in v0.28.7b. The image is textless, uses no fake statistics, and makes infrastructure, labor, privacy, and governance visible without doom framing.
Risk vs Myth risk-myth	risk-vs-myth.png 1672x941 1983 KB	risk-vs-myth.png 1448x815 1631 KB	risk-vs-myth.png 480x270 206 KB	live	Accepted live in v0.28.7b. The split path is readable, non-dashboard, and separates practical safeguards from pale myth imagery without horror or model-consciousness framing.
Model Literate Synthesis model-literate-synthesis	model-literate-synthesis.png 1672x941 2249 KB	model-literate-synthesis.png 1448x815 1771 KB	model-literate-synthesis.png 480x270 212 KB	live	Accepted live in v0.28.7b. The image is textless, non-certifying, and visually connects mechanics, evidence review, and human responsibility.
The Perfect Storm perfect-storm	perfect-storm.png 1672x941 2032 KB	perfect-storm.png 1448x815 1692 KB	perfect-storm.png 480x270 223 KB	live	Accepted live in v0.28.7b. The textless convergence scene is horizontal, uncluttered, and supports the one-spark misconception repair.
Human-Centered AI human-centered-ai	human-centered-ai.png 1672x941 1876 KB	human-centered-ai.png 1448x815 1599 KB	human-centered-ai.png 480x270 217 KB	live	Accepted live in v0.28.7b. The image keeps a human figure central, avoids robot-judge imagery, and uses HTML for all instructional language.

MANIFEST

All Image 2 Candidate Statuses

STAGE	CARD	STATUS	SIZE	TEXT CHECK	NOTES	AUDIT
1	What Is an LLM? what-is-llm	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
1	Pretraining pretraining	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
1	Alignment alignment	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
6	Multimodal AI multimodal	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
6	The Perfect Storm perfect-storm	live	1448x815 1692 KB	verified	Accepted live in v0.28.7b. The textless convergence scene is horizontal, uncluttered, and supports the one-spark misconception repair.	ok
7	Collective Intelligence, Extracted collective-intelligence	live	1448x1086 2192 KB	verified	Accepted live in v0.28.7. Text check: no readable embedded text; abstract source traces are symbolic texture. Uses HTML caption and callouts for all instructional language.	ok
7	Costs We Must Count costs-we-must-count	live	1448x1086 2383 KB	verified	Accepted live in v0.28.7b. The image is textless, uses no fake statistics, and makes infrastructure, labor, privacy, and governance visible without doom framing.	ok
7	Risk vs Myth risk-myth	live	1448x815 1631 KB	verified	Accepted live in v0.28.7b. The split path is readable, non-dashboard, and separates practical safeguards from pale myth imagery without horror or model-consciousness framing.	ok
8	Benefits Worth Taking Seriously benefits-worth-taking-seriously	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
8	Human-Centered AI human-centered-ai	live	1448x815 1599 KB	verified	Accepted live in v0.28.7b. The image keeps a human figure central, avoids robot-judge imagery, and uses HTML for all instructional language.	ok
8	Better AI is a Choice better-ai-choice	missing	n/a n/a	not live	missing verified optimized asset and thumbnail after v0.28.7b source processing	ok
8	Model Literate Synthesis model-literate-synthesis	live	1448x815 1771 KB	verified	Accepted live in v0.28.7b. The image is textless, non-certifying, and visually connects mechanics, evidence review, and human responsibility.	ok


CAPTION REVIEW

HTML Caption and Callout Copy

CARD	CAPTION	CALLOUTS
Costs We Must Count	AI answers depend on infrastructure, energy and water, labor, privacy, governance, and institutional choices.	Infrastructure: Data centers, chips, energy, water, and e-waste are part of the real system. Human systems: Labor, privacy, bias, and information quality depend on deployment choices. Institutions: Governance, incentives, and access shape who benefits and who bears costs.
Risk vs Myth	Model-literate users separate practical risks from unsupported myths, then choose review and safeguards.	Practical risks: Privacy exposure, hallucination, tool misuse, bias, and overreliance need concrete safeguards. Unsupported myths: Conscious chatbots, hidden omniscience, and secret self-training claims need evidence. Review: Risk literacy asks what system, context, tool, and institution are actually involved.
Model Literate Synthesis	Model literacy connects mechanics, evidence checking, and human responsibility without treating fluency as wisdom.	Mechanics: Prompts become tokens, vectors, states, scores, probabilities, and generated tokens. Evidence: RAG, grounding, citations, and sources help only when the support actually matches. Responsibility: People remain accountable for use, review, governance, and decisions.
The Perfect Storm	Data, compute, storage, methods, labor, and incentives converged into modern LLM capability.	Many inputs: Data, compute, storage, and methods each mattered. Human systems: Labor, funding, products, and demand helped move capability into use. Convergence: Modern LLMs came from conditions meeting, not one single breakthrough.
Human-Centered AI	AI output should support human judgment; people remain accountable for decisions, review, and governance.	Support: AI output can help people draft, notice, compare, or review. Judgment: High-stakes decisions still need human responsibility and local context. Accountability: Review, appeal, policy, and governance belong to human institutions.

VISUAL AID

What to picture



Costs We Must Count

The answer is not weightless

AI answers depend on infrastructure, energy and water, labor, privacy, governance, and institutional choices.

- **Infrastructure** Data centers, chips, energy, water, and e-waste are part of the real system.
- **Human systems** Labor, privacy, bias, and information quality depend on deployment choices.
- **Institutions** Governance, incentives, and access shape who benefits and who bears costs.

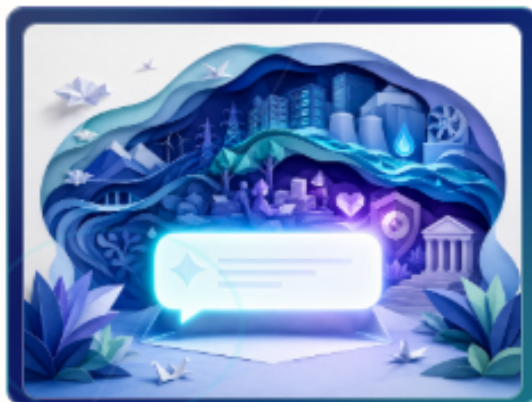
Key takeaway:

Costs vary, but they are real enough to count.

Costs We Must Count: Image 2 visual at 390px.

VISUAL AID

What to picture



Costs We Must Count

The answer is not weightless

AI answers depend on infrastructure, energy and water, labor, privacy, governance, and institutional choices.

- **Infrastructure** Data centers, chips, energy, water, and e-waste are part of the real system.
- **Human systems** Labor, privacy, bias, and information quality depend on deployment choices.
- **Institutions** Governance, incentives, and access shape who benefits and who bears costs.


Key takeaway:

Costs vary, but they are real enough to count.

Costs We Must Count: Image 2 visual at 320px.

VISUAL AID

What to picture



Costs We Must Count

The answer is not weightless

AI answers depend on infrastructure, energy and water, labor, privacy, governance, and institutional choices.

- **Infrastructure** Data centers, chips, energy, water, and e-waste are part of the real system.
- **Human systems** Labor, privacy, bias, and information quality depend on deployment choices.
- **Institutions** Governance, incentives, and access shape who benefits and who bears costs.

Key takeaway:

Costs vary, but they are real enough to count.

Costs We Must Count: Image 2 visual on desktop.

MIDNIGHT LEDGER

33 / 39

1 Before Morning	2 Morning Commute	3 Workday	4 Decision Room
5 Day Repeats	6 Twilight	7 Midnight Ledger	8 New Dawn

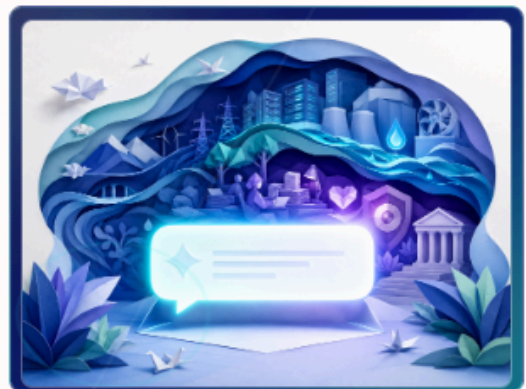
Costs We Must Count

The answer feels weightless. The infrastructure is not.

Generative AI has physical, social, cultural, and ethical costs that must be counted honestly.

VISUAL AID

What to picture



Costs We Must Count


The answer is not weightless

- Home
- Journey**
- Play
- Glossary
- Badge

Costs We Must Count: Journey card smoke at 390px.

VISUAL AID

What to picture



Risk vs Myth

Mechanisms, not myth

Model-literate users separate practical risks from unsupported myths, then choose review and safeguards.

- **Practical risks** Privacy exposure, hallucination, tool misuse, bias, and overreliance need concrete safeguards.
- **Unsupported myths** Conscious chatbots, hidden omniscience, and secret self-training claims need evidence.
- **Review** Risk literacy asks what system, context, tool, and institution are actually involved.

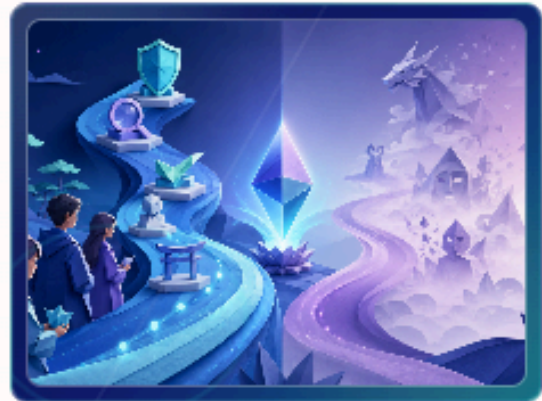
Key takeaway:

Real AI risks usually come from data, context, tools, institutions, and overreliance, not model consciousness or an unexplained shortcut.

Risk vs Myth: Image 2 visual at 390px.

VISUAL AID

What to picture



Risk vs Myth

Mechanisms, not myth

Model-literate users separate practical risks from unsupported myths, then choose review and safeguards.

- **Practical risks** Privacy exposure, hallucination, tool misuse, bias, and overreliance need concrete safeguards.
- **Unsupported myths** Conscious chatbots, hidden omniscience, and secret self-training claims need evidence.
- **Review** Risk literacy asks what system, context, tool, and institution are actually involved.


Key takeaway:

Real AI risks usually come from data, context, tools, institutions, and overreliance, not model consciousness or an unexplained shortcut.

Risk vs Myth: Image 2 visual at 320px.

VISUAL AID

What to picture



Risk vs Myth

Mechanisms, not myth

Model-literate users separate practical risks from unsupported myths, then choose review and safeguards.

- **Practical risks** Privacy exposure, hallucination, tool misuse, bias, and overreliance need concrete safeguards.
- **Unsupported myths** Conscious chatbots, hidden omniscience, and secret self-training claims need evidence.
- **Review** Risk literacy asks what system, context, tool, and institution are actually involved.

Key takeaway:

Real AI risks usually come from data, context, tools, institutions, and overreliance, not model consciousness or an unexplained shortcut.

Risk vs Myth: Image 2 visual on desktop.

MIDNIGHT LEDGER

34 / 39

1 Before Morning	2 Morning Commute	3 Workday	4 Decision Room
5 Day Repeats	6 Twilight	7 Midnight Ledger	8 New Dawn


Risk vs Myth

Practical institutional literacy

Risk literacy means separating practical, mechanism-based AI risks from myths or unsupported claims about what models are.

VISUAL AID

What to picture



Risk vs Myth

Mechanisms, not myth

Model-literate users separate practical risks from unsupported myths, then

Home	Journey	Play	Glossary	Badge
------	---------	------	----------	-------

Risk vs Myth: Journey card smoke at 390px.

VISUAL AID

What to picture



Model Literate Synthesis

Mechanics plus judgment

Model literacy connects mechanics, evidence checking, and human responsibility without treating fluency as wisdom.

- **Mechanics** Prompts become tokens, vectors, states, scores, probabilities, and generated tokens.
- **Evidence** RAG, grounding, citations, and sources help only when the support actually matches.
- **Responsibility** People remain accountable for use, review, governance, and decisions.

Key takeaway:

Mechanics matter, and humans remain responsible.

Model Literate Synthesis: Image 2 visual at 390px.

VISUAL AID

What to picture



Model Literate Synthesis

Mechanics plus judgment

Model literacy connects mechanics, evidence checking, and human responsibility without treating fluency as wisdom.

- **Mechanics** Prompts become tokens, vectors, states, scores, probabilities, and generated tokens.
- **Evidence** RAG, grounding, citations, and sources help only when the support actually matches.
- **Responsibility** People remain accountable for use, review, governance, and decisions.


Key takeaway:

Mechanics matter, and humans remain responsible.

Model Literate Synthesis: Image 2 visual at 320px.

VISUAL AID

What to picture



Model Literate Synthesis

Mechanics plus judgment

Model literacy connects mechanics, evidence checking, and human responsibility without treating fluency as wisdom.

- **Mechanics** Prompts become tokens, vectors, states, scores, probabilities, and generated tokens.
- **Evidence** RAG, grounding, citations, and sources help only when the support actually matches.
- **Responsibility** People remain accountable for use, review, governance, and decisions.

Key takeaway:

Mechanics matter, and humans remain responsible.

Model Literate Synthesis: Image 2 visual on desktop.

NEW DAWN

39 / 39

1 Before Morning	2 Morning Commute	3 Workday	4 Decision Room
5 Day Repeats	6 Twilight	7 Midnight Ledger	8 New Dawn


Model Literate Synthesis

Understand the machine, then decide how to use it

Model literacy means understanding both the mechanics and the human consequences.

VISUAL AID

What to picture



Model Literate Synthesis

Mechanics plus judgment

Home	Journey	Play	Glossary	Badge
------	---------	------	----------	-------

Model Literate Synthesis: Journey card smoke at 390px.

VISUAL AID

What to picture



The Perfect Storm

Why LLMs arrived now

Data, compute, storage, methods, labor, and incentives converged into modern LLM capability.

- **Many inputs** Data, compute, storage, and methods each mattered.
- **Human systems** Labor, funding, products, and demand helped move capability into use.
- **Convergence** Modern LLMs came from conditions meeting, not one single breakthrough.

Key takeaway:

Modern LLMs came from a convergence, not a single spark.

The Perfect Storm: Image 2 visual at 390px.

VISUAL AID

What to picture



The Perfect Storm

Why LLMs arrived now

Data, compute, storage, methods, labor, and incentives converged into modern LLM capability.

- **Many inputs** Data, compute, storage, and methods each mattered.
- **Human systems** Labor, funding, products, and demand helped move capability into use.
- **Convergence** Modern LLMs came from conditions meeting, not one single breakthrough.

Key takeaway:

Modern LLMs came from a convergence, not a single spark.

The Perfect Storm: Image 2 visual at 320px.

VISUAL AID

What to picture



The Perfect Storm

Why LLMs arrived now

Data, compute, storage, methods, labor, and incentives converged into modern LLM capability.

- **Many inputs** Data, compute, storage, and methods each mattered.
- **Human systems** Labor, funding, products, and demand helped move capability into use.
- **Convergence** Modern LLMs came from conditions meeting, not one single breakthrough.

Key takeaway:

Modern LLMs came from a convergence, not a single spark.

The Perfect Storm: Image 2 visual on desktop.

TWILIGHT: THE WIDER LANDSCAPE

31 / 39

1 Before Morning	2 Morning Commute	3 Workday	4 Decision Room
5 Day Repeats	6 Twilight	7 Midnight Ledger	8 New Dawn

The Perfect Storm

Why LLMs arrived now

Modern LLMs became possible when human-created digital data, powerful compute and storage, deep-learning methods, human labor, and economic incentives converged.

VISUAL AID

What to picture



The Perfect Storm

Why LLMs arrived now


Data, compute, storage, methods, labor, and incentives converged into modern LLM capability.

- Home
- Journey
- Play
- Glossary
- Badge

The Perfect Storm: Journey card smoke at 390px.

VISUAL AID

What to picture



Human-Centered AI

Tools should serve dignity

AI output should support human judgment; people remain accountable for decisions, review, and governance.

- **Support** AI output can help people draft, notice, compare, or review.
- **Judgment** High-stakes decisions still need human responsibility and local context.
- **Accountability** Review, appeal, policy, and governance belong to human institutions.

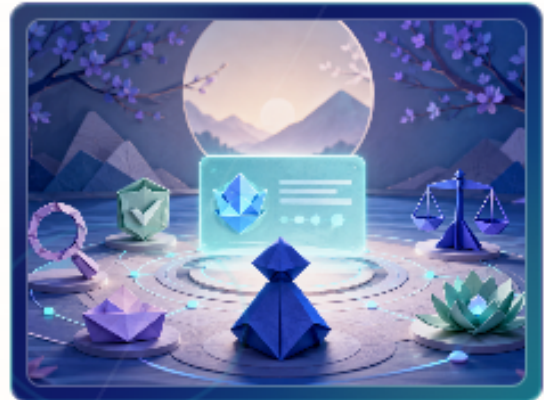
Key takeaway:

Powerful tools still need human purpose.

Human-Centered AI: Image 2 visual at 390px.

VISUAL AID

What to picture



Human-Centered AI

Tools should serve dignity

AI output should support human judgment; people remain accountable for decisions, review, and governance.

- **Support** AI output can help people draft, notice, compare, or review.
- **Judgment** High-stakes decisions still need human responsibility and local context.
- **Accountability** Review, appeal, policy, and governance belong to human institutions.


Key takeaway:

Powerful tools still need human purpose.

Human-Centered AI: Image 2 visual at 320px.

VISUAL AID

What to picture



Human-Centered AI

Tools should serve dignity

AI output should support human judgment; people remain accountable for decisions, review, and governance.

- **Support** AI output can help people draft, notice, compare, or review.
- **Judgment** High-stakes decisions still need human responsibility and local context.
- **Accountability** Review, appeal, policy, and governance belong to human institutions.

Key takeaway:

Powerful tools still need human purpose.

Human-Centered AI: Image 2 visual on desktop.

NEW DAWN

36 / 39

1 Before Morning	2 Morning Commute	3 Workday	4 Decision Room
5 Day Repeats	6 Twilight	7 Midnight Ledger	8 New Dawn


Human-Centered AI

Tools should serve human dignity

Human-centered AI asks whether a system supports dignity, learning, responsibility, creativity, relationships, and the common good.

VISUAL AID

What to picture



Human-Centered AI

Tools should serve dignity

Home	Journey	Play	Glossary	Badge
------	---------	------	----------	-------

Human-Centered AI: Journey card smoke at 390px.

VISUAL AID

What to picture



Prompt to Prediction

Score, choose, append, repeat

Current context enters the model, learned weights shape next-token probabilities, one token is generated, and the response grows token by token.

- **Context enters** The current prompt and response-so-far become the visible input for this run.
- **Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- **One token is generated** The model chooses one response token from the probability cloud.
- **Response grows** That token is appended, then the model runs again.

Key takeaway:

An LLM turns context into next-token probabilities.

Missing asset smoke: What Is an LLM keeps the current visual.

BADGE

Prompt Life: Model Literate

Under construction while the evidence model is reviewed.



Under construction

Pending human review

Prompt Life: Model Literate is under construction. The Journey, checkpoint bank, and evidence model are being tested before any badge is issued.



Home



Journey



Play



Glossary



Badge

Badge page smoke screenshot.

PLAY LAB

Play to learn

Short, calm challenges that make model mechanics visible.

Progress saved on this device. 0 of 5 practice challenges show saved completion or review progress.

Each challenge practices one model-literacy move.

More below ↓

Play challenges

Suggested order follows the Journey.



Glossary Dojo

Ready

Practice concept discrimination.

Practice move: Name concepts

Recommended: Anytime

Anytime

Glossary

Learning cards: All learning cards

Action: Choose, compare, review.

Time: 4 min



Home



Journey



Play



Glossary



Badge

Play page smoke screenshot.

DEFINITIONS PLUS RELATIONSHIPS

Glossary

Use each term with its neighbor. That is where the mental model gets stable.

PRACTICE MODE

Glossary Dojo

Practice twelve terms through definitions, relationships, and common mix-ups.

Practice 12 terms

Search terms

token, tensor, softmax...

A-Z

Learning path

Find terms alphabetically.

Accessibility

Designing tools and materials so more people can perceive, understand, navigate, and participate.

AI can support accessibility through drafting, summarization, translation, captioning, alternate formats, and assistive workflows when reviewed



Home



Journey



Play



Glossary



Badge

Glossary page smoke screenshot.

Command Results

COMMAND	STATUS	NOTES
npm run process:visual-assets	pass	> promptlife@0.28.7-beta.1 process:visual-assets > node scripts/process-journey-visual-assets.mjs Processed costs-we-must-count.png source: public/assets/journey-visuals/v0-28/source/costs-we-must-count.png optimized: public/assets/journey-visuals/v0-28/optimized/costs-we-must-count.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/costs-we-must-count.png Processed risk-vs-myth.png source: public/assets/journey-visuals/v0-28/source/risk-vs-myth.png optimized: public/assets/journey-visuals/v0-28/optimized/risk-vs-myth.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/risk-vs-myth.png Processed model-literate-synthesis.png source: public/assets/journey-visuals/v0-28/source/model-literate-synthesis.png optimized: public/assets/journey-visuals/v0-28/optimized/model-literate-synthesis.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/model-literate-synthesis.png Processed perfect-storm.png source: public/assets/journey-visuals/v0-28/source/perfect-storm.png optimized: public/assets/journey-visuals/v0-28/optimized/perfect-storm.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/perfect-storm.png Processed human-centered-ai.png source: public/assets/journey-visuals/v0-28/source/human-centered-ai.png optimized: public/assets/journey-visuals/v0-28/optimized/human-centered-ai.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/human-centered-ai.png
npm run typecheck	pass	> promptlife@0.28.7-beta.1 typecheck > tsc --noEmit
npm run build	pass	> promptlife@0.28.7-beta.1 build > vite build vite v8.0.14 building client environment for production... [2K transforming...✓ 39 modules transformed. rendering chunks... computing gzip size... dist/index.html 1.03 kB gzip: 0.47 kB dist/assets/index-DMxoEmxG.css 151.39 kB gzip: 29.10 kB dist/assets/index-i5nfDJsG.js 1,078.37 kB gzip: 262.25 kB ✓ built in 342ms [plugin builtin:vite-reporter] (!) Some chunks are larger than 500 kB after minification. Consider: - Using dynamic import() to code-split the application - Use build.rolldownOptions.output.codeSplitting to improve chunking: https://rolldown.rs/reference/OutputOptions.codeSplitting - Adjust chunk size limit for this warning via build.chunkSizeWarningLimit.
npm run build:pages	pass	> promptlife@0.28.7-beta.1 build:pages > GITHUB_PAGES=true vite build vite v8.0.14 building client environment for production... [2K transforming...✓ 39 modules transformed. rendering chunks... computing gzip size... dist/index.html 1.05 kB gzip: 0.47 kB dist/assets/index-DMxoEmxG.css 151.39 kB gzip: 29.10 kB dist/assets/index-B20F0Nuj.js 1,078.40 kB gzip: 262.26 kB ✓ built in 177ms [plugin builtin:vite-reporter] (!) Some chunks are larger than 500 kB after minification. Consider: - Using dynamic import() to code-split the application - Use build.rolldownOptions.output.codeSplitting to improve chunking: https://rolldown.rs/reference/OutputOptions.codeSplitting - Adjust chunk size limit for this warning via build.chunkSizeWarningLimit.
npm run audit:visual-assets	pass	> promptlife@0.28.7-beta.1 audit:visual-assets > node scripts/audit-visual-assets.mjs Visual asset audit passed. Checked 12 manifest entries; live assets before: 1; live assets after: 6; newly live: 5; missing assets: 6; rejected assets: 0; pending a:
npm run audit:visual-aids	pass	> promptlife@0.28.7-beta.1 audit:visual-aids > node scripts/audit-visual-aids.mjs && node scripts/audit-word-wrap.mjs Visual aid audit passed: 39 aids, 0 mismatches, 0 readiness issues. Word-wrap audit pass. Checked 80 files. Unsafe CSS: 0 Known bad splits: 0 Long internal labels: 0 Diagram paragraphs: 0
npm run audit:visual-overflow	pass	> promptlife@0.28.7-beta.1 audit:visual-overflow > node scripts/audit-visual-overflow.mjs Visual overflow audit passed. Checked 39 learner-route visual aids at 320px and 390px.
npm run audit:word-wrap	pass	> promptlife@0.28.7-beta.1 audit:word-wrap > node scripts/audit-word-wrap.mjs Word-wrap audit pass. Checked 80 files. Unsafe CSS: 0 Known bad splits: 0 Long internal labels: 0 Diagram paragraphs: 0
npm run audit:learner-copy	pass	> promptlife@0.28.7-beta.1 audit:learner-copy > node scripts/audit-learner-copy.mjs Learner-copy audit passed. Checked 19 learner-facing source files.
npm run audit:language	pass	> promptlife@0.28.7-beta.1 audit:language > node scripts/audit-language.mjs Language audit passed. Checked 61 current source/docs files for banned wording.

COMMAND	STATUS	NOTES
npm run audit:answers	pass	er identity, not A/B/C/D position. lesson:grounding Grounding Journey checkpoint 4 1 yes 4 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:hallucinations Hallucinations Journey checkpoint 4 1 yes 4 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:how-ai-learns How AI Learns Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:diffusion Diffusion vs Autoregression Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:multimodal Multimodal AI Journey checkpoint 4 1 yes 3 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:perfect-storm The Perfect Storm Journey checkpoint 4 1 yes 3 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:collective-intelligence Collective Intelligence, Extracted Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:costs-we-must-count Costs We Must Count Journey checkpoint 4 1 yes 1 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:risk-myth Risk vs Myth Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:benefits-worth-taking-seriously Benefits Worth Taking Seriously Journey checkpoint 4 1 yes 1 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:human-centered-ai Human-Centered AI Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:better-ai-choice Better AI Is a Choice Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:effective-prompting-literacy Effective Prompting from Model Literacy Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. lesson:model-literate-synthesis Model Literate Synthesis Journey checkpoint 4 1 yes 2 Randomized. Correctness follows answer identity, not A/B/C/D position. exercise:prompt-or-response-label Prompt or Response? label-tokens 2 n/a no n/a Fixed. Token sequence/context order is meaningful. exercise:pick-next-token Pick the Next Token next-token-pick 6 1, 2, 3 yes 1, 4, 5 Randomized. Correctness follows item id. exercise:durable-or-temporary Durable or Temporary? sort-buckets 9 n/a no n/a Fixed. Learner sorts into buckets. exercise:traditions-ingredients Rules, Patterns, or Hybrid? sort-buckets 6 n/a no n/a Fixed. Learner sorts into buckets. exercise:training-loop-durable-step Training Loop: Durable Step toggle-state 4 4 no n/a Fixed. Not a single-choice answer list. exercise:pretraining-patterns-not-recall Pattern Learning, Not Perfect Recall toggle-state 2 2 no n/a Fixed. Not a single-choice answer list. exercise:overfitting-held-out-check Validation Check toggle-state 2 1 no n/a Fixed. Not a single-choice answer list. exercise:fine-tuning-context-decoding-sort Open Run or Durable Change? sort-buckets 5 n/a no n/a Fixed. Learner sorts into buckets. exercise:alignment-method-groups Alignment Method Groups sort-buckets 8 n/a no n/a Fixed. Learner sorts into buckets. exercise:attention-relevance Attention Is Relevance connect-nodes 10 7 no n/a Fixed. Not a single-choice answer list. exercise:context-window-fell-out Context Window: What Fell Out? tap-choice 4 1, 2 yes 2, 3 Randomized. Correctness follows item id. exercise:training-nudge Training Nudge toggle-state 2 1 no n/a Fixed. Not a single-choice answer list. exercise:softmax-funnel Softmax Funnel tap-choice 4 1 yes 4 Randomized. Correctness follows item id. exercise:mlp-feature-reshape MLP Feature Reshape toggle-state 2 2 no n/a Fixed. Not a single-choice answer list. exercise:hidden-state-sort Hidden State Sort sort-buckets 6 n/a no n/a Fixed. Learner sorts into buckets. exercise:open-book-or-learned Open Book or Learned? sort-buckets 8 n/a no n/a Fixed. Learner sorts into buckets. exercise:brain-metaphor-boundary Brain Metaphor Boundary tap-multiple 7 n/a no n/a Fixed. Multi-select grouping is clearer in authored order. prompt-run:prompt-run-prompt-response Prompt or Response? label-tokens 1 n/a no n/a Fixed. Token sequence/context order is meaningful. prompt-run:prompt-run-tokenizer Tokenizer label-tokens 1 n/a no n/a Fixed. Token sequence/context order is meaningful. prompt-run:prompt-run-token-ids Token IDs drag-match 7 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-embedding Embedding Lookup drag-match 7 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-tensor Tensor Stack drag-order 3 n/a no n/a Fixed. Order is the learning objective. prompt-run:prompt-run-attention Attention connect-nodes 10 7 no n/a Fixed. Not a single-choice answer list. prompt-run:prompt-run-mlp MLP Feature Reshape drag-match 5 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-hidden-state Hidden State tap-choice 3 1 yes 2 Randomized. Correctness follows item id. prompt-run:prompt-run-logits Logits tap-choice 4 1 yes 4 Randomized. Correctness follows item id. prompt-run:prompt-run-softmax Softmax Funnel tap-choice 4 1 yes 3 Randomized. Correctness follows item id. prompt-run:prompt-run-sampling Sampling next-token-pick 6 1, 2, 3 yes 3, 4, 5 Randomized. Correctness follows item id. prompt-run:prompt-run-append-repeat Append and Repeat drag-order 4 n/a no n/a Fixed. Order is the learning objective. prompt-run:prompt-run-final-order Put the Run in Order drag-order 12 n/a no n/a Fixed. Order is the learning objective. ## Checkpoint Bank Audit Live Journey cards: 39 Live checkpoint questions: 39 Draft cards: 39 Draft questions: 136 Draft choices: 544 Draft wrong-answer distractors: 408 Draft question-count distribution: {"2":5,"3":12,"4":20,"5":2} Pilot cards: 6 Pilot questions: 18 Pilot choices: 72 Pilot wrong-answer distractors: 54 Pilot model-thinking/application/mechanism questions: 100% Pilot direct-definition questions: 0% Revised active-development cards: 6 Revised active-development questions: 19 Revised active-development choices: 76 Revised active-development wrong-answer distractors: 57 Revised audit-seed correct-answer positions: [4,3,1,4,2,3,1,2,4,1,4,1,1,3,4,3,3,2,2] v0.27.10 active-development cards: 12 v0.27.10 active-development questions: 39 v0.27.10 active-development choices: 156 v0.27.10 active-development wrong-answer distractors: 117 v0.27.10 audit-seed correct-answer positions: [4,1,3,4,2,2,4,3,4,3,3,3,4,3,4,2,4,2,1,2,1,2,4,1,1,4,1,1,4,3,4,4,1,1,1,3,3,1] v0.27.15 active-development cards: 39 v0.27.15 active-development questions: 136 v0.27.15 active-development choices: 544 v0.27.15 active-development wrong-answer distractors: 408 v0.27.15 model-thinking/application/mechanism questions: 136 v0.27.15 direct-definition questions: 0 v0.27.15 audit-seed correct-answer positions: [4,1,2,2,1,4,4,1,2,2,4,3,1,1,3,4,2,2,1,3,2,4,1,3,3,2,4,3,3,4,3,2,1,4,1,2,2,2,1,4,1,1,4,2,4,1,1,3,4,2,1,3,2,3,3,4,4,3,3,4,2,1,3,1,1,2,4,2,2,1,2,1,4,4,3,2,4,3,1,4,3, Checkpoint audit passed.
npm run audit:checkpoints	pass	> promptliffe@0.28.7-beta.1 audit:checkpoints > node scripts/audit-checkpoints.mjs ## Checkpoint Bank Audit Live Journey cards: 39 Live checkpoint questions: 39 Draft cards: 39 Draft questions: 136 Draft choices: 544 Draft wrong-answer distractors: 408 Draft question-count distribution: {"2":5,"3":12,"4":20,"5":2} Pilot cards: 6 Pilot questions: 18 Pilot choices: 72 Pilot wrong-answer distractors: 54 Pilot model-thinking/application/mechanism questions: 100% Pilot direct-definition questions: 0% Revised active-development cards: 6 Revised active-development questions: 19 Revised active-development choices: 76 Revised active-development wrong-answer distractors: 57 Revised audit-seed correct-answer positions: [4,3,1,4,2,3,1,2,4,1,4,1,1,3,4,3,3,2,2] v0.27.10 active-development cards: 12 v0.27.10 active-development questions: 39 v0.27.10 active-development choices: 156 v0.27.10 active-development wrong-answer distractors: 117 v0.27.10 audit-seed correct-answer positions: [4,1,3,4,2,2,4,3,4,3,3,3,4,3,4,2,4,2,1,2,1,2,4,1,1,4,1,1,4,3,4,4,1,1,1,3,3,1] v0.27.15 active-development cards: 39 v0.27.15 active-development questions: 136 v0.27.15 active-development choices: 544 v0.27.15 active-development wrong-answer distractors: 408 v0.27.15 model-thinking/application/mechanism questions: 136 v0.27.15 direct-definition questions: 0 v0.27.15 audit-seed correct-answer positions: [4,1,2,2,1,4,4,1,2,2,4,3,1,1,3,4,2,2,1,3,2,4,1,3,3,2,4,3,3,4,3,2,1,4,1,2,2,2,1,4,1,1,4,2,4,1,1,3,4,2,1,3,2,3,3,4,4,3,3,4,2,1,3,1,1,2,4,2,2,1,2,1,4,4,3,2,4,3,1,4,3, Checkpoint audit passed.

COMMAND	STATUS	NOTES
npm run audit:question-clues	pass	> promptlife@0.28.7-beta.1 audit:question-clues > node scripts/audit-question-clues.mjs Question-clue audit completed. Term-echo candidates: 36 Answer-length warnings: 43 Vague-stem warnings: 0 Wrote docs/journey/checkpoints/question-clue-audit-v0-27-15.md
npm run audit:exercises	pass	> promptlife@0.28.7-beta.1 audit:exercises > node scripts/audit-exercises.mjs Exercise audit passed. Inventoried 39 Journey interactions and exported review PDFs.

BACKLOG

Remaining Image 2 Backlog

Still missing after this pass: **What Is an LLM?, Pretraining, Alignment, Multimodal AI, Benefits Worth Taking Seriously, Better AI Is a Choice.**