

Prompt Life v0.28.9 Image 2 Source Processing

Generated 2026-06-15T15:25:33.458Z. This pass processed the remaining six source files, created optimized and thumbnail derivatives, and made verified visuals live through the Journey visual manifest. Normal learner mode now has verified concept images for all twelve manifest entries.

12
MANIFEST ENTRIES

12
LIVE ASSETS

6
NEWLY LIVE

0
STILL MISSING

0
REJECTED

SUMMARY

Six new assets were accepted

The new live visuals are **What Is an LLM?**, **Pretraining**, **Alignment**, **Multimodal AI**, **Benefits Worth Taking Seriously**, **Better AI Is a Choice**. The six already-live visuals remain live. No assets were rejected, none remain missing, and no normal learner-facing implementation placeholder appears in this pass.

SOURCE MAPPING

Source Files, Optimized Files, and Thumbnails

CARD	SOURCE	OPTIMIZED	THUMBNAIL	STATUS	REVIEW NOTES
What Is an LLM? what-is-llm	what-is-an-llm.png 1672x941 1843 KB	what-is-an-llm.png 1448x815 1542 KB	what-is-an-llm.png 480x270 189 KB	live	Accepted live in v0.28.9. The textless prompt-to-model-to-token scene supports current context plus learned weights without implying consciousness, hidden search, or permanent learning.
Pretraining pretraining	pretraining.png 1672x941 2419 KB	pretraining.png 1448x815 1929 KB	pretraining.png 480x270 239 KB	live	Accepted live in v0.28.9. The broad example-stream scene conveys scale and reusable patterns without presenting sources as a searchable archive.
Alignment alignment	alignment.png 1672x941 1987 KB	alignment.png 1448x815 1621 KB	alignment.png 480x270 202 KB	live	Accepted live in v0.28.9. The path-and-guardrail scene supports behavior shaping without implying moral agency or a conscious model.
Multimodal AI multimodal	multimodal-ai.png 1672x941 1839 KB	multimodal-ai.png 1448x815 1552 KB	multimodal-ai.png 480x270 203 KB	live	Accepted live in v0.28.9. The text/image/audio/video streams are clear and avoid eye, brain, surveillance, or human-sensation imagery.
Benefits Worth Taking Seriously benefits-worth-taking-seriously	benefits-worth-taking-seriously.png 1672x941 2534 KB	benefits-worth-taking-seriously.png 1448x815 1973 KB	benefits-worth-taking-seriously.png 480x270 237 KB	live	Accepted live in v0.28.9. The guarded tool-garden scene shows practical benefit categories with a non-identifiable reviewer and avoids utopian claims.
Better AI Is a Choice better-ai-choice	better-ai-choice.png 1672x941 2200 KB	better-ai-choice.png 1448x815 1797 KB	better-ai-choice.png 480x270 229 KB	live	Accepted live in v0.28.9. The branching-path scene supports responsible choice and avoids inevitability or fake control-panel imagery.

MANIFEST

All Image 2 Candidate Statuses

STAGE	CARD	STATUS	SIZE	TEXT CHECK	NOTES	AUDIT
1	What Is an LLM? what-is-llm	live	1448x815 1542 KB	verified	Accepted live in v0.28.9. The textless prompt-to-model-to-token scene supports current context plus learned weights without implying consciousness, hidden search, or permanent learning.	ok
1	Pretraining pretraining	live	1448x815 1929 KB	verified	Accepted live in v0.28.9. The broad example-stream scene conveys scale and reusable patterns without presenting sources as a searchable archive.	ok
1	Alignment alignment	live	1448x815 1621 KB	verified	Accepted live in v0.28.9. The path-and-guardrail scene supports behavior shaping without implying moral agency or a conscious model.	ok
6	Multimodal AI multimodal	live	1448x815 1552 KB	verified	Accepted live in v0.28.9. The text/image/audio/video streams are clear and avoid eye, brain, surveillance, or human-sensation imagery.	ok
6	The Perfect Storm perfect-storm	live	1448x815 1692 KB	verified	Accepted live in v0.28.7b. The textless convergence scene is horizontal, uncluttered, and supports the one-spark misconception repair.	ok
7	Collective Intelligence, Extracted collective-intelligence	live	1448x1086 2192 KB	verified	Accepted live in v0.28.7. Text check: no readable embedded text; abstract source traces are symbolic texture. Uses HTML caption and callouts for all instructional language.	ok
7	Costs We Must Count costs-we-must-count	live	1448x1086 2383 KB	verified	Accepted live in v0.28.7b. The image is textless, uses no fake statistics, and makes infrastructure, labor, privacy, and governance visible without doom framing.	ok
7	Risk vs Myth risk-myth	live	1448x815 1631 KB	verified	Accepted live in v0.28.7b. The split path is readable, non-dashboard, and separates practical safeguards from pale myth imagery without horror or model-consciousness framing.	ok
8	Benefits Worth Taking Seriously benefits-worth-taking-seriously	live	1448x815 1973 KB	verified	Accepted live in v0.28.9. The guarded tool-garden scene shows practical benefit categories with a non-identifiable reviewer and avoids utopian claims.	ok
8	Human-Centered AI human-centered-ai	live	1448x815 1599 KB	verified	Accepted live in v0.28.7b. The image keeps a human figure central, avoids robot-judge imagery, and uses HTML for all instructional language.	ok
8	Better AI Is a Choice better-ai-choice	live	1448x815 1797 KB	verified	Accepted live in v0.28.9. The branching-path scene supports responsible choice and avoids inevitability or fake control-panel imagery.	ok
8	Model Literate Synthesis model-literate-synthesis	live	1448x815 1771 KB	verified	Accepted live in v0.28.7b. The image is textless, non-certifying, and visually connects mechanics, evidence review, and human responsibility.	ok

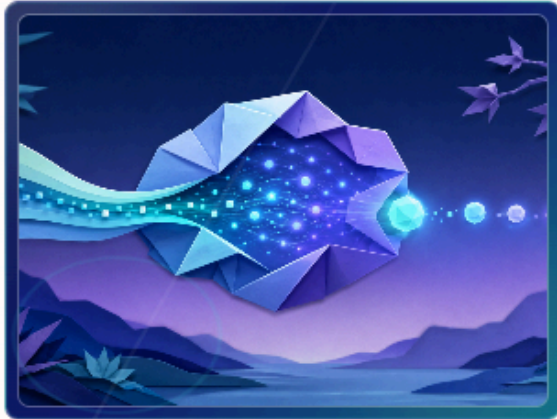
CAPTION REVIEW

HTML Caption and Callout Copy

CARD	CAPTION	CALLOUTS
What Is an LLM?	An LLM uses current context and learned weights to generate one response token at a time.	Context enters: The current prompt and response so far provide the visible input for this run. Weights shape scores: Learned weights from training shape next-token scores without acting like a searchable database. Tokens emerge: The response grows as one selected token is appended, then the model runs again.
Pretraining	Pretraining shapes broad reusable patterns from many examples before ordinary use.	Many examples: The training loop repeats across large collections of examples. Weights change: Prediction error drives durable weight updates during training. Not an archive: Pretraining does not create a perfect searchable memory of every source.
Alignment	Alignment shapes behavior through training, feedback, policy, evaluation, and runtime controls.	Behavior shaping: Training and feedback can encourage preferred response patterns. Runtime controls: Policies, filters, and system design can steer behavior during use. No conscience: Alignment does not give the model moral understanding or a truth guarantee.
Multimodal AI	Multimodal systems can represent or process multiple media types together without human-like perception.	Multiple media: Text, images, audio, video, or code can enter some systems. Shared representations: The system encodes media into model-usable forms or linked spaces. Not sensation: Processing media is not the same as seeing, hearing, or experiencing as a person does.
Benefits Worth Taking Seriously	AI benefits can be real and useful when bounded by context, safeguards, evidence, and human review.	Useful supports: Drafting, translation, summarization, accessibility, and coding help can be valuable with review. Bounded use: The value depends on task fit, evidence, safeguards, and context. Human review: People still judge quality, stakes, and consequences.
Better AI Is a Choice	Better AI outcomes depend on design, data choices, evaluation, incentives, governance, and human review.	Design choices: Task fit, model size, deployment, and safeguards shape outcomes. Institutional choices: Governance, incentives, evaluation, and labor planning affect who benefits. Human review: Better outcomes require accountable choices, not a single switch.

VISUAL AID

What to picture



What Is an LLM?

Score, choose, append, repeat

An LLM uses current context and learned weights to generate one response token at a time.

- **Context enters** The current prompt and response so far provide the visible input for this run.
- **Weights shape scores** Learned weights from training shape next-token scores without acting like a searchable database.
- **Tokens emerge** The response grows as one selected token is appended, then the model runs again.

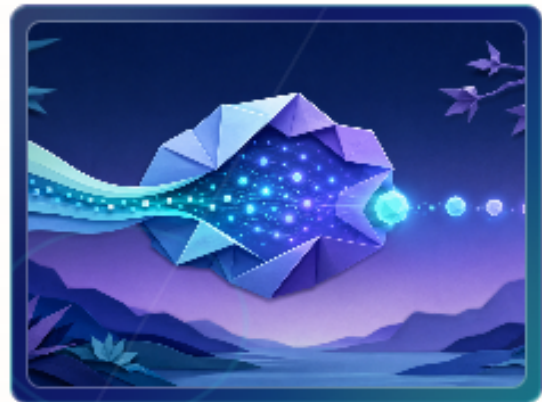
Key takeaway:

An LLM turns context into next-token probabilities.

What Is an LLM?: Image 2 visual at 390px.

VISUAL AID

What to picture



What Is an LLM?

Score, choose, append, repeat

An LLM uses current context and learned weights to generate one response token at a time.

- **Context enters** The current prompt and response so far provide the visible input for this run.
- **Weights shape scores** Learned weights from training shape next-token scores without acting like a searchable database.
- **Tokens emerge** The response grows as one selected token is appended, then the model runs again.

Key takeaway:

An LLM turns context into next-token probabilities.

What Is an LLM?: Image 2 visual at 320px.

VISUAL AID

What to picture



What Is an LLM?

Score, choose, append, repeat

An LLM uses current context and learned weights to generate one response token at a time.

- **Context enters** The current prompt and response so far provide the visible input for this run.
- **Weights shape scores** Learned weights from training shape next-token scores without acting like a searchable database.
- **Tokens emerge** The response grows as one selected token is appended, then the model runs again.

Key takeaway:

An LLM turns context into next-token probabilities.

What is an LLM?: Image 2 visual on desktop.

VISUAL AID

What to picture



Pretraining

Scale, not perfect recall

Pretraining shapes broad reusable patterns from many examples before ordinary use.

- **Many examples** The training loop repeats across large collections of examples.
- **Weights change** Prediction error drives durable weight updates during training.
- **Not an archive** Pretraining does not create a perfect searchable memory of every source.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Pretraining: Image 2 visual at 390px.

VISUAL AID

What to picture



Pretraining

Scale, not perfect recall

Pretraining shapes broad reusable patterns from many examples before ordinary use.

- **Many examples** The training loop repeats across large collections of examples.
- **Weights change** Prediction error drives durable weight updates during training.
- **Not an archive** Pretraining does not create a perfect searchable memory of every source.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Pretraining: Image 2 visual at 320px.

VISUAL AID

What to picture



Pretraining

Scale, not perfect recall

Pretraining shapes broad reusable patterns from many examples before ordinary use.

- **Many examples** The training loop repeats across large collections of examples.
- **Weights change** Prediction error drives durable weight updates during training.
- **Not an archive** Pretraining does not create a perfect searchable memory of every source.

Key takeaway:

Pretraining builds broad capability, not perfect source recall.

Pretraining: Image 2 visual on desktop.

VISUAL AID

What to picture



Alignment

Shaping behavior, not conscience

Alignment shapes behavior through training, feedback, policy, evaluation, and runtime controls.

- **Behavior shaping** Training and feedback can encourage preferred response patterns.
- **Runtime controls** Policies, filters, and system design can steer behavior during use.
- **No conscience** Alignment does not give the model moral understanding or a truth guarantee.

Key takeaway:

Alignment shapes behavior, but it is not built-in morality.

Alignment: Image 2 visual at 390px.

VISUAL AID

What to picture



Alignment

Shaping behavior, not conscience

Alignment shapes behavior through training, feedback, policy, evaluation, and runtime controls.

- **Behavior shaping** Training and feedback can encourage preferred response patterns.
- **Runtime controls** Policies, filters, and system design can steer behavior during use.
- **No conscience** Alignment does not give the model moral understanding or a truth guarantee.

Key takeaway:

Alignment shapes behavior, but it is not built-in morality.

Alignment: Image 2 visual at 320px.

VISUAL AID

What to picture



Alignment

Shaping behavior, not conscience

Alignment shapes behavior through training, feedback, policy, evaluation, and runtime controls.

- **Behavior shaping** Training and feedback can encourage preferred response patterns.
- **Runtime controls** Policies, filters, and system design can steer behavior during use.
- **No conscience** Alignment does not give the model moral understanding or a truth guarantee.

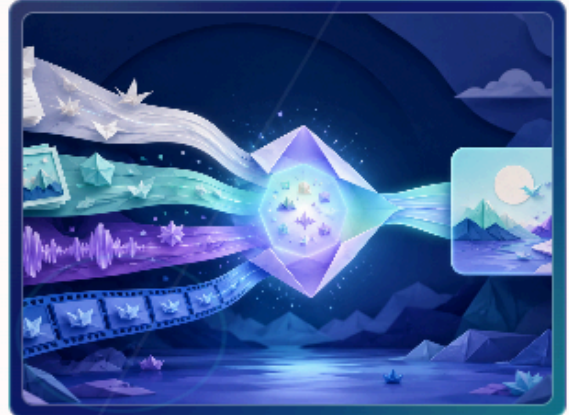
Key takeaway:

Alignment shapes behavior, but it is not built-in morality.

Alignment: Image 2 visual on desktop.

VISUAL AID

What to picture



Multimodal AI

Input, representation, output

Multimodal systems can represent or process multiple media types together without human-like perception.

- **Multiple media** Text, images, audio, video, or code can enter some systems.
- **Shared representations** The system encodes media into model-usable forms or linked spaces.
- **Not sensation** Processing media is not the same as seeing, hearing, or experiencing as a person does.

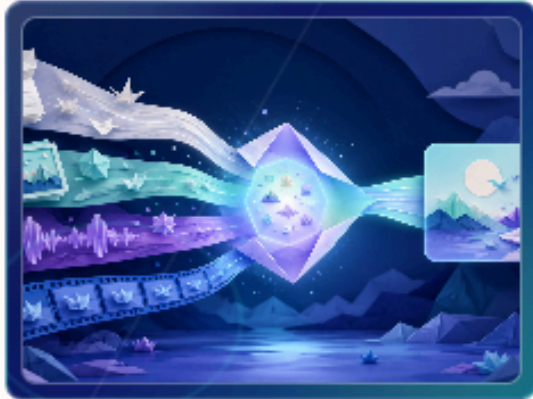
Key takeaway:

Multimodal means more than one media type, not human-like perception.

Multimodal AI: Image 2 visual at 390px.

VISUAL AID

What to picture



Multimodal AI

Input, representation, output

Multimodal systems can represent or process multiple media types together without human-like perception.

- **Multiple media** Text, images, audio, video, or code can enter some systems.
- **Shared representations** The system encodes media into model-usable forms or linked spaces.
- **Not sensation** Processing media is not the same as seeing, hearing, or experiencing as a person does.


Key takeaway:

Multimodal means more than one media type, not human-like perception.

Multimodal AI: Image 2 visual at 320px.

VISUAL AID

What to picture



Multimodal AI

Input, representation, output

Multimodal systems can represent or process multiple media types together without human-like perception.

- **Multiple media** Text, images, audio, video, or code can enter some systems.
- **Shared representations** The system encodes media into model-usable forms or linked spaces.
- **Not sensation** Processing media is not the same as seeing, hearing, or experiencing as a person does.

Key takeaway:

Multimodal means more than one media type, not human-like perception.

Multimodal AI: Image 2 visual on desktop.

VISUAL AID

What to picture



Benefits Worth Taking Seriously

Benefits without utopia

AI benefits can be real and useful when bounded by context, safeguards, evidence, and human review.

- **Useful supports** Drafting, translation, summarization, accessibility, and coding help can be valuable with review.
- **Bounded use** The value depends on task fit, evidence, safeguards, and context.
- **Human review** People still judge quality, stakes, and consequences.

Key takeaway:

Benefits can be real and still bounded.

Benefits Worth Taking Seriously: Image 2 visual at 390px.

VISUAL AID

What to picture



Benefits Worth Taking Seriously

Benefits without utopia

AI benefits can be real and useful when bounded by context, safeguards, evidence, and human review.

- **Useful supports** Drafting, translation, summarization, accessibility, and coding help can be valuable with review.
- **Bounded use** The value depends on task fit, evidence, safeguards, and context.
- **Human review** People still judge quality, stakes, and consequences.


Key takeaway:

Benefits can be real and still bounded.

Benefits Worth Taking Seriously: Image 2 visual at 320px.

VISUAL AID

What to picture



Benefits Worth Taking Seriously

Benefits without utopia

AI benefits can be real and useful when bounded by context, safeguards, evidence, and human review.

- **Useful supports** Drafting, translation, summarization, accessibility, and coding help can be valuable with review.
- **Bounded use** The value depends on task fit, evidence, safeguards, and context.
- **Human review** People still judge quality, stakes, and consequences.

Key takeaway:

Benefits can be real and still bounded.

Benefits Worth Taking Seriously: Image 2 visual on desktop.

VISUAL AID

What to picture



Better AI Is a Choice

Responsible AI is chosen

Better AI outcomes depend on design, data choices, evaluation, incentives, governance, and human review.

- **Design choices** Task fit, model size, deployment, and safeguards shape outcomes.
- **Institutional choices** Governance, incentives, evaluation, and labor planning affect who benefits.
- **Human review** Better outcomes require accountable choices, not a single switch.

Key takeaway:

Harms are shaped by choices, not destiny.

Better AI Is a Choice: Image 2 visual at 390px.

VISUAL AID

What to picture



Better AI Is a Choice

Responsible AI is chosen

Better AI outcomes depend on design, data choices, evaluation, incentives, governance, and human review.

- **Design choices** Task fit, model size, deployment, and safeguards shape outcomes.
- **Institutional choices** Governance, incentives, evaluation, and labor planning affect who benefits.
- **Human review** Better outcomes require accountable choices, not a single switch.

Key takeaway:

Harms are shaped by choices, not destiny.

Better AI Is a Choice: Image 2 visual at 320px.

VISUAL AID

What to picture



Better AI Is a Choice

Responsible AI is chosen


Better AI outcomes depend on design, data choices, evaluation, incentives, governance, and human review.

- **Design choices** Task fit, model size, deployment, and safeguards shape outcomes.
- **Institutional choices** Governance, incentives, evaluation, and labor planning affect who benefits.
- **Human review** Better outcomes require accountable choices, not a single switch.

Key takeaway:

Harms are shaped by choices, not destiny.

Better AI Is a Choice: Image 2 visual on desktop.




What Is an LLM?

A learned next-token prediction system

A large language model is a learned system that turns current context into probabilities for the next token.

VISUAL AID

What to picture



What Is an LLM?

Score, choose, append, repeat

An LLM uses current context and learned weights to generate one response token at a time.



Home



Journey



Play



Glossary



Badge

What is an LLM?: Journey card smoke at 390px.



Pretraining

Broad learning before normal use

Pretraining is the broad first training phase where a model learns general patterns by repeatedly predicting tokens across large datasets.

VISUAL AID

What to picture



Pretraining

Scale, not perfect recall

Pretraining shapes broad reusable patterns from many examples before



Home



Journey



Play



Glossary



Badge

Pretraining: Journey card smoke at 390px.



Alignment

Shaping behavior toward human goals

Alignment is the effort to shape AI behavior toward human instructions, preferences, safety boundaries, and values.

VISUAL AID

What to picture



Alignment

Shaping behavior, not conscience

Alignment shapes behavior through training feedback policy evaluation and



Home



Journey



Play



Glossary



Badge

Alignment: Journey card smoke at 390px.




Multimodal AI

Multiple media types together

Multimodal AI works with more than one kind of input or output, such as text, images, audio, video, or code.

VISUAL AID

What to picture



Multimodal AI

Input, representation, output

Multimodal systems can represent or process multiple media types together without human-like perception.



Home



Journey



Play



Glossary



Badge

Multimodal AI: Journey card smoke at 390px.



The Perfect Storm

Why LLMs arrived now

Modern LLMs became possible when human-created digital data, powerful compute and storage, deep-learning methods, human labor, and economic incentives converged.

VISUAL AID


What to picture



The Perfect Storm

Why LLMs arrived now

Data, compute, storage, methods, labor



The Perfect Storm: Journey card smoke at 390px.




Collective Intelligence, Extracted

No creators, no model

Generative AI depends on patterns learned from human-created language, art, code, documents, research, and culture.

VISUAL AID

What to picture



Collective Intelligence, Extracted: Journey card smoke at 390px.




Costs We Must Count

The answer feels weightless. The infrastructure is not.

Generative AI has physical, social, cultural, and ethical costs that must be counted honestly.

VISUAL AID

What to picture



Costs We Must Count

The answer is not weightless



Home



Journey



Play



Glossary



Badge

Costs We Must Count: Journey card smoke at 390px.




Risk vs Myth

Practical institutional literacy

Risk literacy means separating practical, mechanism-based AI risks from myths or unsupported claims about what models are.

VISUAL AID

What to picture



Risk vs Myth

Mechanisms, not myth

Model-literate users separate practical risks from unsupported myths, then



Home



Journey



Play



Glossary



Badge

Risk vs Myth: Journey card smoke at 390px.

NEW DAWN

35 / 39



Benefits Worth Taking Seriously

Real help, not utopian hype

AI benefits are strongest when models amplify human judgment rather than replacing accountability.

VISUAL AID


What to picture



Benefits Worth Taking Seriously

Benefits without utopia

AI benefits can be real and useful when



Benefits Worth Taking Seriously: Journey card smoke at 390px.

NEW DAWN

36 / 39




Human-Centered AI

Tools should serve human dignity

Human-centered AI asks whether a system supports dignity, learning, responsibility, creativity, relationships, and the common good.


VISUAL AID

What to picture



Human-Centered AI

Tools should serve dignity



Human-Centered AI: Journey card smoke at 390px.

NEW DAWN

37 / 39



Better AI Is a Choice

Design, policy, and incentives matter

The costs are shaped by design, deployment, governance, and incentives.

VISUAL AID


What to picture



Better AI Is a Choice

Responsible AI is chosen

Better AI outcomes depend on design,



Better AI Is a Choice: Journey card smoke at 390px.

NEW DAWN

39 / 39




Model Literate Synthesis

Understand the machine, then decide how to use it

Model literacy means understanding both the mechanics and the human consequences.

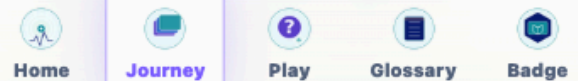
VISUAL AID

What to picture



Model Literate Synthesis

Mechanics plus judgment



Model Literate Synthesis: Journey card smoke at 390px.

VISUAL AID

What to picture



Collective Intelligence, Extracted No creators, no model

Human-created traces flow into learned model patterns while provenance and responsibility remain human questions.

- **Human traces** Writing, code, art, research, forums, and documentation can leave learnable traces.
- **Collection questions** Provenance, consent, copyright, attribution, and compensation matter.
- **Model limit** The model absorbs statistical patterns without responsibility.

Key takeaway:

The model did not create its abilities alone.

Already-live smoke: Collective Intelligence still renders its concept image.

BADGE

Prompt Life: Model Literate

Under construction while the evidence model is reviewed.



Under construction

Pending human review

Prompt Life: Model Literate is under construction. The Journey, checkpoint questions, and evidence model are being tested before any badge is issued.



Home



Journey



Play



Glossary



Badge

Badge page smoke screenshot.

PLAY LAB

Play to learn

Short, calm challenges that make model mechanics visible.

Progress saved on this device. 0 of 5 practice challenges show saved completion or review progress.

Each challenge practices one model-literacy move.

More below ↓

Play challenges

Suggested order follows the Journey.



Glossary Dojo

Ready

Practice concept discrimination.

Practice move: Name concepts

Recommended: Anytime

Anytime

Glossary

Learning cards: All learning cards

Action: Choose, compare, review.

Time: 4 min



Home



Journey



Play



Glossary



Badge

Play page smoke screenshot.

DEFINITIONS PLUS RELATIONSHIPS

Glossary

Use each term with its neighbor. That is where the mental model gets stable.

PRACTICE MODE

Glossary Dojo

Practice twelve terms through definitions, relationships, and common mix-ups.

Practice 12 terms

Search terms

token, tensor, softmax...

A-Z

Learning path

Find terms alphabetically.

Accessibility

Designing tools and materials so more people can perceive, understand, navigate, and participate.

AI can support accessibility through drafting, summarization, translation, captioning, alternate formats, and assistive workflows when reviewed



Home



Journey



Play



Glossary



Badge

Glossary page smoke screenshot.

Command Results

COMMAND	STATUS	NOTES
npm run process:visual-assets	pass	> promptlife@0.28.9 process:visual-assets > node scripts/process-journey-visual-assets.mjs Processed what-is-an-llm.png source: public/assets/journey-visuals/v0-28/source/what-is-an-llm.png optimized: public/assets/journey-visuals/v0-28/optimized/what-is-an-llm.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/what-is-an-llm.png Processed pretraining.png source: public/assets/journey-visuals/v0-28/source/pretraining.png optimized: public/assets/journey-visuals/v0-28/optimized/pretraining.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/pretraining.png Processed alignment.png source: public/assets/journey-visuals/v0-28/source/alignment.png optimized: public/assets/journey-visuals/v0-28/optimized/alignment.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/alignment.png Processed multimodal-ai.png source: public/assets/journey-visuals/v0-28/source/multimodal-ai.png optimized: public/assets/journey-visuals/v0-28/optimized/multimodal-ai.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/multimodal-ai.png Processed benefits-worth-taking-seriously.png source: public/assets/journey-visuals/v0-28/source/benefits-worth-taking-seriously.png optimized: public/assets/journey-visuals/v0-28/optimized/benefits-worth-taking-seriously.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/benefits-worth-taking-seriously.png Processed better-ai-choice.png source: public/assets/journey-visuals/v0-28/source/better-ai-choice.png optimized: public/assets/journey-visuals/v0-28/optimized/better-ai-choice.png thumbnail: public/assets/journey-visuals/v0-28/thumbnails/better-ai-choice.png
npm run typecheck	pass	> promptlife@0.28.9 typecheck > tsc --noEmit
npm run build	pass	> promptlife@0.28.9 build > vite build vite v8.0.14 building client environment for production... [2K transforming...✓ 39 modules transformed. rendering chunks... computing gzip size... dist/index.html1.03 kB gzip: 0.46 kB dist/assets/index-DMxoEmxG.css151.39 kB gzip: 29.10 kB dist/assets/index-DX0H-pqL.js1,078.31 kB gzip: 262.22 kB ✓ built in 414ms [plugin builtin:vite-reporter] (!) Some chunks are larger than 500 kB after minification. Consider: - Using dynamic import() to code-split the application - Use build.rolldownOptions.output.codeSplitting to improve chunking: https://rolldown.rs/reference/OutputOptions.codeSplitting - Adjust chunk size limit for this warning via build.chunkSizeWarningLimit.
npm run build:pages	pass	> promptlife@0.28.9 build:pages > GITHUB_PAGES=true vite build vite v8.0.14 building client environment for production... [2K transforming...✓ 39 modules transformed. rendering chunks... computing gzip size... dist/index.html1.05 kB gzip: 0.47 kB dist/assets/index-DMxoEmxG.css151.39 kB gzip: 29.10 kB dist/assets/index-CjX-jPpL.js1,078.34 kB gzip: 262.23 kB ✓ built in 183ms [plugin builtin:vite-reporter] (!) Some chunks are larger than 500 kB after minification. Consider: - Using dynamic import() to code-split the application - Use build.rolldownOptions.output.codeSplitting to improve chunking: https://rolldown.rs/reference/OutputOptions.codeSplitting - Adjust chunk size limit for this warning via build.chunkSizeWarningLimit.
npm run audit:visual-assets	pass	> promptlife@0.28.9 audit:visual-assets > node scripts/audit-visual-assets.mjs Visual asset audit passed. Checked 12 manifest entries; live assets before: 6; live assets after: 12; newly live: 6; missing assets: 0; rejected assets: 0; pending i
npm run audit:visual-aids	pass	> promptlife@0.28.9 audit:visual-aids > node scripts/audit-visual-aids.mjs && node scripts/audit-word-wrap.mjs Visual aid audit passed: 39 aids, 0 mismatches, 0 readiness issues. Word-wrap audit pass. Checked 82 files. Unsafe CSS: 0 Known bad splits: 0 Long internal labels: 0 Diagram paragraphs: 0
npm run audit:visual-overflow	pass	> promptlife@0.28.9 audit:visual-overflow > node scripts/audit-visual-overflow.mjs Visual overflow audit passed. Checked 39 learner-route visual aids at 320px and 390px.
npm run audit:word-wrap	pass	> promptlife@0.28.9 audit:word-wrap > node scripts/audit-word-wrap.mjs Word-wrap audit pass. Checked 82 files. Unsafe CSS: 0 Known bad splits: 0 Long internal labels: 0 Diagram paragraphs: 0
npm run audit:learner-copy	pass	> promptlife@0.28.9 audit:learner-copy > node scripts/audit-learner-copy.mjs Learner-copy audit passed. Checked 19 learner-facing source files.

COMMAND	STATUS	NOTES
npm run audit:learner-leaks	pass	> promptliffe@0.28.9 audit:learner-leaks > node scripts/audit-learner-leaks.mjs Learner leak audit passed. Static records: 45; allowed: 45; rendered screens: 7.
npm run audit:language	pass	> promptliffe@0.28.9 audit:language > node scripts/audit-language.mjs Language audit passed. Checked 61 current source/docs files for banned wording.
npm run audit:answers	pass	/a Fixed. Not a single-choice answer list. exercise:pretraining-patterns-not-recall Pattern Learning, Not Perfect Recall toggle-state 2 2 no n/a Fixed. Not a single-choice answer list. exercise:overfitting-held-out-check Validation Check toggle-state 2 1 no n/a Fixed. Not a single-choice answer list. exercise:fine-tuning-context-decoding-sort Open Run or Durable Change? sort-buckets 5 n/a no n/a Fixed. Learner sorts into buckets. exercise:alignment-method-groups Alignment Method Groups sort-buckets 8 n/a no n/a Fixed. Learner sorts into buckets. exercise:attention-relevance Attention Is Relevance connect-nodes 10 7 no n/a Fixed. Not a single-choice answer list. exercise:context-window-fell-out Context Window: What Fell Out? tap-choice 4 1, 2 yes 2, 3 Randomized. Correctness follows item id. exercise:training-nudge Training Nudge toggle-state 2 1 no n/a Fixed. Not a single-choice answer list. exercise:softmax-funnel Softmax Funnel tap-choice 4 1 yes 4 Randomized. Correctness follows item id. exercise:mlp-feature-reshape MLP Feature Reshape toggle-state 2 2 no n/a Fixed. Not a single-choice answer list. exercise:hidden-state-sort Hidden State Sort sort-buckets 6 n/a no n/a Fixed. Learner sorts into buckets. exercise:open-book-or-learned Open Book or Learned? sort-buckets 8 n/a no n/a Fixed. Learner sorts into buckets. exercise:brain-metaphor-boundary Brain Metaphor Boundary tap-multiple 7 n/a no n/a Fixed. Multi-select grouping is clearer in authored order. prompt-run:prompt-run-prompt-response Prompt or Response? label-tokens 1 n/a no n/a Fixed. Token sequence/context order is meaningful. prompt-run:prompt-run-tokenizer Tokenizer label-tokens 1 n/a no n/a Fixed. Token sequence/context order is meaningful. prompt-run:prompt-run-token-ids Token IDs drag-match 7 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-embedding Embedding Lookup drag-match 7 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-tensor Tensor Stack drag-order 3 n/a no n/a Fixed. Order is the learning objective. prompt-run:prompt-run-attention Attention connect-nodes 10 7 no n/a Fixed. Not a single-choice answer list. prompt-run:prompt-run-mlp MLP Feature Reshape drag-match 5 n/a no n/a Fixed. Learner matches paired items. prompt-run:prompt-run-hidden-state Hidden State tap-choice 3 1 yes 2 Randomized. Correctness follows item id. prompt-run:prompt-run-logits Logits tap-choice 4 1 yes 4 Randomized. Correctness follows item id. prompt-run:prompt-run-softmax Softmax Funnel tap-choice 4 1 yes 3 Randomized. Correctness follows item id. prompt-run:prompt-run-sampling Sampling next-token-pick 6 1, 2, 3 yes 3, 4, 5 Randomized. Correctness follows item id. prompt-run:prompt-run-append-repeat Append and Repeat drag-order 4 n/a no n/a Fixed. Order is the learning objective. prompt-run:prompt-run-final-order Put the Run in Order drag-order 12 n/a no n/a Fixed. Order is the learning objective. ## Checkpoint Bank Audit Live Journey cards: 39 Live checkpoint questions: 39 Draft cards: 39 Draft questions: 136 Draft choices: 544 Draft wrong-answer distractors: 408 Draft question-count distribution: {"2":5,"3":12,"4":20,"5":2} Pilot cards: 6 Pilot questions: 18 Pilot choices: 72 Pilot wrong-answer distractors: 54 Pilot model-thinking/application/mechanism questions: 100% Pilot direct-definition questions: 0% Revised active-development cards: 6 Revised active-development questions: 19 Revised active-development choices: 76 Revised active-development wrong-answer distractors: 57 Revised audit-seed correct-answer positions: [4,3,1,4,2,3,1,2,4,1,4,1,1,3,4,3,3,2,2] v0.27.10 active-development cards: 12 v0.27.10 active-development questions: 39 v0.27.10 active-development choices: 156 v0.27.10 active-development wrong-answer distractors: 117 v0.27.10 audit-seed correct-answer positions: [4,1,3,4,2,2,4,3,4,3,3,4,3,4,2,4,2,1,2,1,2,4,1,1,4,1,1,4,3,4,4,4,1,1,1,3,3,1] v0.27.15 active-development cards: 39 v0.27.15 active-development questions: 136 v0.27.15 active-development choices: 544 v0.27.15 active-development wrong-answer distractors: 408 v0.27.15 model-thinking/application/mechanism questions: 136 v0.27.15 direct-definition questions: 0 v0.27.15 audit-seed correct-answer positions: [4,1,2,2,1,4,4,1,2,2,4,3,1,1,3,4,2,2,1,3,2,4,1,3,3,2,4,3,3,4,3,2,1,4,1,2,2,2,1,4,1,1,4,2,4,1,1,3,4,2,1,3,2,3,3,4,4,3,3,4,2,1,3,1,1,2,4,2,2,1,2,1,4,4,3,2,4,3,1,4,3, Checkpoint audit passed.
npm run audit:checkpoints	pass	> promptliffe@0.28.9 audit:checkpoints > node scripts/audit-checkpoints.mjs ## Checkpoint Bank Audit Live Journey cards: 39 Live checkpoint questions: 39 Draft cards: 39 Draft questions: 136 Draft choices: 544 Draft wrong-answer distractors: 408 Draft question-count distribution: {"2":5,"3":12,"4":20,"5":2} Pilot cards: 6 Pilot questions: 18 Pilot choices: 72 Pilot wrong-answer distractors: 54 Pilot model-thinking/application/mechanism questions: 100% Pilot direct-definition questions: 0% Revised active-development cards: 6 Revised active-development questions: 19 Revised active-development choices: 76 Revised active-development wrong-answer distractors: 57 Revised audit-seed correct-answer positions: [4,3,1,4,2,3,1,2,4,1,4,1,1,3,4,3,3,2,2] v0.27.10 active-development cards: 12 v0.27.10 active-development questions: 39 v0.27.10 active-development choices: 156 v0.27.10 active-development wrong-answer distractors: 117 v0.27.10 audit-seed correct-answer positions: [4,1,3,4,2,2,4,3,4,3,3,4,3,4,2,4,2,1,2,1,2,4,1,1,4,1,1,4,3,4,4,4,1,1,1,3,3,1] v0.27.15 active-development cards: 39 v0.27.15 active-development questions: 136 v0.27.15 active-development choices: 544 v0.27.15 active-development wrong-answer distractors: 408 v0.27.15 model-thinking/application/mechanism questions: 136 v0.27.15 direct-definition questions: 0 v0.27.15 audit-seed correct-answer positions: [4,1,2,2,1,4,4,1,2,2,4,3,1,1,3,4,2,2,1,3,2,4,1,3,3,2,4,3,3,4,3,2,1,4,1,2,2,2,1,4,1,1,4,2,4,1,1,3,4,2,1,3,2,3,3,4,4,3,3,4,2,1,3,1,1,2,4,2,2,1,2,1,4,4,3,2,4,3,1,4,3, Checkpoint audit passed.

COMMAND	STATUS	NOTES
npm run audit:question- clues	pass	> promptlife@0.28.9 audit:question-clues > node scripts/audit-question-clues.mjs Question-clue audit completed. Term-echo candidates: 36 Answer-length warnings: 43 Vague-stem warnings: 0 Wrote docs/journey/checkpoints/question-clue-audit-v0-27-15.md
npm run audit:exercises	pass	> promptlife@0.28.9 audit:exercises > node scripts/audit-exercises.mjs Exercise audit passed. Inventoried 39 Journey interactions and exported review PDFs.

BACKLOG

Remaining Visual Strategy Backlog

Still missing from the Image 2 manifest after this pass: none.