

Play Games UX + Mechanics Audit

Current-state audit of the five visible Play challenges: Glossary Dojo, Attention Match, Probability Picker, Context Stack, and Prompt Run. This pass did not implement gameplay changes.

Executive Findings

Strongest Today

Glossary Dojo is the most complete loop. **Context Stack** is the clearest small model-mechanics game.

Needs A Build Pass

Probability Picker is still a placeholder, and **Attention Match** still uses the old Attention Weave mechanic.

Capstone Risk

Prompt Run remains instructionally essential, but it still feels closer to a long guided worksheet than a replayable challenge.

Mobile Result

390px and 320px Play landing checks showed no horizontal overflow. Cards are tall at 320px, but vertical scrolling is obvious.

Scope

Reviewed `src/features/play`, the challenge registry, shared Play components, shared storage, Glossary Dojo integration, Context Stack, Attention Match, Probability Picker, Prompt Run, Badge language, reset behavior, and Play styles. Screenshots were captured from the local app at `http://127.0.0.1:5173/?debug=1&v=0268-play-audit`.

No Journey order, badge logic, checkpoint randomization, dependencies, generated assets, or core game mechanics were changed in this pass.

Current Slate

Challenge	Current route/component	Gameplay loop	Shared shell?	Shared storage?	Recommendation
Glossary Dojo	<code>GlossaryDojoGame</code> via <code>gameId === glossary-dojo</code>	Start or resume a 12-question round; choose glossary matches; review feedback; repeat or review missed terms.	Yes: <code>PlayChallengeShell</code> , status pill, progress rail, feedback panel.	Yes for summary attempts/completions; Dojo mastery remains in <code>promptlife.glossaryDojo.v1</code> .	Keep. Preserve loop; polish only around stale state, mobile feedback, and copy.
Attention Match	<code>AttentionWeave</code> rendered with title "Attention Match"	Tap one token, then another, to draw relevance arcs for one pronoun-style dependency.	No. Uses legacy <code>GameHeader</code> , <code>weave-board</code> , and old game patterns.	Completion bridges to <code>attention-match</code> ; legacy <code>attention-weave</code> insight is preserved.	Rebuild as Attention Match v1 with several short dependency examples and shared Play shell.

Challenge	Current route/component	Gameplay loop	Shared shell?	Shared storage?	Recommendation
Probability Picker	Disabled card in the normal slate; internal <code>FoundationReadyScreen</code> exists.	No playable loop yet. Card teaches the future slot: logits -> softmax -> sample one token.	Card uses shared component. Placeholder route uses shared shell but is not reachable from the disabled card.	No fake attempts or completions; reports <code>Coming soon</code> .	Build next. It fills the largest current mechanics gap: probability is not truth.
Context Stack	<code>ContextStack</code> via <code>gameId === context-stack</code>	Tap cards into a limited context window across three rounds; older cards fall out; retrieved evidence must enter context.	Yes: shared shell, status pill, progress rail, board, choice buttons, feedback, completion panel.	Yes: starts, progress, completions, review-suggested outcomes, and misconception tags write to <code>promptlife.playChallenges.v1</code> .	Keep. Small v2.1 polish later; not the next urgent rebuild.
Prompt Run	<code>PromptRunScreen</code> via <code>trace-one-prompt</code> or <code>prompt-run</code>	Walk through 12 steps plus final ordering challenge; each step uses an exercise before Continue unlocks.	No. Uses legacy <code>GameHeader</code> , <code>ExerciseShell</code> , and Prompt Run-specific progress UI.	Shared completion is written at final completion; detailed step progress remains in <code>promptlife:v1:promptRunProgress</code> .	Revise as Prompt Run v2 after Probability Picker and Attention Match.

Retired Handling

Token Pipeline Relay and How AI Learns are absent from the normal Play slate. They remain in `retiredPlayChallenges` for compatibility. If internal state sets `gameId` to one of those retired IDs, the app renders a calm compatibility screen that says progress was not lost. There is no visible normal Play card for either retired activity.

Score Matrix

Scores are 1-5 where 5 is strongest, except maintenance burden where 1 is low/easy and 5 is high/costly.

Challenge	Learning clarity	Mechanics fit	Misconception value	Mobile comfort	Replay value	Calm/fun factor	Badge relevance	Maintenance burden
Glossary Dojo	5	5	5	4	5	4	2	4
Attention Match	3	3	4	3	2	3	2	2
Probability Picker	3	1	3	5	1	2	3	1
Context Stack	5	5	5	4	4	4	3	3
Prompt Run	4	3	5	3	2	3	5	4

Game-By-Game Notes

Glossary Dojo - Keep

Dojo has the cleanest learning loop: real glossary terms, close distractors, explicit feedback, repeat/review modes, and independent mastery. It preserves old Dojo progress while also writing shared Play summary progress. Mobile risk is mainly vertical: answer choices can run under the bottom nav until the learner scrolls, but controls remain usable.

Attention Match - Rebuild

The concept is right, but the current action is one legacy Attention Weave example. It should become a small set of sentence cards: choose a target token, connect it to the most relevant prior token, then explain that attention is relevance between positions, not awareness.

Probability Picker - Build

This is the most important empty slot. It should teach logits, softmax, probabilities, sampling, uncertainty, and why likely does not mean true. The current disabled card is coherent and does not create fake progress.

Context Stack - Keep

Context Stack now has a strong v2 loop: place cards, see overflow, keep key context visible, and use retrieved evidence only after it enters the window. It already uses the shared Play shell and storage helpers. Future polish should focus on shorter card labels at 320px and slightly more varied replay.

Prompt Run - Revise

Prompt Run carries the whole "day in the life of a prompt" story and badge relevance. Its risk is experience shape: 13 steps plus exercises can feel like a long lesson path. A v2 should make the prompt visibly move through fewer stronger scenes, while keeping the same core distinctions.

Recommended Build Order

- 1. **Probability Picker v1.** Fill the empty slate slot and teach the Decision Room mechanics that no current Play game practices directly.
- 2. **Attention Match v1.** Replace the compatibility bridge with a purpose-built shared-shell challenge.
- 3. **Prompt Run v2.** Keep the capstone but shorten the feel, add stronger continuity, and align it with shared Play language.

Context Stack no longer needs to be first because its v2 refactor is already in place. A later Context Stack v2.1 pass can tune labels, replay, and visual density.

Unified Play Interaction Language

Pattern	Recommendation
Action verbs	Name, connect, choose, place, trace. Show these as practice moves, but keep the sentence after them concrete.
Status pills	Ready, In progress, Review suggested, Completed, Coming soon, Retired. Avoid score-like status labels.
Stat row	Played, Current, Progress saved, Review suggested, Completed. Replace future "Best" wording with "Progress" or "Last round" where possible.
Board zones	Use consistent zones: incoming/source, working/choice area, current tray/window, fell-out/history, feedback, reflection.
Cards/chips/trays	Cards are movable ideas; chips are terms or tokens; trays are temporary holding areas; rails show progress through a round, not performance.
Feedback panels	Correct: "Good distinction" or "Insight strengthened." Review: "This choice reveals a common mix-up." Completion: "Insight unlocked" or "Round completed."
Storage	Record attempts on meaningful start, update progress during rounds, complete only when an insight or round finishes, and preserve legacy keys by bridging summaries.

Learner-Facing Copy Issues Found

- Best: Not yet , Best result , and Best: N of 12 are mild score-like remnants. Prefer Progress , Last round , or Current .
- Back to games appears on Prompt Run and Attention Match while the landing now says Play challenges. Prefer Back to Play .
- Attention Match still says "source token" and "target token" before learners have much help in that screen. A rebuild should add a one-line setup.
- Probability Picker has an internal coming-soon screen, but the visible disabled card is the only reachable placeholder from the normal slate.
- Badge stats still say 0 of 5 play challenges . The criterion panel clarifies Play is practice history, but a future pass could make the stat label even less requirement-like.
- No punitive language, leaderboards, timers, streaks, or score emphasis were found in the normal Play surface.

Mobile QA

- 390px Play landing: no horizontal overflow; five-card slate renders in order.
- 320px Play landing: no horizontal overflow; cards are tall, but vertical scrolling is clear.
- 390px route checks: Glossary Dojo, Attention Match, Context Stack, Prompt Run, and Badge loaded without horizontal overflow in captured states.
- Bottom nav stayed usable in captured states. Some lower controls require vertical scroll, which is expected in the current design.
- Reduced-motion support is present in shared animation utilities and CSS via `prefers-reduced-motion`; this was source-inspected rather than browser-emulated in this pass.

Storage Audit

- Created/read/written by Play foundation: `promptlife.playChallenges.v1`.
- Preserved/read/written by Dojo: `promptlife.glossaryDojo.v1`.
- Preserved/read by Prompt Run: `promptlife:v1:promptRunProgress` and legacy `pl.promptRunProgress`.
- Preserved/read for old game insights: `promptlife:v1:gameInsights` and legacy `pl.gameInsights`.
- Preserved/read for old flags: `promptlife:v1:traceComplete`, `promptlife:v1:learningTourComplete`, and their legacy `pl.*` counterparts.
- Reset flow source confirms Play challenge progress and Glossary Dojo practice are cleared by the existing Start over/debug clear paths.

Verification

- `npm run typecheck`: passed.
- `npm run build`: passed with the existing Vite large-chunk warning.
- `npm run build:pages`: passed with the existing Vite large-chunk warning.
- `npm run audit:answers`: passed. Total surfaces: 69; randomized surfaces: 46; excluded fixed-order surfaces: 23.

Known Issues


The Vite production build still emits the existing large-chunk warning. Probability Picker is not playable yet. Attention Match is still the renamed legacy Attention Weave interaction. Prompt Run is accurate but heavy for a Play experience.

Files Changed


- docs/play/PLAY_FOUNDATION_LOG.md
- docs/play/screenshots/v0-26-8-play-games-audit-screenshots.json
- docs/play/screenshots/v0-26-8-play-games-audit-play-landing-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-play-landing-320.png
- docs/play/screenshots/v0-26-8-play-games-audit-glossary-dojo-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-attention-match-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-probability-picker-card-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-context-stack-top-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-context-stack-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-prompt-run-top-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-prompt-run-390.png
- docs/play/screenshots/v0-26-8-play-games-audit-badge-390.png
- docs/play/prompt-life-v0-26-8-play-games-audit-report.html
- docs/play/prompt-life-v0-26-8-play-games-audit-report.pdf

Screenshots

Play landing at 390px



Play landing at 320px



Practice glossary terms

New round. Choose the best glossary match, then use the feedback to sharpen the distinction.

Question 11 of 1210 answered

TERM TO DESCRIPTION

What is Input context?

Choose the best description.

A Everything currently visible to the model: prompt tokens, conversation history, retrieved documents, and already-generated response tokens.

B AI systems that create new text, images, audio, code, video, or other media.

C The limited amount of tokens or media the model can consider at once.

D A token generated by the model during response generation.

Home

Journey

Play

Glossary

Badge

Tap a source token, then a target token. This preview asks which token it depends on.

Status

Choose a source token.

Ajealousdogchaseda

startledcatbecauseithissed

Connect 'it' to the token it most likely refers to.

Reset

Open attention in glossary

Reflection

Which two tokens did you connect, and why did that relation matter?

One sentence is enough.

Home


Journey

Play


Glossary

Badge

Probability Picker card at 390px



Context Stack top at 390px



Prompt Run top at 390px

[Back to games](#)

Prompt Run

1 of 13

Guide one prompt through the model from visible context to one selected token.

Sample prompt: User prompt: Give me one sentence with two household pets in conflict.

STEP 1

Prompt or Response?

Separate the complete user prompt from generated response tokens.

Hint

Hints are here when you want a nudge.
No penalty, no shame.

Get hint

EXERCISE

Prompt or Response?

Separate the complete user prompt from generated response tokens.

Home

Journey

Play

Glossary

Badge

Badge at 390px

PROGRESS

Prompt Life: Model Literate

Track your understanding of the prompt-to-response path.

1 of 39 lessons

0 of 5 play challenges

0 of 17 exercises

0 of 13 Prompt Run

0 reflections

Badge criteria

Home

Journey

Play

Glossary

Badge