

Content Repair and Mobile Shell Report

v0.6 turns Prompt Life into a clearer mobile-first learning path: each lesson now teaches the same way, the app shell scrolls more reliably on phones, the bottom nav has a safe-area-aware frosted gradient, and the review/PDF routes are ready for content QA.

What Is Done

- Normalized the Journey to the requested 25 lessons, including Prompt vs Response, Vectors, and How AI Learns.
- Added Brain Bridge to every lesson so brain metaphors include explicit limits.
- Added reusable visual aids and review routes: `/review/lesson-cards` and `/review/visual-aids`.
- Moved reusable exercises out of Journey lessons so Journey teaches and Play practices.
- Repaired mobile scrolling, safe-area padding, bottom nav gradient, and route/lesson scroll reset behavior.
- Added PWA metadata and a lesson-card PDF exporter.

How To Run Locally

```
npm install
```

```
npm run dev
```

Open `http://localhost:5173/`.

Verification

```
npm run typecheck : passed.
```

```
npm run build : passed.
```

```
npm run export:lesson-pdf : passed.
```

Responsive smoke check: no horizontal overflow at 320, 390, or 430px; desktop review route renders 25 lesson cards.

Files Changed

Core app: `src/data/content.ts` , `src/main.tsx` , `src/styles/global.css` , `src/components/VisualAids.tsx` , `index.html` , `public/manifest.webmanifest` , `package.json` , `scripts/export-lesson-pdf.mjs` .

Docs and artifacts: `README.md` , `docs/ACCESSIBILITY_QA.md` , `docs/REVIEW_NOTES.md` , `docs/CONTENT_REPAIR_V0_6.md` , `docs/VISUAL_AIDS_V0_6.md` , `docs/CONTENT_QA_CHECKLIST.md` , `docs/review/LESSON_CONTENT_MATRIX.md` , `docs/review/prompt-life-lesson-cards-v0-6.pdf` , and v0.6 screenshots.

Next Three Improvements

- Add automated Playwright route tests for lessons, review routes, progress reset, and bottom-nav visibility.
- Add an axe accessibility pass and strengthen the glossary drawer focus trap.
- Add a small service worker only after the v0.6 content shell is stable enough to benefit from offline caching.

Screenshots



A DAY IN THE LIFE OF A PROMPT

Prompt Life

Model literacy for the AI era: clear mechanics, useful metaphors, and a little play.



CURRENT PATH

Start at the beginning

0% complete. Next stop: What Is an LLM?

0%

Start journey

Home

Journey

Play

Glossary

Badge

Home: compact hero, current path, and visible start button.

1

Before Morning

What an LLM is, where it came from, and which processes can durably change weights.

The prompt has not arrived yet. The model is being shaped before ordinary use.

FOCUS
Separate durable training from temporary inference.

KEEP CLEAR
Training and fine-tuning can update weights; inference normally does not.

What Is an LLM?
A learned next-token prediction system

Two AI Traditions
Rules versus learned patterns

Training
Prediction error updates weights

Pretraining
Broad durable learning before

Home

Journey

Play

Glossary

Badge

Journey: section intro card and lesson rows.



What Is an LLM?

A learned next-token prediction system

A large language model is a learned system that predicts likely next tokens from the context it can see.

Prompt

Response


Token

Weight

Inference

VISUAL AID


What to picture



Lesson top: stage indicator, definition, and wrapped term chips.

VISUAL AID

What to picture



Prompt to Prediction

A prompt enters the current context; learned weights shape next-token probabilities.

CORE IDEA

Core idea

A large language model is a learned system that predicts likely next tokens from the context it can see.

WHERE IT HAPPENS

It sits between the current input context and the next response token.

WHY IT MATTERS

This definition lowers the mystery; the

Lesson visual aid: reusable SVG/CSS diagram with HTML caption.

BRAIN BRIDGE

Helpful, then limited

HELPFUL COMPARISON

Like a brain, many small signals can combine into useful patterns.


WHERE IT BREAKS

Unlike a brain, an LLM has no lived body, feelings, personal goals, or awareness.

TINY INTERACTION

Activate the prompt patterns

FEATURE CLOUD



Home Journey Play Glossary Badge

Brain Bridge: helpful comparison plus explicit limit.

CHECKPOINT

Which description is most accurate?

- A A learned prediction system
- B A conscious reader
- C A hand-written rulebook

REFLECT

Teach it back in one sentence.

Example: Inference uses the model's weights but does not rewrite them.

Answer checkpoint to continue

Home Journey Play Glossary Badge

Bottom nav: safe-area-aware frosted gradient over scrollable content.

BEFORE MORNING

1 / 25



1

Before Morning

2

3

4

5

6

What Is an LLM?

A learned next-token prediction system

A large language model is a learned system that predicts likely next tokens from the context it can see.

Prompt

Response

Token

Weight

Inference



Home



Journey



Play



Glossary



Badge

320px check: chips wrap and nav labels remain present.

PROGRESS

Prompt Life: Model Literate

Track your understanding of the prompt-to-response path.



0

of 25 lessons

0

of 3 play insights

0

of 10 exercises

0

of 13 Prompt Run

0

reflections



Home



Journey



Play



Glossary




Badge

Badge: cleaned copy, progress stats, and v0.6.0 label lower on the page.

What Is an LLM?

A learned next-token prediction system



Prompt to Prediction

A prompt enters the current context; learned weights shape next-token probabilities.

DEFINITION

A large language model is a learned system that predicts likely next tokens from the context it can see.

WHERE

It sits between the current input context and the next response token.

WHY


This definition lowers the mystery: the model is powerful because training shaped many weights, not because it has a hidden mind.

RELATIONSHIP

During inference, the current context enters a forward pass, one generated response token is chosen, and that token

Review route: one lesson card per print page.

Prompt to Prediction



Prompt to Prediction

A prompt enters the current context; learned weights shape next-token probabilities.

TRADITIONS

Rules and Learned Patterns



Rules and Learned Patterns

Symbolic AI follows explicit rules; deep learning uses learned parameters.

Visual aid gallery: reusable diagrams for all lessons.