

Before Morning Implementation Report

Date: June 5, 2026

This report wraps the implementation summary and UI screenshots for the v0.17.1 Before Morning polish pass.

What Changed

- Bumped the visible Badge version to `v0.17.1` .
- Removed the duplicate visible `Core idea` heading after the lesson-card pill.
- Removed repeated path-label chips from individual Journey rows.
- Revised the first seven Before Morning cards for clearer mobile-first teaching.
- Added/refined tiny interactions for prompt trace, rule/pattern sorting, training steps, pretraining myths, overfitting curves, fine-tuning contrast, and alignment grouping.
- Added internal exercise metadata for the revised lesson architecture.
- Added a text-only future Image 2 asset plan. No generated PNG assets were created.

Verification

Passed

`npm run typecheck`

Passed

`npm run build`

Passed

`npm run audit:answers`

Passed

`npm run build:pages`

67

audited answer surfaces

The build still reports the existing Vite large-chunk warning.

Files Changed

- `src/main.tsx`
- `src/data/content.ts`
- `src/data/exercises.ts`
- `src/components/VisualAids.tsx`
- `src/styles/global.css`
- `docs/REVIEW_NOTES.md`
- `docs/stage-audits/v0-17-before-morning/IMAGE_ASSET_PLAN.md`
- `docs/stage-audits/v0-17-before-morning/IMPLEMENTATION_REPORT_V0_17_1.md`
- `docs/stage-audits/v0-17-before-morning/screenshot-index.md`

Next Three v0.2 Improvements

1. Create the four planned textless Image 2 assets and keep instructional text in HTML callouts.
2. Add focused regression tests for the new tiny interactions and answer randomization helper.
3. Add a repeatable mobile screenshot harness for 320px, 390px, and 430px checks.

Journey overview. First seven Journey rows no longer repeat path labels inside each card row.

MAIN PATH

Today in Prompt Life

Follow one prompt from before training through inference, generation, and model literacy.

Jump to a stage of the prompt's day.

1 Before Morning Model shaped before use	2 Morning Commute Text becomes numbers	3 Workday Layers process context	4 Decision Room One token is chosen
5 The Day Repeats Context grows and expires	6 Twilight: The Wider Landscape Wider AI landscape	7 Midnight Ledger Costs and shadows	8 New Dawn Better human use

PATH VIEW

Choose a path

All cards follow the full day. Essential is the shortest path. Deep adds mechanics. Ethics follows costs, choices, and human consequences.

All**Essential****Deep****Ethics**

Showing all 38 Journey cards across 8 stages.

GUIDED COMPARISONS

Run and compare

Use these when the map feels abstract: one guides a prompt through the model, the other compares how AI systems learn.

Prompt Run**How AI Learns**

 **Home** **Journey** **Play** **Glossary** **Badge**

What Is an LLM? Revised prompt trace showing the next-token cloud.

HELPFUL COMPARISON

Like a person using context to anticipate what word might come next.

WHERE IT BREAKS

A person has a body, goals, lived experience, and awareness. The model has learned weights and temporary computations.

TINY INTERACTION

Tap through one prompt

1. Context

2. Token cloud

3. Chosen

4. Appended

Next-token cloud

floor, room, and tiles are plausible. quantum and elephant are much weaker in this local context.

floor

room

tiles

counter

quantum

elephant

Context enters, a next-token cloud appears, floor is chosen, then the appended response becomes the next context.



Home



Journey



Play




Glossary



Badge

Rationalists vs Empiricists. Revised rules and learned patterns visual.



Training. Durable update step highlighted.

WHERE IT BREAKS

The model is not reflecting on mistakes. An optimizer changes numbers to reduce error.

TINY INTERACTION

Tap through the training loop

Predict

Compare

Loss

Update weights

Update weights: The optimizer changes weights.
This is the durable step.

Insight unlocked: this is training, not ordinary inference.

Predict, compare, loss, update weights. The update is the durable learning step.

CHECKPOINT

What makes training different from inference?

A Training writes the final answer

Home

Journey

Play

Glossary

Badge

Pretraining. Perfect-recall myth distinction.

Helpful, then limited

HELPFUL COMPARISON

Like broad education shaping later intuition.

WHERE IT BREAKS

Humans connect learning to lived experience. A model stores statistical structure in parameters.

TINY INTERACTION

Pattern learning or recall?

Broad pattern learning

Perfect recall myth

Myth to drop

Perfect recall is the myth. Seeing text during training does not mean the model stores every source as a searchable memory.

Pretraining is broad durable pattern learning, not a perfect memory of every source.

CHECKPOINT

During pretraining, what changes?

A Nothing changes



Home



Journey



Play



Glossary



Badge

Overfitting. Generalizing curve and held-out examples distinction.

HELPFUL COMPARISON

Like memorizing practice questions without understanding the principle.

WHERE IT BREAKS

The model is not consciously memorizing. Optimization can make it rely on brittle numerical patterns.

TINY INTERACTION

Choose the better curve

Overfit curve

Generalizing curve

training dots

held-out dots

New examples handled better

Insight unlocked: evaluation needs held-out examples.

The smoother curve may miss a few old dots, but it carries the pattern to held-out examples.

CHECKPOINT

Which is a sign of overfitting?



Home



Journey



Play



Glossary



Badge

Fine-Tuning. Durable training versus temporary context steering versus decoding choice.

Training changes output patterns.

TINY INTERACTION

Sort the steering moves

fine-tuning examples	Durable training
prompt instructions	Temporary context steering
retrieved PDF	Temporary context steering
sampling next token	Decoding choice
adapter weights	Durable training


Durable training	fine-tuning examples, adapter weights
Temporary context steering	prompt instructions, retrieved PDF
Decoding choice	sampling next token

Insight unlocked: Fine-tuning shapes future behavior. Prompting and RAG steer the current context. Sampling chooses one next token.

Fine-tuning is durable training. Prompting and RAG steer the current run. Sampling chooses the next token.

Home Journey Play Glossary Badge

Alignment. Durable shaping, runtime steering, and evaluation grouped.



Checkpoint randomization. Answer ordering still comes from the randomized choice system.

system prompt, policy filter, tool rule

Evaluation
safety test, red-team review

Insight unlocked: Alignment is a layered effort: some methods train behavior, some steer a run, and some test the system.

Alignment can use durable shaping, runtime steering, and evaluation. It shapes behavior, not moral agency.

CHECKPOINT

What does alignment try to do?

A Permanently solve all AI risk

B Shape model behavior toward human goals, instructions, and safety boundaries

C Guarantee every answer is true

D Make the model conscious

Not quite. Alignment is a set of technical and human processes, not consciousness or a perfect safety guarantee. Choose another answer to retry.

REFLECT

Teach it back in one sentence.

Home


Journey

Play

Glossary

Badge

320px Alignment spot check. Chips and bucket labels wrap at narrow width.



Badge version. Single badge remains and version appears under Start over.

PROGRESS

Prompt Life: Model Literate

Track your understanding of the prompt-to-response path.



0
of 38 lessons

0
of 3 play insights

0
of 16 exercises

0
of 13 Prompt Run

0
reflections

Badge criterion

Home

Journey

Play

Glossary

Badge