

Generated Asset Integration Report

Date: June 5, 2026

This report wraps the implementation summary and screenshots for the Before Morning generated-asset integration pass.

What Changed

- Bumped the visible app version to `v0.17.3`.
- Added `src/data/visualAssets.ts` as a registry for the four provided textless Before Morning PNG assets.
- Integrated the provided assets into `What Is an LLM?`, `Pretraining`, `Fine-Tuning`, and `Alignment`.
- Kept instructional labels, captions, callouts, accessible descriptions, and review metadata in HTML/app code.
- Updated the visual-aid review gallery with generated-asset filename, path, type, alt text, and text-handling metadata.
- Preserved Journey order, progress rules, checkpoint randomization, games, and the single badge model.

Files Changed

- `src/data/visualAssets.ts`
- `src/components/VisualAids.tsx`
- `src/data/content.ts`
- `src/main.tsx`
- `src/styles/global.css`
- `public/assets/generated/before-morning/` four provided PNG assets
- `docs/REVIEW_NOTES.md`
- `docs/stage-audits/v0-17-before-morning/README.md`
- `docs/stage-audits/v0-17-before-morning/card-inventory.md`
- `docs/stage-audits/v0-17-before-morning/recommendations.md`
- `docs/stage-audits/v0-17-before-morning/IMAGE_ASSET_PLAN.md`
- `docs/stage-audits/v0-17-before-morning/screenshot-index.md`
- `docs/stage-audits/v0-17-before-morning/GENERATED_ASSET_INTEGRATION_V0_17_3.md`
- `docs/stage-audits/v0-17-before-morning/screenshots/` v0.17.3 screenshots

Assets Integrated

ASSET	LESSON	VISUAL AID ID
<code>before-morning-llm-cloud.png</code>	What Is an LLM?	<code>before-morning-llm-cloud</code>
<code>before-morning-pretraining-landscape.png</code>	Pretraining	<code>before-morning-pretraining-landscape</code>
<code>before-morning-finetuning-path.png</code>	Fine-Tuning	<code>before-morning-finetuning-path</code>
<code>before-morning-alignment-garden.png</code>	Alignment	<code>before-morning-alignment-garden</code>

Visual Aid Behavior

The four generated-image visual aids render the provided PNG inside a stable 4:3 frame, then use app-rendered titles, captions, callouts, and key takeaways below the image. The visual-aid review gallery lists asset metadata and keeps the text-handling rule explicit.

Where `LLMs Fit` remains `ai-family-tree`, a coded SVG/HTML visual aid. No generated PNG was added for that topology card.

Accessibility Handling

- Each generated image has dedicated `alt` text in the asset registry.
- Each visual aid has an accessible description and HTML callouts.
- The review gallery exposes filename, path, visual type, alt text, and text-handling notes.
- The image frame includes a visible fallback note if an asset fails to load.
- No instructional text is baked into the PNG files.

Verification

Passed

`npm run typecheck`

Passed

`npm run build`

Passed

`npm run build:pages`

Passed

`npm run audit:answers`

68

audited answer surfaces

46

randomized surfaces

Browser QA passed for generated visual aids at 390px, Alignment at 320px, visual-aid gallery metadata, Learn mode, Preview mode, Review mode, Badge version, and console errors.

Browser console errors observed: none.

The build still reports the existing Vite large-chunk warning.

Known Issues

- The existing Vite large-chunk warning remains.
- The generated PNGs are 4:3 assets; the app intentionally preserves that ratio for stable mobile and PDF layout.

Next Three Recommended Steps

1. Add a lightweight visual regression script for the four generated visual-aid IDs at 320px and 390px.
2. User-test whether learners understand that image-backed visuals are conceptual scenes and that formal teaching text lives in HTML.
3. Clean stale old visual-aid IDs from historical v0.6 docs only if those docs are meant to stay current rather than archival.

Run Locally

1. `npm install`
2. `npm run dev -- --host 0.0.0.0`
3. Open `http://localhost:5173/` on this computer or the Vite network URL on a phone on the same Wi-Fi.

Visual gallery overview. The review route lists v0.17.3 and the first generated asset entry.

PROMPT LIFE V0.17.3

Visual Aid Gallery

Reusable SVG/CSS diagrams for the current lesson path, with mobile and PDF readability checks.

BEFORE-MORNING-LLM-CLOUD

Prompt to Prediction

Lesson: What Is an LLM?

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-llm-cloud.png

Type: generated-image



Prompt to Prediction

Score, choose, append, repeat

A textless model-cloud scene is paired with HTML callouts for context, learned

What Is an LLM? The LLM cloud asset renders with generated-image metadata and HTML callouts.

BEFORE-MORNING-LLM-CLOUD

Prompt to Prediction

Lesson: What Is an LLM?

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-llm-cloud.png

Type: generated-image



Prompt to Prediction

Score, choose, append, repeat

A textless model-cloud scene is paired with HTML callouts for context, learned weights, one generated token, and append-repeat.

- 1 Context enters** The current prompt and response-so-far become the visible input for this run.
- 2 Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- 3 One token is generated** The model chooses one response token from the probability cloud.

Pretraining. The pretraining landscape asset renders with textless art and HTML explanation.

BEFORE-MORNING-PRETRAINING-LANDSCAPE

Broad Pretraining

Lesson: Pretraining **Pattern:** generatedImage

Variant: generated-image

Asset: before-morning-pretraining-landscape.png

Type: generated-image



Broad Pretraining

Scale, not perfect recall

A textless folded landscape is paired with HTML callouts for many examples, repeated updates, broad patterns, and the perfect-recall limit.

- 1 Many examples flow** Pretraining repeats the training loop across enormous collections of examples.
- 2 Updates shape weights** Prediction errors drive repeated weight updates during training.
- 3 Broad patterns form** The model develops grammar, style, associations, and task patterns before normal use.
- 4 Not perfect recall** Broad exposure does

Fine-Tuning. The fine-tuning path asset renders inside the same 4:3 frame with app-rendered callouts.

BEFORE-MORNING-FINETUNING-PATH

Fine-Tuning Path

Lesson: Fine-Tuning

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-finetuning-path.png

Type: generated-image



Fine-Tuning Path

Durable adaptation after pretraining

A textless targeted-path scene is paired with HTML callouts for pretrained base, targeted data, future responses, and durable adaptation.

- 1 Pretrained base** Fine-tuning begins with a model already shaped by broad pretraining.
- 2 Targeted data** Examples, domain data, or preference data guide the additional training step.
- 3 Future responses** The training changes how the model is likely to behave later.
- 4 More than a prompt** Fine-tuning is durable adaptation, unlike one temporary prompt or retrieved document.

Alignment. The alignment garden asset renders with guardrail-oriented callouts.

BEFORE-MORNING-ALIGNMENT-GARDEN


Alignment Landscape

Lesson: Alignment **Pattern:** generatedImage

Variant: generated-image

Asset: before-morning-alignment-garden.png

Type: generated-image



Alignment Landscape

Shaping behavior, not conscience

A textless alignment garden is paired with HTML callouts for preferred behavior, guardrails, feedback and policies, and the no-conscience limit.

- 1 Preferred behavior** Training and system design can encourage more useful response patterns.
- 2 Guardrails** Runtime safeguards can reduce risky paths during use.
- 3 Feedback and policies** Human feedback, evaluations, policies, and filters can shape behavior.
- 4 No conscience** Alignment does not give the model moral understanding or a truth

320px Alignment check. The generated image and first callouts remain readable at a narrow mobile width.

BEFORE-MORNING-ALIGNMENT-GARDEN

Alignment Landscape

Lesson: Alignment

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-alignment-garden.png

Type: generated-image

Alignment Landscape

Shaping behavior, not conscience

A textless alignment garden is paired with HTML callouts for preferred behavior, guardrails, feedback and policies, and the no-conscience limit.

- 1 Preferred behavior** Training and system design can encourage more useful response patterns.
- 2 Guardrails** Runtime safeguards can reduce risky paths during use.

Badge version. The single badge remains Prompt Life: Model Literate, with **v0.17.3** right-aligned under Start over.

