

Visual Loading + Key Terms Collapse Report

Date: June 5, 2026

This report wraps the v0.17.4 implementation summary and screenshots for the generated-asset loading fix and the two-row Key terms chip collapse.

What Changed

- Bumped the visible app version to `v0.17.4`.
- Added a base-aware public asset helper so generated PNG paths work in local dev and GitHub Pages builds under `/promptlife/`.
- Updated the generated visual-asset registry to use that helper for the four existing Before Morning PNGs.
- Added a reusable `KeyTermsChips` lesson panel after the visual aid. It defaults to two rows, with `Show all terms` and `Show fewer`.
- Prioritized the `Where LLMs Fit` chip order so the most important topology terms appear before overflow.
- Kept `Where LLMs Fit` on the coded SVG/HTML `ai-family-tree` visual. No generated PNG was added for it.
- Compacted the lesson-stage timeline under 340px so the first visual aid appears sooner on the narrowest phone check.
- Updated review notes, screenshot index, and the v0.17.4 stage audit note.

Files Changed

- `src/utils/assetPath.ts`
- `src/data/visualAssets.ts`
- `src/components/VisualAids.tsx`
- `src/main.tsx`
- `src/styles/global.css`
- `docs/REVIEW_NOTES.md`
- `docs/stage-audits/v0-17-before-morning/GENERATED_ASSET_LOADING_FIX_V0_17_4.md`
- `docs/stage-audits/v0-17-before-morning/screenshot-index.md`
- `docs/stage-audits/v0-17-before-morning/screenshots/ v0.17.4 screenshots`
- `docs/stage-audits/v0-17-before-morning/prompt-life-v0-17-4-visual-loading-key-terms-report.html`
- `docs/stage-audits/v0-17-before-morning/prompt-life-v0-17-4-visual-loading-key-terms-report.pdf`

Asset Path Fix

`publicAssetPath(path)` now derives public asset URLs from `import.meta.env.BASE_URL`. That means the same registry produces the right URL shape for each build target:

TARGET	EXAMPLE GENERATED PNG PATH
Local dev	<code>/assets/generated/before-morning/before-morning-llm-cloud.png</code>
GitHub Pages build	<code>/promptlife/assets/generated/before-morning/before-morning-llm-cloud.png</code>

The generated-image fallback now reads: `Visual asset unavailable`. The callouts below still explain the concept.

Generated Image Load Status

4 / 4

Before Morning generated images loaded in local dev browser QA.

1448 x 1086

Natural image size observed for each generated asset.

0

Generated-asset 404s observed in local dev or mounted Pages-style preview.

Passed

Mounted Pages-style check served JS and generated PNGs from `/promptlife/`.

Chip Collapse Behavior

- Glossary chips now render through a semantic `Key terms` section after the visual aid.
- Collapsed mode clips the list after two rows with a soft fade.
- The toggle is a real button with `aria-expanded` and `aria-controls`.
- Expansion is local UI state only; it does not touch lesson progress, badge progress, or checkpoint state.
- At 320px and 390px, `Where LLMs Fit` now shows the coded visual before the Key terms panel, then keeps the chips tidy below it.

Where LLMs Fit Visual Status

`Where LLMs Fit` still uses the coded SVG/HTML AI family tree. It has no generated image element and no generated PNG asset. The branch interaction and callouts appear before the Key terms panel so the visual is no longer buried by glossary chips.

Verification

- `npm run typecheck` : passed.
- `npm run build` : passed with the existing Vite large-chunk warning.
- `npm run build:pages` : passed with the existing Vite large-chunk warning.
- `npm run audit:answers` : passed; 68 audited surfaces, 46 randomized, 22 fixed-order exclusions.
- Browser QA: passed for generated images, Key terms collapse/expand, coded AI family tree, visual gallery, 320px/390px/430px mobile sweeps, local dev, and mounted Pages-style preview.

Known Issues

- The existing Vite large-chunk warning remains.
- The mounted Pages-style check proves local build paths; the deployed GitHub Pages CDN and iPhone cache still need a fresh-device check after push/deploy.
- During repeated local dev reload automation, Vite logged one React `createRoot` HMR message. No generated asset 404s appeared, and the mounted Pages build had no console errors.

Next Three Recommended Steps

1. Add a lightweight Pages-build asset-path check that fails if generated PNGs resolve to root `/assets/...` in a GitHub Pages build.
2. Add a narrow visual-regression capture for `Where LLMs Fit` collapsed and expanded terms at 320px and 390px.
3. After deployment, test the iPhone URL with `?v=0174` or clear site data to avoid stale cached bundles.

Run Locally

From the repository root:

1. `npm install`
2. `npm run dev -- --host 0.0.0.0`
3. Open the local or network URL printed by Vite. For phone testing on the same Wi-Fi, use the network URL.

For a Pages-style local build smoke test, run `npm run build:pages` and serve `dist` under a `/promptlife/` mount.

What Is an LLM? Generated visual loads in the lesson at 390px.

BEFORE-MORNING-LLM-CLOUD

Prompt to Prediction

Lesson: What Is an LLM?

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-llm-cloud.png

Type: generated-image



Prompt to Prediction

Score, choose, append, repeat

A textless model-cloud scene is paired with HTML callouts for context, learned weights, one generated token, and append-repeat.

- 1 Context enters** The current prompt and response-so-far become the visible input for this run.
- 2 Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- 3 One token is generated** The model chooses one response token from the probability cloud.

Pretraining. Generated landscape visual loads in the lesson at 390px.

BEFORE-MORNING-PRETRAINING-LANDSCAPE

Broad Pretraining

Lesson: Pretraining

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-pretraining-landscape.png

Type: generated-image



Broad Pretraining

Scale, not perfect recall

A textless folded landscape is paired with HTML callouts for many examples, repeated updates, broad patterns, and the perfect-recall limit.

- 1 Many examples flow** Pretraining repeats the training loop across enormous collections of examples.
- 2 Updates shape weights** Prediction errors drive repeated weight updates during training.
- 3 Broad patterns form** The model develops grammar, style, associations, and task patterns before normal use.
- 4 Not perfect recall** Broad exposure does

Fine-Tuning. Generated path visual loads in the lesson at 390px.

BEFORE-MORNING-FINETUNING-PATH

Fine-Tuning Path

Lesson: Fine-Tuning

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-finetuning-path.png

Type: generated-image



Fine-Tuning Path

Durable adaptation after pretraining

A textless targeted-path scene is paired with HTML callouts for pretrained base, targeted data, future responses, and durable adaptation.

- 1 Pretrained base** Fine-tuning begins with a model already shaped by broad pretraining.
- 2 Targeted data** Examples, domain data, or preference data guide the additional training step.
- 3 Future responses** The training changes how the model is likely to behave later.
- 4 More than a prompt** Fine-tuning is durable adaptation, unlike one temporary prompt or retrieved document.

Alignment. Generated garden visual loads in the lesson at 390px.

BEFORE-MORNING-ALIGNMENT-GARDEN

Alignment Landscape


Lesson: Alignment

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-alignment-garden.png

Type: generated-image



Alignment Landscape

Shaping behavior, not conscience

A textless alignment garden is paired with HTML callouts for preferred behavior, guardrails, feedback and policies, and the no-conscience limit.

- 1 Preferred behavior** Training and system design can encourage more useful response patterns.
- 2 Guardrails** Runtime safeguards can reduce risky paths during use.
- 3 Feedback and policies** Human feedback, evaluations, policies, and filters can shape behavior.
- 4 No conscience** Alignment does not give the model moral understanding or a truth

Where LLMs Fit. Key terms collapse to two rows in a panel after the coded visual.

from data; deep learning is one branch.

- 3 **Generative AI** Systems create new media, including language, images, audio, code, or video.
- 4 **LLMs** Language/code generators that usually produce response tokens autoregressively.
- 5 **Diffusion and multimodal** Other generative branches can denoise patterns or work across media types.

Key takeaway:
An LLM is one kind of generative AI inside the broader AI family.

Key terms
Tap a term to open the glossary. [Show all terms](#)

AI Machine learning

Deep learning Generative AI

CORE IDEA

AI is the broad field of building systems that perform tasks associated with intelligence. Machine learning is a major branch of AI where systems learn patterns from data. Deep learning is a kind of machine learning that uses large neural networks. Generative AI creates new text, images, audio, code, or other media. LLMs are generative AI systems

Home **Journey** Play Glossary Badge

Where LLMs Fit expanded. The same chip list expands below the visual with a Show fewer control.

2

Machine learning Systems learn patterns from data; deep learning is one branch.

3

Generative AI Systems create new media, including language, images, audio, code, or video.

4

LLMs Language/code generators that usually produce response tokens autoregressively.

5

Diffusion and multimodal Other generative branches can denoise patterns or work across media types.

Key takeaway:

An LLM is one kind of generative AI inside the broader AI family.

Key terms

Tap a term to open the glossary.

Show fewer

AI

Machine learning

Deep learning

Generative AI

LLM

Diffusion model

Multimodal AI

Symbolic AI

Classical machine learning

Rule-based AI

Foundation model

CORE IDEA

Home

Journey

Play


Glossary

Badge

AI family tree. The coded SVG/HTML visual remains visible and was not replaced by a generated PNG.

VISUAL AID

What to picture



AI Family Tree

Where LLMs fit

AI branches into rule-based systems and machine learning; deep learning includes generative AI, including LLMs and diffusion models.

- 1 **AI** The broad field, not a synonym for LLM.
- 2 **Machine learning** Systems learn patterns from data; deep learning is one branch.
- 3 **Generative AI** Systems create new media, including language, images, audio, code, or video.
- 4 **LLMs** Language/code generators that usually produce response tokens autoregressively.
- 5 **Diffusion and multimodal** Other generative branches can denoise patterns or work across media types.

Key takeaway:

An LLM is one kind of generative AI inside the broader AI family.

Home

Journey

Play

Glossary

Badge

Visual gallery. Generated assets remain visible in the review route with metadata.

BEFORE-MORNING-LLM-CLOUD

Prompt to Prediction

Lesson: What Is an LLM?

Pattern: generatedImage

Variant: generated-image

Asset: before-morning-llm-cloud.png

Type: generated-image



Prompt to Prediction

Score, choose, append, repeat

A textless model-cloud scene is paired with HTML callouts for context, learned weights, one generated token, and append-repeat.

- 1 Context enters** The current prompt and response-so-far become the visible input for this run.
- 2 Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- 3 One token is generated** The model chooses one response token from the probability cloud.

320px check. The AI family tree reaches the first viewport before the Key terms panel at a narrow mobile width.

