

PROMPT LIFE INTERNAL AUDIT

v0.19 Morning Commute Stage Audit

Date: 2026-06-06

Stage: Morning Commute

Cards: Inference, Prompt vs Response, Tokenization, Token IDs, Embeddings, Vectors, Tensors.

Scope: audit-only. No live curriculum changes, generated assets, new games, heavy 3D libraries, progress changes, checkpoint randomization changes, or badge changes.

Result: The stage is accurate and in a sound order. The next pass should focus on visual specificity and tiny interactions, especially Vectors and Tensors.

Findings

WHAT IS STRONG

- Inference clearly says fixed weights are used, not durably updated.
- Prompt vs Response correctly separates given context from generated tokens.
- Embeddings are defined as learned starting vectors, not definitions or memories.
- Checkpoint feedback is clear and randomized.


WHAT NEEDS WORK

- Vectors and Tensors remain the most abstract cards.
- Tokenization should show uneven chunks more explicitly.
- Token IDs need a stronger lookup-key visual.
- Tensors should become a lightweight CSS 3D or coded axis inspector.


NEXT IMPLEMENTATION DIRECTION

Keep all seven cards. Repair coded visuals and tiny interactions before adding any generated Morning Commute PNGs. If one Image 2 asset is later generated, Embeddings is the best candidate because a textless lookup-table/meaning-cloud image could reduce abstraction while HTML keeps labels accessible.

Selected Screenshots



Journey top: stage navigation remains readable at 390px.



Morning Commute overview: actual card order is Inference, Prompt vs Response, Tokenization, Token IDs, Embeddings, Vectors, Tensors.

Mechanics Screens

VISUAL AID

What to picture



Forward Pass

Inference creates temporary states while durable weights stay fixed.

- 1 **Point 1** Current context enters the model.
- 2 **Point 2** Fixed weights are used, not updated.
- 3 **Point 3** Temporary hidden states lead to next-token scores.

Inference: accurate forward-pass visual; next pass should emphasize temporary activations over fixed weights.

The model repeatedly samples next tokens from probabilities.

PROMPT VS RESPONSE

The response grows one token at a time.

The model does not process the prompt once and then write a whole answer.

1. The current context enters the model.
2. The model predicts probabilities for the next token.
3. One token is chosen.
4. That token is appended to the context.
5. The model runs again.
6. The response grows one token at a time.

Prompt or Response?

- ☐ Prompt = given to the model.
- ☒ Response = generated by the model.
- ☐ Context = everything currently visible.

USER PROMPT

Prompt token: Give me one sentence with two household pets in conflict.

RESPONSE SO FAR

Response token: A



Home



Journey



Play



Glossary



Badge

Prompt vs Response: role distinction is clear; demo is tall on mobile.

TINY INTERACTION

Tap to split the response

TOKENIZER

Split response

A jealous dog chased a startled cat across the kitchen floor.

Tap once to split the generated response into model-readable chunks.

The model does not carry raw sentences through the transformer. It starts with token IDs.

Tokenization: interaction should show uneven chunks and punctuation/space cases.

CHECKPOINT

What is an embedding in this app?

A A learned starting vector for a token ID

B A finished sentence

C A secret English definition

D A human memory


Insight unlocked. An embedding is a learned starting vector before context reshapes it.

Embeddings: checkpoint reinforces learned starting vector.

Dense-Card Spot Checks

VISUAL AID

What to picture



Feature Vector

A vector is a list of numerical features, not a sentence.

- 1 **Point 1** Labels such as animal-ish, grammar role, tone, position, and context clue are simplified.
- 2 **Point 2** Real features are distributed across many dimensions.

Vectors at 320px: readable, but concept remains abstract and simplified labels need care.

VISUAL AID

What to picture



Tensor Block

Tensors organize many token vectors for layer-by-layer processing.

- 1 **Point 1** Token positions run down one axis.
- 2 **Point 2** Feature dimensions run across another axis.
- 3 **Point 3** Batch x tokens x features is an advanced shape note.

Tensors at 320px: readable; strongest candidate for CSS 3D or coded axis inspection.

Prompt to Prediction

Score, choose, append, repeat

A textless model-cloud scene is paired with HTML callouts for context, learned weights, one generated token, and append-repeat.

- 1 **Context enters** The current prompt and response-so-far become the visible input for this run.
- 2 **Weights shape probabilities** Learned weights, built earlier during training, shape next-token scores.
- 3 **One token is generated** The model chooses one response token from the probability cloud.
- 4 **Response grows** That token is appended, then the model runs again.

Key takeaway:

An LLM turns context into next-token probabilities.

LEARNING OBJECTIVE

Show current context flowing through learned weights into one generated token.

ASSET PATH

/assets/generated/before-morning/before-morning-llm-cloud.png

ALT TEXT

An abstract folded-paper prompt ribbon flows into a glowing model cloud, and one token-like bead exits.

ACCESSIBLE DESCRIPTION

A textless conceptual image shows a prompt

Regression QA: Before Morning generated assets still load in the review gallery.