

The Day Repeats Stage Audit

Documentation-only review of Autoregression, Context Window, RAG and Retrieval, Grounding, and Hallucinations. No live curriculum, progress, games, badge, dependency, generated asset, checkpoint randomization, or Glossary Dojo behavior was changed.

63 screenshots

390px primary QA

320px and 430px spot checks

Build passed

Stage Audited

The Day Repeats teaches how a response grows after one token is chosen: append the token, run again, grow temporary context, retrieve evidence when a RAG system is present, ground what can be supported, and recognize hallucinations when fluent output outruns evidence.

Stage focus

How responses grow and context changes.

Keep clear

Context is temporary, not permanent memory.

Cards audited

Autoregression, Context Window, RAG and Retrieval, Grounding, Hallucinations.

Screenshot safety

Lesson captures used Preview mode in a temporary browser profile.

Major Findings

- The stage order is correct: Autoregression -> Context Window -> RAG and Retrieval -> Grounding -> Hallucinations.
- Autoregression and Context Window are accurate but still use the older slim lesson schema.
- RAG and Retrieval is the strongest current card: retrieval plus context, not training or permanent memory.
- Grounding and Hallucinations handle evidence, truth, uncertainty, and intent carefully.
- The next pass should make the complete loop more visible across the stage.

Visual Findings

- RAG, Grounding, and Hallucinations remain readable at 320px and 390px.
- Autoregression and Context Window are readable but too generic for their teaching burden.
- No horizontal overflow was reported in the screenshot manifest.
- No required Image 2 asset is recommended before the coded/SVG implementation pass.


Specific Answers

1. Autoregression is partly distinct from Sampling; the next pass should show Sampling chooses and Autoregression loops.
2. Context Window mostly avoids permanent-memory confusion, but needs an explicit durable-vs-temporary field.
3. RAG is clearly retrieval plus context, not training.
4. Grounding is clearly evidence support, not a truth guarantee.
5. Hallucinations are framed as unsupported fluent output, not lying or consciousness.
6. Autoregression is the visual that most needs repair.
7. Autoregression most needs a tiny interaction, with Context Window close behind.
8. No cards are too technical; Autoregression and Context Window are more thin than technical.
9. No cards are out of order.
10. No required Image 2 asset is justified now; Hallucinations is only an optional later textless bridge-scene candidate.
11. The stage prepares learners for later risk and responsibility stages by connecting temporary context, retrieval, grounding, hallucinations, and human judgment.

Recommended Assets

Card	Coded SVG / HTML	Image 2	Tiny interaction
Autoregression	Choose-token, append, rerun, context grows.	None.	Tap-step append loop.
Context Window	Tray with prompt, history, retrieved text, response-so-far, and old context falling out.	None.	Push cards into a limited window.
RAG and Retrieval	Keep current lanes; add lane-highlight states.	None.	Tap prompt -> retriever -> notes -> context -> response.
Grounding	Claim-to-evidence support map.	None.	Match claims to evidence.
Hallucinations	Bridge with supported and unsupported claims.	Optional later textless bridge scene.	Mark missing supports.

Representative Screenshots



Stage overview with Autoregression current and the Day Repeats card sequence visible.

VISUAL AID

What to picture



Append and Repeat

The chosen token is appended, then the model runs again.

Key terms

[Show all terms](#)

Autoregression

Completion

CORE IDEA


Autoregression means response generation by repeated next-token prediction: choose one token, append it, and run again.



Autoregression is readable at 320px, but the loop should show response-so-far and context growth more directly.

VISUAL AID

What to picture



Temporary Context Window

Only visible context can influence the next token.

Key terms

[Show all terms](#)

Context window

Input context

CORE IDEA

A context window is the limited amount of tokens or media the model can currently use.

WHERE IT HAPPENS



Home

Journey

Play


Glossary

Badge

Context Window is readable at 320px, but should show old context falling out and retrieved/context card types more clearly.

VISUAL AID

What to picture



Open-Book Retrieval

Retrieval plus context, not training

Retrieved notes enter the context before response tokens are generated.

- 1 Ask** The user prompt starts the run.
- 2 Retrieve** A search system finds relevant outside material.
- 3 Add to context** Retrieved notes become temporary context tokens.
- 4 Generate** The model still generates response tokens one at a time.
- 5 Weights stay fixed** RAG does not normally update model weights.

Key takeaway:

RAG is retrieval plus context, not training.

Key terms

Tap a term to open the glossary.

[Show all terms](#)



Home



Journey



Play



Glossary




Badge

RAG is the current gold-standard visual in this stage: prompt, retriever, notes, context tray, response, and fixed weights.

VISUAL AID

What to picture



Evidence Anchor

Tying answers to evidence

Evidence enters context so the generated response can stay connected to sources.

- 1 Evidence** Retrieved documents, data, citations, or tool results can enter the current run.
- 2 Context** The evidence becomes visible input, not permanent model memory.
- 3 Answer** The response should stay connected to the evidence while still being reviewed.
- 4 Limit** Grounding helps, but it can fail when retrieval is poor or the model misuses evidence.

Key takeaway:

Grounding helps tie an answer to evidence; it does not guarantee truth.




[Home](#)[Journey](#)[Play](#)[Glossary](#)[Badge](#)

Grounding remains readable and correctly says evidence helps but does not guarantee truth.

VISUAL AID

What to picture



Unsupported Bridge

Fluent is not always grounded

A response can look smooth while missing evidence supports.

- 1 Fluent surface** The output may read smoothly and confidently.
- 2 Missing support** Some claims may lack evidence, citation, or retrieved context.
- 3 Generation** Likely-token generation is not the same as truth verification.
- 4 Review** Grounding, retrieval, uncertainty, and human review reduce risk but do not erase it.

Key takeaway:

Fluency is not evidence.

Key terms

[Show all terms](#)



[Home](#)[Journey](#)[Play](#)[Glossary](#)[Badge](#)

Hallucinations visual keeps the key distinction clear: fluency is not evidence.

CHECKPOINT

What does RAG usually do?

- A** Permanently updates model weights
- B** Retrieves outside information and places it into the current context
- C** Guarantees every answer is true
- D** Makes the model aware of documents

Not quite. RAG is retrieval plus context. It is not training, permanent memory, consciousness, or a truth guarantee. Choose another answer to retry.

REFLECT

Teach it back in one sentence.

Reflection is available in Learn mode.



Home



Journey



Play



Glossary



Badge

Wrong-answer feedback remains non-shaming and reinforces RAG as retrieval plus context.

Verification

- `npm run typecheck`: passed.
- `npm run build`: passed with the existing Vite large-chunk warning.
- `npm run build:pages`: passed with the existing Vite large-chunk warning.
- `npm run audit:answers`: passed.
- `stage-audit.json`: valid JSON.

Known Issues

- The existing Vite large-chunk warning remains.
- Autoregression and Context Window still need implementation repairs.
- RAG, Grounding, and Hallucinations interactions remain generic until the next implementation pass.

Recommended Next Prompt

Implement a v0.22.1 The Day Repeats repair pass. Do not add new Journey cards, games, generated PNG assets, heavy 3D libraries, dependencies, badge changes, progress-rule changes, checkpoint-randomization changes, Glossary Dojo changes, or Journey order changes. Keep the order: Autoregression, Context Window, RAG and Retrieval, Grounding, Hallucinations. Convert Autoregression and Context Window from the older slim schema to the richer lesson architecture with explicit core explanation, model lifecycle, durable-vs-temporary notes, prompt-vs-response/current-context notes, misconceptions, and clearer feedback. Preserve the strong RAG, Grounding, and Hallucinations distinctions while giving them more specific tiny interactions. Update coded SVG/HTML visuals so Autoregression shows choose-token -> append -> rerun, Context Window shows prompt/conversation/retrieved/response-so-far cards plus old context falling out, RAG keeps retrieval-plus-context-not-training, Grounding shows claim-to-evidence support, and Hallucinations shows fluent output outrunning support. Keep all text in HTML, keep visuals readable at 320px and 390px, and run `npm run typecheck`, `npm run build`, `npm run build:pages`, and `npm run audit:answers`.