

Prompt Life v0.25.2 New Dawn Implementation Report

Date: 2026-06-07 | Project: Prompt Life | Stage: New Dawn

Bottom line: The five existing New Dawn cards were polished with clearer coded visuals, diagnostic tiny interactions, a stronger final synthesis chain, and source-aware wording. No cards, games, generated PNG assets, dependencies, progress rules, checkpoint randomization, badge model, or Glossary Dojo behavior were changed.

What Changed

Card	Improvement
Benefits Worth Taking Seriously	Benefit-tier coded visual, evidence-tier sort, new Benefits glossary term, safer checkpoint wording with human review.
Human-Centered AI	Accountability-flow visual and student-support-note scenario choice that keeps review and responsibility with people.
Better AI Is a Choice	Better-AI control panel and lever selection for an internal policy assistant scenario.
Effective Prompting from Model Literacy	Expanded context tray plus prompt-builder with audience/context, constraints, examples, evidence, uncertainty, format, and review.
Model Literate Synthesis	Full-chain map and tap-order capstone from prompt context through tokens, IDs, embeddings, hidden states, logits, probabilities, sampling, RAG/grounding, review, and accountability.

Files Changed

- src/data/content.ts
- src/main.tsx
- src/components/VisualAids.tsx
- src/styles/global.css
- docs/REVIEW_NOTES.md
- docs/stage-audits/v0-25-new-dawn/source-review.md
- docs/stage-audits/v0-25-new-dawn/IMAGE_ASSET_PLAN.md
- docs/stage-audits/v0-25-new-dawn/TINY_INTERACTION_PLAN.md
- docs/stage-audits/v0-25-new-dawn/IMPLEMENTATION_REPORT_V0_25_2.md
- docs/stage-audits/v0-25-new-dawn/implementation-screenshots/*
- docs/stage-audits/v0-25-new-dawn/prompt-life-v0-25-2-new-dawn-implementation-report.html
- docs/stage-audits/v0-25-new-dawn/prompt-life-v0-25-2-new-dawn-implementation-report.pdf

Visual Changes

- Benefits uses evidence-tier buckets: useful now with review, plausible with safeguards, speculative/hype.
- Human-Centered AI uses a concrete accountability flow.
- Better AI uses a lever/control-panel visual.
- Effective Prompting uses a richer context tray.
- Synthesis uses a full-chain map rather than the previous generic map/compass/lantern visual.

Tiny Interactions

- `benefit-tier-sort` : sort claims by confidence.
- `human-centered-scenario` : choose the accountable deployment.
- `better-ai-levers` : select better design/governance levers.
- `prompt-builder` : build a better current-context prompt.
- `synthesis-chain` : put the prompt-to-accountability story in order.

Glossary Updates

Added **Benefits**. Updated New Dawn learning-path grouping and key-term display priority so Benefits, Uncertainty, Source review, Human review, Model literacy, and final-chain terms appear sensibly without mobile chip overflow.

Source Discipline

No precise statistics or new external claims were added. Benefits and mitigations use cautious language such as can, may, with review, depends on, when appropriate, and should be evaluated.

Badge and Synthesis

The one-badge model remains: **Prompt Life: Model Literate**. The synthesis card now better prepares the learner for that badge by connecting machine mechanics with human review, governance, and accountability.

Verification

- `npm run typecheck` : passed.
- `npm run build` : passed, with the existing Vite large-chunk warning.
- `npm run build:pages` : passed, with the existing Vite large-chunk warning.
- `npm run audit:answers` : passed. Journey checkpoints remain randomized by answer identity.
- Screenshot manifest: 16 screenshots, 390px and 320px, horizontal overflow: false.

Known Issues

- The app still has the existing Vite large bundle warning.

- New Dawn remains internally source-reviewed but not learner-cited.
- Tap/cycle interactions are used instead of drag-only controls for mobile and keyboard accessibility.
- Image 2 asset generation should wait until the coded New Dawn relationships are reviewed.


Readiness

New Dawn is ready for whole-app QA and game review. It is not yet ready for new Image 2 asset generation beyond planning.

Screenshots

New Dawn overview

New Dawn overview in Journey | no horizontal overflow detected



Benefits visual

Benefits tier visual | no horizontal overflow detected

VISUAL AID

What to picture

Benefit confidence tiers

Useful now + review Access Translate Summarize Search/RAG	Plausible + safeguards Tutoring Triage Workflow	Speculative / hype Solves all Replaces judgment No costs
--	---	--

Real help needs evidence + review

Benefit Tiers

Benefits without utopia

AI benefits are easiest to trust when claims are tiered by evidence, context, safeguards, and human review.

1

Useful now with review Accessibility, translation support, summarization, search/RAG, drafting, and coding assistance can help under human review.

2

Plausible with safeguards Tutoring support, research triage, brainstorming, and workflow support should be evaluated in context.

3

Speculative / hype Broad utopia, replacement, or no-cost claims should not be stated as fact.

Key takeaway:

Benefits can be real and still bounded.

Home

Journey

Play

Glossary

Badge

file:///Users/kevinhegg/Desktop/promptlife/docs/stage-audits/v0-25-new-dawn/prompt-life-v0-25-2-new-dawn-implementation-report.html

5/19

Benefits interaction

Benefits evidence-tier interaction success state | no horizontal overflow detected

TINY INTERACTION

Sort benefit strength

AI can assist drafts, search, translation, and access under human review.
Useful now with review

AI will automatically solve education.
Speculative / hype

AI has human judgment now.
Speculative / hype

AI benefits mean no costs matter.
Speculative / hype

Useful now with review
AI can assist drafts, search, translation, and access under human review.

Plausible with safeguards
empty

Speculative / hype
AI will automatically solve education., AI has human judgment now., AI benefits mean no costs matter.

Insight strengthened. Benefits can be real and bounded when humans review evidence and context.

Home

Journey

Play

Glossary

Badge

Human-Centered AI visual

Human-centered accountability visual | no horizontal overflow detected

VISUAL AID

What to picture



Accountability Flow

Tools should serve dignity

A human-centered deployment treats AI output as support inside a decision context, then keeps review, governance, and accountable action with people.

- 1 Human judgment** People remain accountable for high-stakes decisions and institutional use.
- 2 Dignity** Speed, profit, and automation should not outrank persons, learning, or relationships.
- 3 Model limit** A model can sound ethical without moral understanding.

Key takeaway:

Powerful tools still need human purpose.

Key terms

Tap a term to open the glossary.

[Show all terms](#)



Home



Journey



Play



Glossary



Badge

Human-Centered AI interaction

Human-centered scenario correct feedback | no horizontal overflow detected

TINY INTERACTION

Center the human purpose

An AI tool helps summarize student support notes. Which deployment keeps people accountable?

AI summary

student support context

human review

accountable action

Use summary as a draft; human reviews before action.

Automatically decide support needs with no review.

Hide model limitations from staff.

Treat fluent text as proof.

Insight strengthened. Human-centered AI keeps responsibility with people and institutions.

Choose the deployment that keeps people, review, and institutional accountability visible.

Home

Journey

Play

Glossary


Badge

Better AI visual

Better-AI control panel visual | no horizontal overflow detected

VISUAL AID

What to picture



The diagram shows a 'Better-AI control panel' with several buttons: Provenance, Privacy, Efficient, RAG, Review, Eval, Labor, Sustain, Govern, and Task fit. A green arrow points from 'Govern' to 'Task fit'.

Better-AI Control Panel

Responsible AI is chosen

AI outcomes are shaped by design, deployment, governance, incentives, and institutional choices.

- 1 Technical choices** Smaller models, efficient inference, distillation, RAG, and better hardware use can fit some tasks.
- 2 Data choices** Provenance, consent, licensing, creator compensation, and privacy-preserving deployment matter.
- 3 Institutional choices** Human review, policy, labor transition planning, public-interest models, and independent evaluation shape outcomes.

Key takeaway:
Harms are shaped by choices, not destiny.

Key terms [Show all terms](#)


Home


Journey



Play


Glossary


Badge

Better AI interaction

Better-AI levers success state | no horizontal overflow detected




Effective Prompting visual

Prompting context tray visual | no horizontal overflow detected

VISUAL AID

What to picture



The diagram illustrates the components of a prompt and their relationship to the current context and response. Eight components (Task, Context, Limits, Example, Evidence, Uncertain, Format, and Review) are shown in yellow boxes, with arrows pointing to a central blue box labeled 'Current context this run'. An arrow from this box points to a yellow box labeled 'Response'.

Context Tray

Prompting steers this run

Good prompts pack task, context, constraints, examples, evidence needs, uncertainty, review, and format into the current run.

- 1 Prompt parts** Task, context, constraints, examples, evidence needs, uncertainty, review, and format shape the current input.
- 2 Evidence and uncertainty** Source needs, retrieved context, and uncertainty requests can improve reviewability.
- 3 Boundary** Prompting usually changes context, not weights.

Key takeaway:
Prompting is context design for one run.

Key terms [Show all terms](#)
Tap a term to open the glossary


Home


Journey



Play


Glossary


Badge


Effective Prompting interaction

Prompt builder complete state with constraints and examples | no horizontal overflow detected




Synthesis visual

Synthesis full-chain visual | no horizontal overflow detected




Synthesis interaction

Synthesis chain complete state | no horizontal overflow detected




Synthesis visual at 320px

Synthesis full-chain visual at 320px | no horizontal overflow detected



Synthesis interaction at 320px

Synthesis chain interaction at 320px | no horizontal overflow detected



Synthesis key takeaway at 320px

Synthesis key takeaway clearance at 320px | no horizontal overflow detected

Model literacy connects prompt, tokens, IDs, embeddings, hidden states, logits, probabilities, sampling, append-and-repeat generation, RAG, grounding, review, and accountability.

- 1 Machine chain** Prompt context becomes tokens and IDs, then embeddings, hidden states, logits, probabilities, and sampled response tokens.
- 2 Evidence chain** Context windows, RAG, and grounding can add evidence for this run, but they do not guarantee truth.
- 3 Human chain** Benefits, costs, review, governance, and accountability remain human responsibilities.

Key takeaway:

Mechanics matter, and humans remain responsible.

Key terms

[Show all terms](#)

Model literacy

Prompt

Token

Token ID

CORE IDEA



Home



Journey



Play



Glossary



Badge

Synthesis final action

Synthesis final action clear of bottom nav | no horizontal overflow detected

A Mechanics alone settle every use decision

B Fluent output proves wisdom

C Mechanics matter, and humans remain responsible

D Ethics replaces learning the mechanics

Insight unlocked: model literacy joins mechanism with judgment.

REFLECT

Teach it back in one sentence.

Example: Inference uses the model's weights but does not rewrite them.

Finish and view badge

 Home

 Journey

 Play

 Glossary

 Badge

Badge single model

Badge screen still shows one badge model | no horizontal overflow detected

